

MLP in MNIST dataset -- Keras Implementation

Importing packages

In [1]:

```
# if you keras is not using tensorflow as backend set "KERAS_BACKEND=tensorflow" use this c
import tensorflow as tf
from keras.utils import np_utils
from keras.datasets import mnist
import seaborn as sns
from keras.initializers import RandomNormal
import matplotlib.pyplot as plt
```

Using TensorFlow backend.

Function for plotting dynamic graph

In [2]:

```
%matplotlib notebook
import matplotlib.pyplot as plt
import numpy as np
import time
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    fig.canvas.draw()
```

Loading mnist data

In [3]:

```
# the data, shuffled and split between train and test sets
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

In [4]:

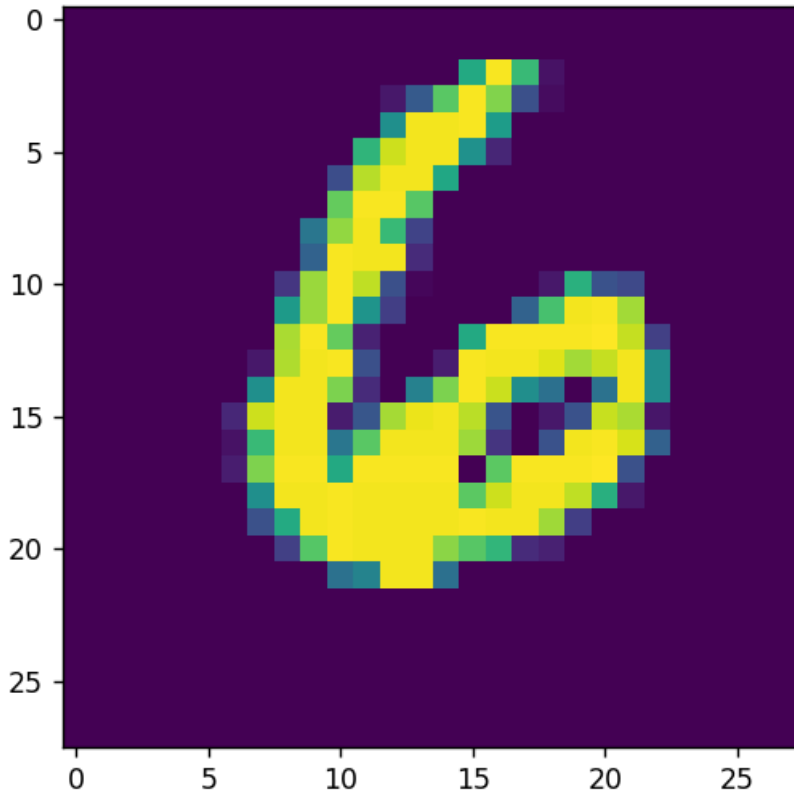
```
print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d, %d)" % (X_train.shape[1], X_train.shape[2]))
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d, %d)" % (X_test.shape[1], X_test.shape[2]))
```

Number of training examples : 60000 and each image is of shape (28, 28)
 Number of training examples : 10000 and each image is of shape (28, 28)

In [5]:

```
plt.imshow(X_train[299])
```

Figure 1



Out[5]:

```
<matplotlib.image.AxesImage at 0x1f0adaa87f0>
```

In [6]:

```
# if you observe the input shape its 2 dimensional vector  
# for each image we have a (28*28) vector  
# we will convert the (28*28) vector into single dimensional vector of 1 * 784
```

```
X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])  
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])
```

In [7]:

```
print(X_train.shape, y_train)  
print(X_test.shape, y_test)
```

```
(60000, 784) [5 0 4 ... 5 6 8]  
(10000, 784) [7 2 1 ... 4 5 6]
```

In [8]:

```
# after converting the input images from 3d to 2d vectors
```

```
print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d)"%  
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d)"%(
```

Number of training examples : 60000 and each image is of shape (784)

Number of training examples : 10000 and each image is of shape (784)

sample data

In [9]:

```
# An example data point
print(X_train[0])
```

```
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  3  18  18  18  126  136  175  26  166  255
247 127  0  0  0  0  0  0  0  0  0  0  0  0  30  36  94  154
170 253 253 253 253 253 225 172 253 242 195  64  0  0  0  0  0  0
  0  0  0  0  0  49 238 253 253 253 253 253 253 253 251 93  82
 82  56  39  0  0  0  0  0  0  0  0  0  0  0  18 219 253
253 253 253 253 198 182 247 241  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  80 156 107 253 253 205 11  0  43 154
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  14  1 154 253 90  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  139 253 190  2  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  11 190 253 70  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  35 241
225 160 108  1  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  81 240 253 253 119 25  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  45 186 253 253 150 27  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  16 93 252 253 187
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  249 253 249 64  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  46 130 183 253
253 207  2  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  39 148 229 253 253 253 250 182  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  24 114 221 253 253 253
253 201 78  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0 23 66 213 253 253 253 253 198 81  2  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  18 171 219 253 253 253 253 195
80  9  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
55 172 226 253 253 253 253 244 133 11  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  136 253 253 253 212 135 132 16
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0]
```

Normalizing pixel values

In [9]:

```
# if we observe the above matrix each cell is having a value between 0-255  
# before we move to apply machine learning algorithms lets try to normalize the data  
#  $X \Rightarrow (X - X_{min}) / (X_{max} - X_{min}) = X / 255$   
  
X_train = X_train/255  
X_test = X_test/255
```

After normalization (pixel value ranges from 0-1)

In [10]:

```
# example data point after normlizing
print(X_train[0])
```

```
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  3  18  18  18  126  136  175  26  166  255
247 127  0  0  0  0  0  0  0  0  0  0  0  0  30  36  94  154
170 253 253 253 253 253 225 172 253 242 195  64  0  0  0  0  0  0
  0  0  0  0  0  49 238 253 253 253 253 253 253 253 251 93  82
 82  56  39  0  0  0  0  0  0  0  0  0  0  0  18 219 253
253 253 253 253 198 182 247 241  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  80 156 107 253 253 205 11  0  43 154
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  14  1 154 253 90  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  139 253 190  2  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  11 190 253 70  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  35 241
225 160 108  1  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  81 240 253 253 119 25  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  45 186 253 253 150 27  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  16 93 252 253 187
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  249 253 249 64  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  46 130 183 253
253 207  2  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  39 148 229 253 253 253 250 182  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  24 114 221 253 253 253
253 201 78  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0 23 66 213 253 253 253 253 198 81  2  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  18 171 219 253 253 253 253 195
80  9  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
55 172 226 253 253 253 253 244 133 11  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  136 253 253 253 212 135 132 16
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0]
```

Encoding class labels

In [11]:

```
# here we are having a class number for each image
print("Class label of first image :", y_train[0])

# Lets convert this into a 10 dimensional vector
# ex: consider an image is 5 convert it into 5 => [0, 0, 0, 0, 0, 1, 0, 0, 0, 0]
# this conversion needed for MLPs

Y_train = np_utils.to_categorical(y_train, 10)
Y_test = np_utils.to_categorical(y_test, 10)

print("After converting the output into a vector : ",Y_train[0])
```

Class label of first image : 5

After converting the output into a vector : [0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]

Model 1

In [12]:

```
from keras.models import Sequential
from keras.layers import Dense, Activation, Dropout

# some model parameters

output_dim = 10
input_dim = X_train.shape[1]

batch_size = 128
nb_epoch = 20
```

In [13]:

```

model1 = Sequential()
model1.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.01)))
model1.add(Dropout(0.5))

model1.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.01)))
model1.add(Dropout(0.5))

model1.add(Dense(output_dim, activation='softmax'))

print(model1.summary())

model1.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model1.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

```

WARNING:tensorflow:From C:\Users\vansh\Anaconda3\envs\Deep_learning\lib\site-packages\keras\backend\tensorflow_backend.py:74: The name tf.get_default_graph is deprecated. Please use tf.compat.v1.get_default_graph instead.

WARNING:tensorflow:From C:\Users\vansh\Anaconda3\envs\Deep_learning\lib\site-packages\keras\backend\tensorflow_backend.py:517: The name tf.placeholder is deprecated. Please use tf.compat.v1.placeholder instead.

WARNING:tensorflow:From C:\Users\vansh\Anaconda3\envs\Deep_learning\lib\site-packages\keras\backend\tensorflow_backend.py:4115: The name tf.random_normal is deprecated. Please use tf.random.normal instead.

WARNING:tensorflow:From C:\Users\vansh\Anaconda3\envs\Deep_learning\lib\site-packages\keras\backend\tensorflow_backend.py:133: The name tf.placeholder_with_default is deprecated. Please use tf.compat.v1.placeholder_with_default instead.

WARNING:tensorflow:From C:\Users\vansh\Anaconda3\envs\Deep_learning\lib\site-packages\keras\backend\tensorflow_backend.py:3445: calling dropout (from tensorflow.python.ops.nn_ops) with keep_prob is deprecated and will be removed in a future version.

Instructions for updating:

Please use `rate` instead of `keep_prob`. Rate should be set to `rate = 1 - keep_prob`.

WARNING:tensorflow:From C:\Users\vansh\Anaconda3\envs\Deep_learning\lib\site-packages\keras\backend\tensorflow_backend.py:4138: The name tf.random_uniform is deprecated. Please use tf.random.uniform instead.

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 512)	401920
dropout_1 (Dropout)	(None, 512)	0
dense_2 (Dense)	(None, 128)	65664
dropout_2 (Dropout)	(None, 128)	0
dense_3 (Dense)	(None, 10)	1290

Total params: 468,874
Trainable params: 468,874
Non-trainable params: 0

None

WARNING:tensorflow:From C:\Users\vansh\Anaconda3\envs\Deep_learning\lib\site-packages\keras\optimizers.py:790: The name tf.train.Optimizer is deprecated. Please use tf.compat.v1.train.Optimizer instead.

WARNING:tensorflow:From C:\Users\vansh\Anaconda3\envs\Deep_learning\lib\site-packages\tensorflow\python\ops\math_grad.py:1250: add_dispatch_support.<locals>.wrapper (from tensorflow.python.ops.array_ops) is deprecated and will be removed in a future version.

Instructions for updating:

Use tf.where in 2.0, which has the same broadcast rule as np.where

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 5s 80us/step - loss: 12.5850
- acc: 0.2181 - val_loss: 9.8416 - val_acc: 0.3886

Epoch 2/20

60000/60000 [=====] - 3s 48us/step - loss: 10.1937
- acc: 0.3669 - val_loss: 8.5453 - val_acc: 0.4693

Epoch 3/20

60000/60000 [=====] - 3s 49us/step - loss: 9.3444 -
acc: 0.4199 - val_loss: 8.0839 - val_acc: 0.4981

Epoch 4/20

60000/60000 [=====] - 3s 47us/step - loss: 9.1723 -
acc: 0.4304 - val_loss: 8.0043 - val_acc: 0.5030

Epoch 5/20

60000/60000 [=====] - 3s 48us/step - loss: 8.3629 -
acc: 0.4807 - val_loss: 7.0649 - val_acc: 0.5615

Epoch 6/20

60000/60000 [=====] - 3s 47us/step - loss: 8.1204 -
acc: 0.4960 - val_loss: 7.0859 - val_acc: 0.5599

Epoch 7/20

60000/60000 [=====] - 3s 48us/step - loss: 7.8107 -
acc: 0.5152 - val_loss: 6.6105 - val_acc: 0.5896

Epoch 8/20

60000/60000 [=====] - 3s 47us/step - loss: 7.9137 -
acc: 0.5088 - val_loss: 6.6857 - val_acc: 0.5851

Epoch 9/20

60000/60000 [=====] - 3s 49us/step - loss: 8.2320 -
acc: 0.4890 - val_loss: 7.5361 - val_acc: 0.5324

Epoch 10/20

60000/60000 [=====] - 3s 49us/step - loss: 8.0205 -
acc: 0.5021 - val_loss: 7.1956 - val_acc: 0.5535

Epoch 11/20

60000/60000 [=====] - 3s 48us/step - loss: 8.0914 -
acc: 0.4978 - val_loss: 6.2347 - val_acc: 0.6126

Epoch 12/20

60000/60000 [=====] - 3s 48us/step - loss: 7.3911 -
acc: 0.5411 - val_loss: 6.2756 - val_acc: 0.6105

Epoch 13/20

60000/60000 [=====] - 3s 48us/step - loss: 7.5374 -
acc: 0.5321 - val_loss: 8.4569 - val_acc: 0.4751

Epoch 14/20

60000/60000 [=====] - 3s 48us/step - loss: 6.9246 -
acc: 0.5702 - val_loss: 5.3398 - val_acc: 0.6686

Epoch 15/20

60000/60000 [=====] - 3s 48us/step - loss: 6.5722 -
acc: 0.5921 - val_loss: 5.6229 - val_acc: 0.6509

```
Epoch 16/20
60000/60000 [=====] - 3s 49us/step - loss: 6.5891 -
acc: 0.5909 - val_loss: 5.2502 - val_acc: 0.6742
Epoch 17/20
60000/60000 [=====] - 3s 49us/step - loss: 6.2670 -
acc: 0.6110 - val_loss: 4.5675 - val_acc: 0.7165
Epoch 18/20
60000/60000 [=====] - 3s 48us/step - loss: 6.1065 -
acc: 0.6209 - val_loss: 4.6888 - val_acc: 0.7088
Epoch 19/20
60000/60000 [=====] - 3s 48us/step - loss: 6.3317 -
acc: 0.6070 - val_loss: 7.4789 - val_acc: 0.5359
Epoch 20/20
60000/60000 [=====] - 3s 48us/step - loss: 6.6072 -
acc: 0.5899 - val_loss: 5.4056 - val_acc: 0.6646
```

train and test loss vs Epochs

In [16]:

```
score = model1.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=0)

# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

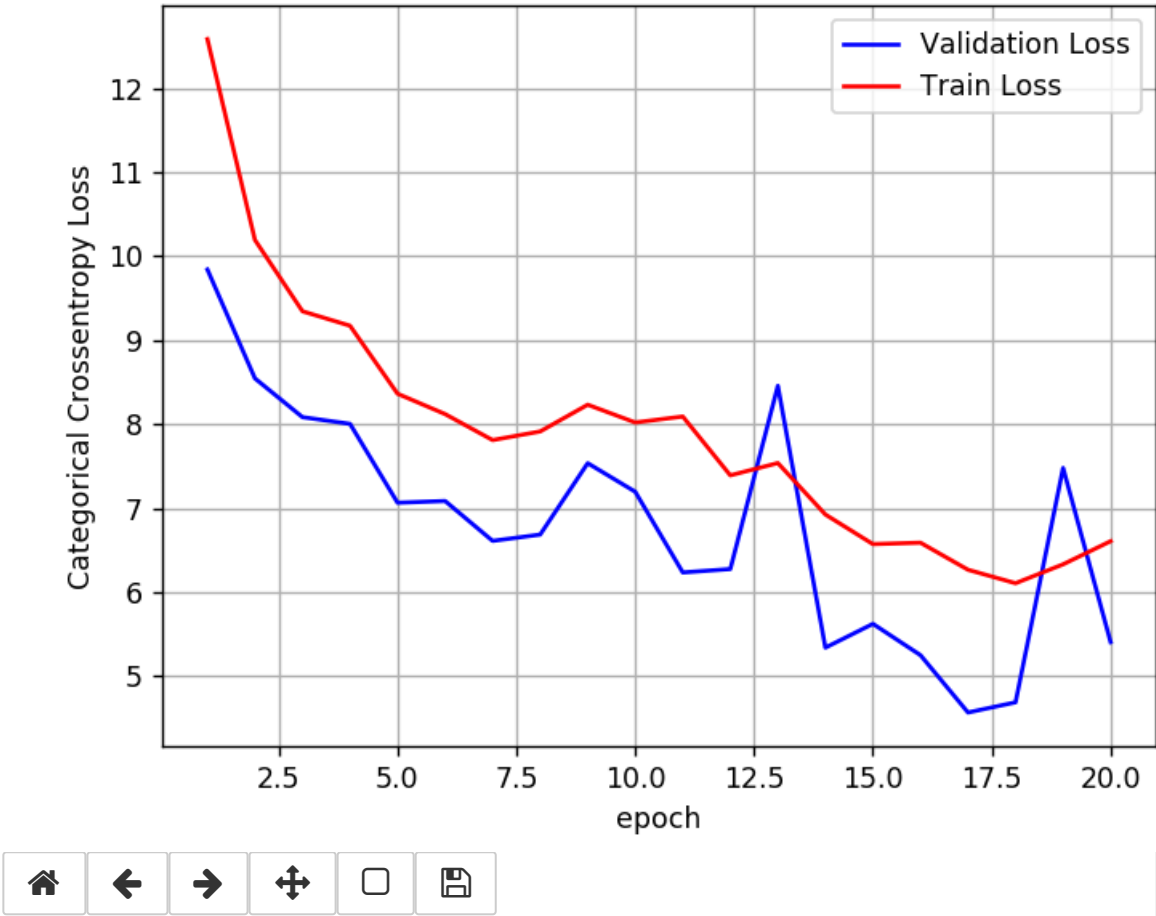
# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 5.405633866882324

Test accuracy: 0.6646

Figure 4



Violin plots of weights

In [17]:

```

w_after = model1.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

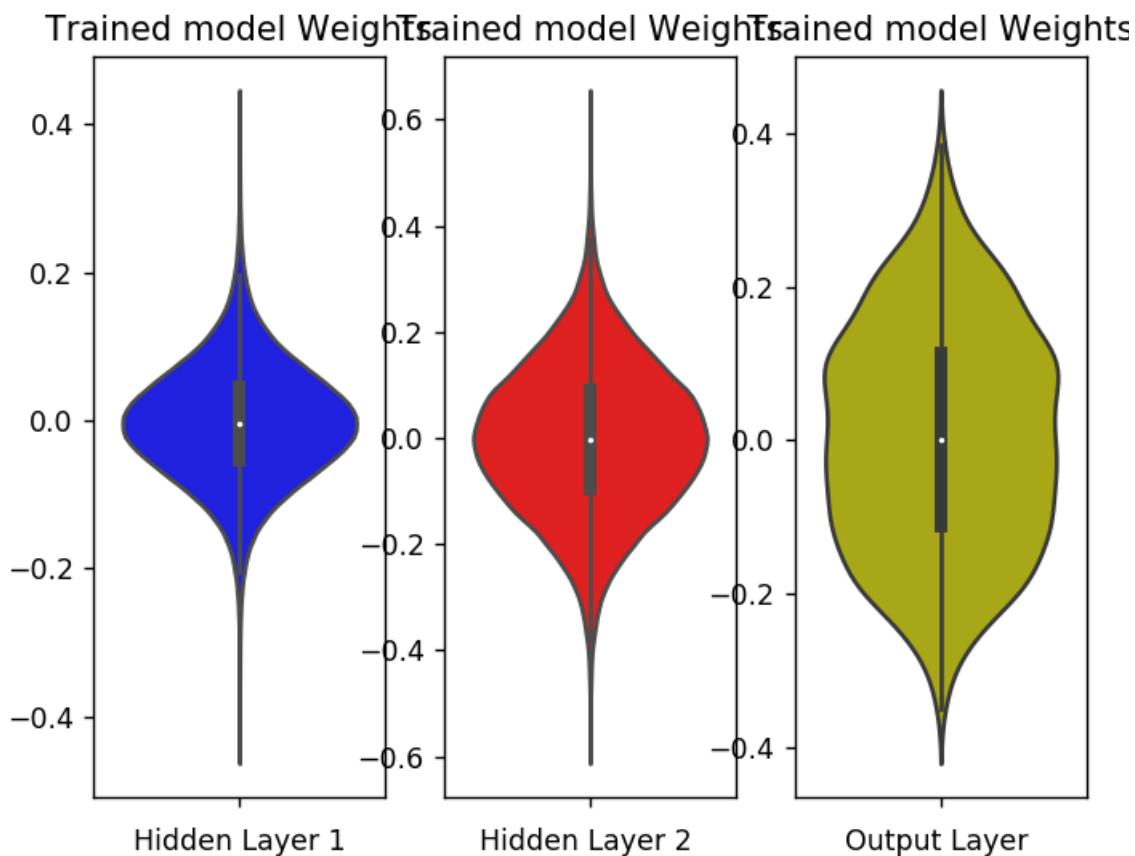
fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```

Figure 5



1. Weights at hidden layer 1 ranges (-0.2 to 0.2).
2. Weights at hidden layer 2 have more variance (-0.4 to 0.4).
3. Weights at output layer are off-centered (-0.4 to 0.4).

Model 2

In [18]:

```

model2 = Sequential()
model2.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.01)))
model2.add(Dense(128, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.01)))
model2.add(Dropout(0.5))

model2.add(Dense(64, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.01)))
model2.add(Dropout(0.5))

model2.add(Dense(output_dim, activation='softmax'))

print(model2.summary())

model2.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model2.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

```

Layer (type)	Output Shape	Param #
dense_4 (Dense)	(None, 512)	401920
dense_5 (Dense)	(None, 128)	65664
dropout_3 (Dropout)	(None, 128)	0
dense_6 (Dense)	(None, 64)	8256
dropout_4 (Dropout)	(None, 64)	0
dense_7 (Dense)	(None, 10)	650
Total params: 476,490		
Trainable params: 476,490		
Non-trainable params: 0		

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 4s 61us/step - loss: 13.4442
 - acc: 0.1647 - val_loss: 12.6372 - val_acc: 0.2158

Epoch 2/20

60000/60000 [=====] - 3s 51us/step - loss: 11.6720
 - acc: 0.2752 - val_loss: 9.9816 - val_acc: 0.3807

Epoch 3/20

60000/60000 [=====] - 3s 51us/step - loss: 10.5429
 - acc: 0.3455 - val_loss: 9.9067 - val_acc: 0.3851

Epoch 4/20

60000/60000 [=====] - 3s 53us/step - loss: 10.5595
 - acc: 0.3446 - val_loss: 9.9687 - val_acc: 0.3814

Epoch 5/20

60000/60000 [=====] - 3s 53us/step - loss: 10.3295
 - acc: 0.3588 - val_loss: 9.9273 - val_acc: 0.3839

Epoch 6/20

60000/60000 [=====] - 3s 50us/step - loss: 10.0814
 - acc: 0.3743 - val_loss: 9.9309 - val_acc: 0.3836

```
Epoch 7/20
60000/60000 [=====] - 3s 51us/step - loss: 10.2240
- acc: 0.3655 - val_loss: 9.8755 - val_acc: 0.3872
Epoch 8/20
60000/60000 [=====] - 3s 50us/step - loss: 9.8280 -
acc: 0.3900 - val_loss: 8.9418 - val_acc: 0.4450
Epoch 9/20
60000/60000 [=====] - 3s 50us/step - loss: 9.4864 -
acc: 0.4113 - val_loss: 9.7467 - val_acc: 0.3952
Epoch 10/20
60000/60000 [=====] - 3s 52us/step - loss: 9.4932 -
acc: 0.4108 - val_loss: 9.0106 - val_acc: 0.4409
Epoch 11/20
60000/60000 [=====] - 3s 54us/step - loss: 8.9193 -
acc: 0.4464 - val_loss: 8.6056 - val_acc: 0.4660
Epoch 12/20
60000/60000 [=====] - 3s 52us/step - loss: 8.9756 -
acc: 0.4430 - val_loss: 8.4394 - val_acc: 0.4764
Epoch 13/20
60000/60000 [=====] - 3s 51us/step - loss: 8.9882 -
acc: 0.4421 - val_loss: 8.3750 - val_acc: 0.4804
Epoch 14/20
60000/60000 [=====] - 3s 51us/step - loss: 8.7018 -
acc: 0.4599 - val_loss: 8.2476 - val_acc: 0.4883
Epoch 15/20
60000/60000 [=====] - 3s 51us/step - loss: 8.6530 -
acc: 0.4630 - val_loss: 8.2132 - val_acc: 0.4902
Epoch 16/20
60000/60000 [=====] - 3s 54us/step - loss: 8.6741 -
acc: 0.4617 - val_loss: 8.0635 - val_acc: 0.4996
Epoch 17/20
60000/60000 [=====] - 3s 52us/step - loss: 9.0568 -
acc: 0.4379 - val_loss: 8.3073 - val_acc: 0.4846
Epoch 18/20
60000/60000 [=====] - 3s 53us/step - loss: 8.8086 -
acc: 0.4533 - val_loss: 8.2073 - val_acc: 0.4908
Epoch 19/20
60000/60000 [=====] - 3s 53us/step - loss: 8.8016 -
acc: 0.4538 - val_loss: 8.3379 - val_acc: 0.4827
Epoch 20/20
60000/60000 [=====] - 3s 53us/step - loss: 8.7108 -
acc: 0.4595 - val_loss: 8.5877 - val_acc: 0.4672
```

train and test loss vs Epochs

In [19]:

```
score = model2.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=0)

# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

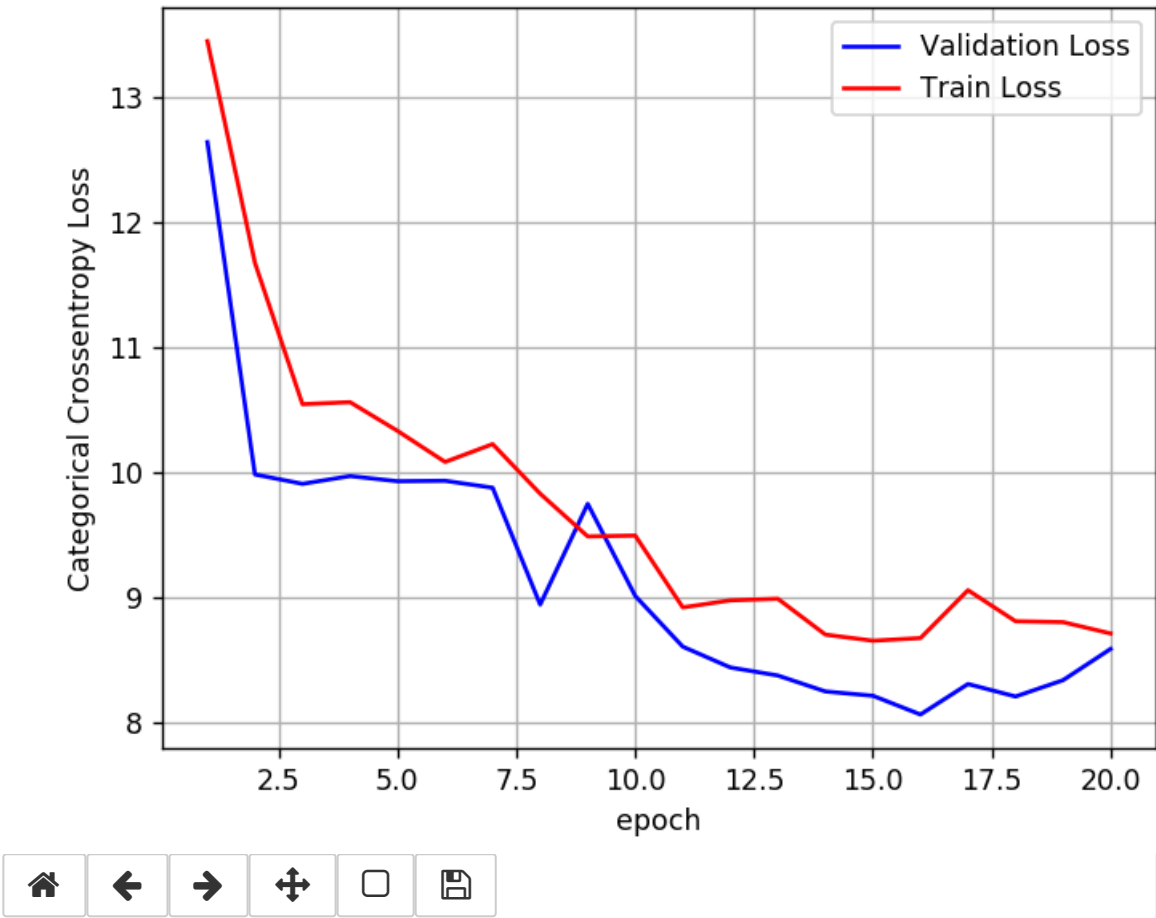
# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 8.587725143432618

Test accuracy: 0.4672

Figure 6



Visualizing weights with violin plot

In [20]:

```

w_after = model2.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

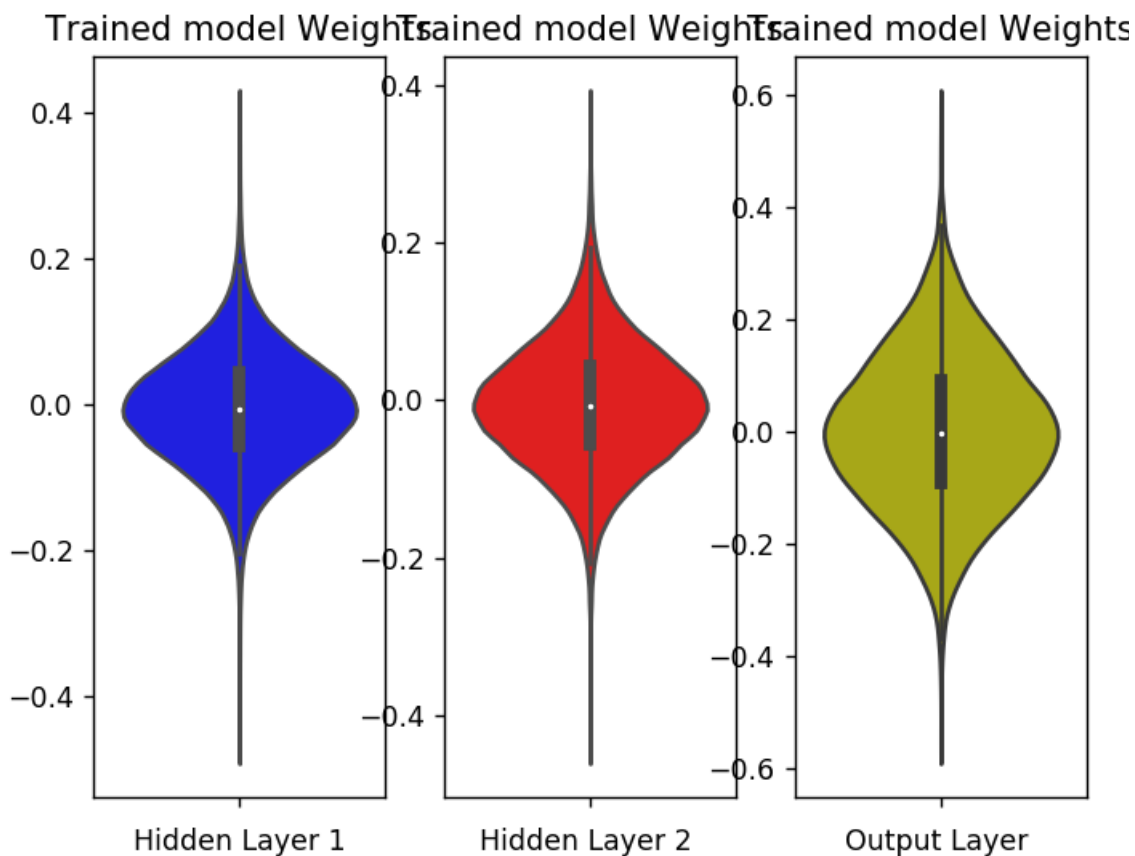
fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```

Figure 7



1. Weights at hidden layer 1 ranges(-0.2 to 0.2).
2. Weights at hidden layer 2 ranges (-0.3 to 0.3).
3. Weights at output layer ranges (-0.4 to 0.4).

Model 3

In [21]:

```

from keras.initializers import he_normal
from keras.layers.normalization import BatchNormalization

model3 = Sequential()

model3.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=he_normal))
model3.add(Dense(128, activation='relu', input_shape=(input_dim,), kernel_initializer=he_normal))
model3.add(Dropout(0.5))

model3.add(Dense(128, activation='relu', input_shape=(input_dim,), kernel_initializer=he_normal))
model3.add(Dense(64, activation='relu', input_shape=(input_dim,), kernel_initializer=he_normal))
model3.add(Dense(32, activation='relu', kernel_initializer=he_normal(seed=None)))

model3.add(Dropout(0.5))

model3.add(BatchNormalization())

model3.add(Dense(output_dim, activation='softmax'))

print(model3.summary())

model3.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model3.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_val, Y_val))

```

Layer (type)	Output Shape	Param #
dense_8 (Dense)	(None, 512)	401920
dense_9 (Dense)	(None, 128)	65664
dropout_5 (Dropout)	(None, 128)	0
dense_10 (Dense)	(None, 128)	16512
dense_11 (Dense)	(None, 64)	8256
dense_12 (Dense)	(None, 32)	2080
dropout_6 (Dropout)	(None, 32)	0
batch_normalization_1 (Batch Normalization)	(None, 32)	128
dense_13 (Dense)	(None, 10)	330
Total params: 494,890		
Trainable params: 494,826		
Non-trainable params: 64		

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 5s 91us/step - loss: 0.9664 -
 acc: 0.7120 - val_loss: 0.2113 - val_acc: 0.9447

```
Epoch 2/20
60000/60000 [=====] - 4s 70us/step - loss: 0.3508 -
acc: 0.9156 - val_loss: 0.1334 - val_acc: 0.9623
Epoch 3/20
60000/60000 [=====] - 4s 70us/step - loss: 0.2472 -
acc: 0.9411 - val_loss: 0.1125 - val_acc: 0.9684
Epoch 4/20
60000/60000 [=====] - 4s 71us/step - loss: 0.1941 -
acc: 0.9531 - val_loss: 0.1159 - val_acc: 0.9704
Epoch 5/20
60000/60000 [=====] - 4s 71us/step - loss: 0.1642 -
acc: 0.9596 - val_loss: 0.0865 - val_acc: 0.9775
Epoch 6/20
60000/60000 [=====] - 4s 71us/step - loss: 0.1388 -
acc: 0.9652 - val_loss: 0.0983 - val_acc: 0.9747
Epoch 7/20
60000/60000 [=====] - 4s 71us/step - loss: 0.1247 -
acc: 0.9679 - val_loss: 0.0996 - val_acc: 0.9764
Epoch 8/20
60000/60000 [=====] - 4s 71us/step - loss: 0.1184 -
acc: 0.9687 - val_loss: 0.0945 - val_acc: 0.9766
Epoch 9/20
60000/60000 [=====] - 4s 70us/step - loss: 0.1051 -
acc: 0.9713 - val_loss: 0.1006 - val_acc: 0.9777
Epoch 10/20
60000/60000 [=====] - 4s 71us/step - loss: 0.0984 -
acc: 0.9736 - val_loss: 0.0903 - val_acc: 0.9790
Epoch 11/20
60000/60000 [=====] - 4s 71us/step - loss: 0.0850 -
acc: 0.9769 - val_loss: 0.0896 - val_acc: 0.9794
Epoch 12/20
60000/60000 [=====] - 4s 71us/step - loss: 0.0816 -
acc: 0.9781 - val_loss: 0.0981 - val_acc: 0.9798
Epoch 13/20
60000/60000 [=====] - 4s 74us/step - loss: 0.0774 -
acc: 0.9793 - val_loss: 0.0944 - val_acc: 0.9796
Epoch 14/20
60000/60000 [=====] - 4s 71us/step - loss: 0.0719 -
acc: 0.9802 - val_loss: 0.1023 - val_acc: 0.9785
Epoch 15/20
60000/60000 [=====] - 4s 72us/step - loss: 0.0708 -
acc: 0.9800 - val_loss: 0.1031 - val_acc: 0.9792
Epoch 16/20
60000/60000 [=====] - 4s 74us/step - loss: 0.0713 -
acc: 0.9799 - val_loss: 0.1086 - val_acc: 0.9784
Epoch 17/20
60000/60000 [=====] - 4s 75us/step - loss: 0.0647 -
acc: 0.9822 - val_loss: 0.0949 - val_acc: 0.9806
Epoch 18/20
60000/60000 [=====] - 4s 74us/step - loss: 0.0642 -
acc: 0.9815 - val_loss: 0.1004 - val_acc: 0.9797
Epoch 19/20
60000/60000 [=====] - 4s 73us/step - loss: 0.0576 -
acc: 0.9833 - val_loss: 0.1076 - val_acc: 0.9807
Epoch 20/20
60000/60000 [=====] - 4s 71us/step - loss: 0.0567 -
acc: 0.9835 - val_loss: 0.1011 - val_acc: 0.9815
```

Train and test loss vs epochs

In [22]:

```
score = model3.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=0)

# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

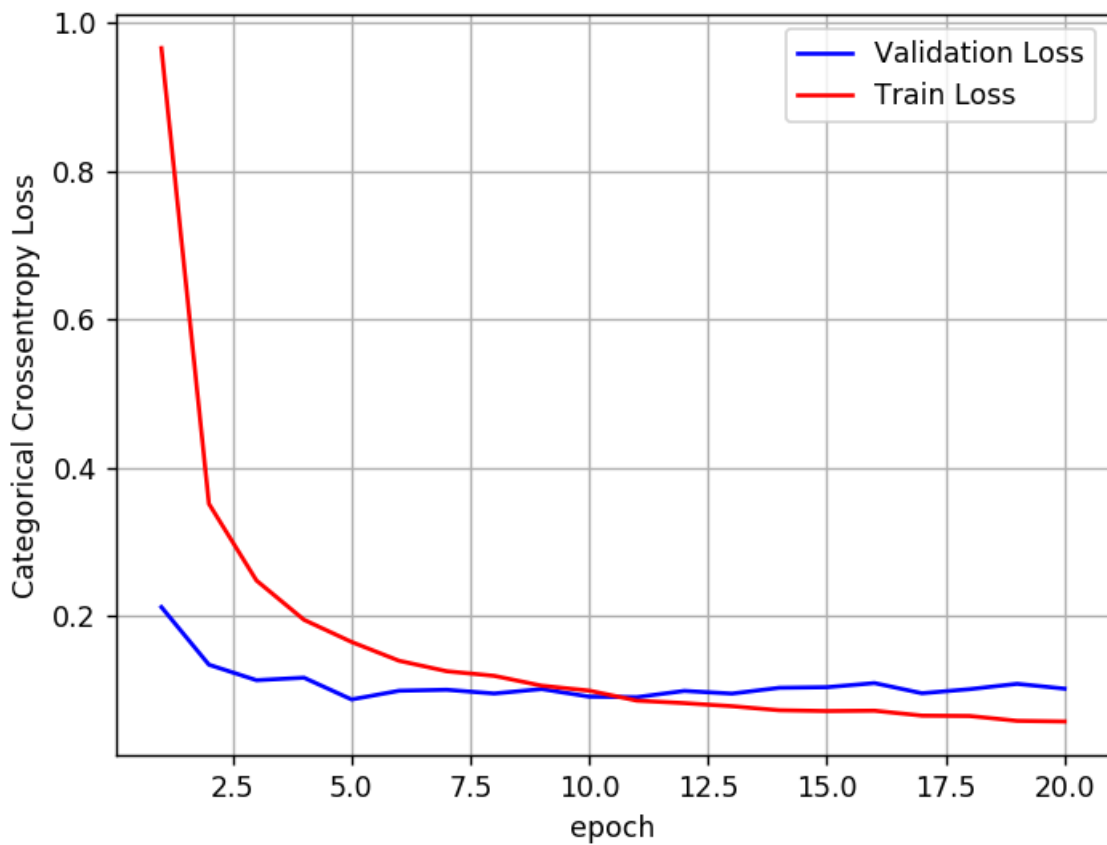
# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.10108382558252847

Test accuracy: 0.9815

Figure 8



1. validation loss was not decreasing after 2 epochs while train loss was decreasing slowly.

Visualizing weights with violin plot

In [23]:

```

w_after = model3.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

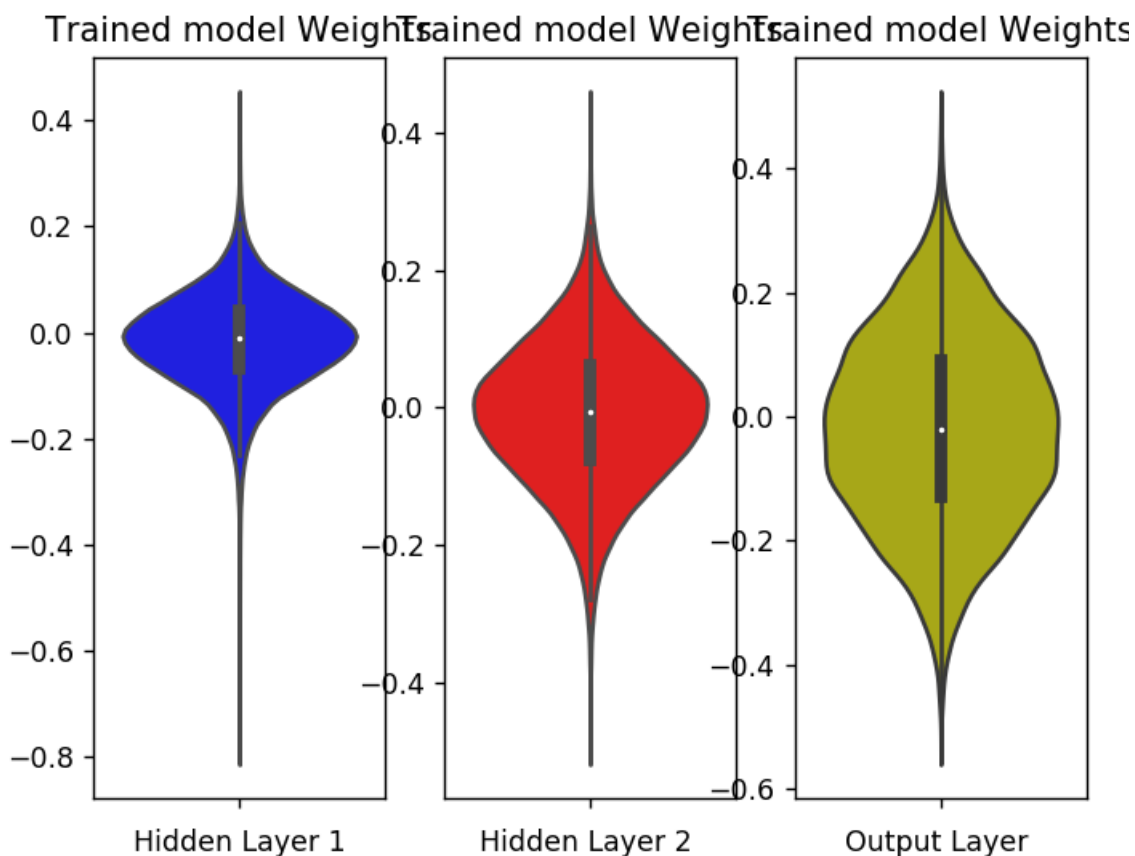
fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```

Figure 9



1. Weights at hidden layer 1 ranges(-0.2 to 0.2).
2. Weights at hidden layer 2 ranges (-0.3 to 0.3).
3. Weights at output layer ranges (-0.6 to 0.4).

In [24]:

```
from prettytable import PrettyTable
x = PrettyTable()
x.field_names = ["model", "No. dense layers", "dropout layers and rate", "weight initilization", "Train accuracy", "Test accuracy"]
x.add_row(["Model 1", "2 (512,128)", "2 (0.5)", "Random", "58.99%", "66.46%"])
x.add_row(["Model 2", "3 (512,128,64)", "2 (0.5)", "Random", "45.9%", "46.7%"])
x.add_row(["Model 3", "5(512,128,128,64,32)", "2 (0.5)", "He_normal", "98%", "98%"])

print(x)
```

```
+-----+-----+-----+-----+
+-----+-----+-----+-----+
| model | No. dense layers | dropout layers and rate | weight initilization | Train accuracy | Test accuracy |
+-----+-----+-----+-----+
+-----+-----+-----+-----+
| Model 1 | 2 (512,128) | 2 (0.5) | Random | 58.99% | 66.46% |
| Model 2 | 3 (512,128,64) | 2 (0.5) | Random | 45.9% | 46.7% |
| Model 3 | 5(512,128,128,64,32) | 2 (0.5) | He_normal | 98% | 98% |
+-----+-----+-----+-----+
+-----+-----+-----+-----+
```

Obeservations:

1. Model 1 has high train loss and low validation loss.
2. Model 1 has performed worst.
3. Slight decrease in model accuracy has been obeserved while adding dropouts.
4. He_normal weight initialization has slightly fasten the convergence of the model.
5. Model 4 is slightly overfitting with train loss of 0.05 and test loss of 0.1.