

Netflix Movie Recommendation System- Case Study

9

1.1 Problem Description

Netflix is all about connecting people to the movies they love. To help customers find those movies, they developed world-class movie recommendation system: CinematchSM. Its job is to predict whether someone will enjoy a movie based on how much they liked or disliked other movies. Netflix use those predictions to make personal movie recommendations based on each customer's unique tastes. And while **Cinematch** is doing pretty well, it can always be made better.

Now there are a lot of interesting alternative approaches to how Cinematch works that netflix haven't tried. Some are described in the literature, some aren't. We're curious whether any of these can beat Cinematch by making better predictions. Because, frankly, if there is a much better approach it could make a big difference to our customers and our business.

Credits: <https://www.netflixprize.com/rules.html>

1.2 Problem Statement

Netflix provided a lot of anonymous rating data, and a prediction accuracy bar that is 10% better than what Cinematch can do on the same training data set. (Accuracy is a measurement of how closely predicted ratings of movies match subsequent actual ratings.)

1.4 Real world/Business Objectives and constraints

Objectives:

1. Predict the rating that a user would give to a movie that he has not yet rated.
2. Minimize the difference between predicted and actual rating (RMSE and MAPE)

Constraints:

1. Some form of interpretability.

2. Machine Learning Problem

2.1 Data

2.2.2 Performance metric

- Mean Absolute Percentage Error: https://en.wikipedia.org/wiki/Mean_absolute_percentage_error
- Root Mean Square Error: https://en.wikipedia.org/wiki/Root-mean-square_deviation

2.2.3 Machine Learning Objective and Constraints

1. Minimize RMSE.
2. Try to provide some interpretability.

In [19]:

```
# this is just to know how much time will it take to run this entire ipython notebook
from datetime import datetime
# globalstart = datetime.now()
import pandas as pd
import numpy as np
import matplotlib
matplotlib.use('nbagg')

import matplotlib.pyplot as plt
plt.rcParams.update({'figure.max_open_warning': 0})

import seaborn as sns
sns.set_style('whitegrid')
import os
from scipy import sparse
from scipy.sparse import csr_matrix

from sklearn.decomposition import TruncatedSVD
from sklearn.metrics.pairwise import cosine_similarity
import random
```

3. Exploratory Data Analysis

3.1 Preprocessing

3.1.1 Converting / Merging whole data to required format: u_i , m_j , r_{ij}

In [20]:

```

start = datetime.now()
if not os.path.isfile('data.csv'):
    # Create a file 'data.csv' before reading it
    # Read all the files in netflix and store them in one big file('data.csv')
    # We re reading from each of the four files and appendig each rating to a global file '
    data = open('data.csv', mode='w')

    row = list()
    files=['combined_data_1.txt','combined_data_2.txt',
           'combined_data_3.txt', 'combined_data_4.txt']
    for file in files:
        print("Reading ratings from {}".format(file))
        with open(file) as f:
            for line in f:
                del row[:] # you don't have to do this.
                line = line.strip()
                if line.endswith(':'):
                    # All below are ratings for this movie, until another movie appears.
                    movie_id = line.replace(':', '')
                else:
                    row = [x for x in line.split(',')]
                    row.insert(0, movie_id)
                    data.write(','.join(row))
                    data.write('\n')
            print("Done.\n")
    data.close()
print('Time taken :', datetime.now() - start)

```

Time taken : 0:00:00

In [21]:

```

print("creating the dataframe from data.csv file..")
df = pd.read_csv('data.csv', sep=',',
                 names=['movie', 'user', 'rating', 'date'])
df.date = pd.to_datetime(df.date)
print('Done.\n')

# we are arranging the ratings according to time.
print('Sorting the dataframe by date..')
df.sort_values(by='date', inplace=True)
print('Done..')

```

creating the dataframe from data.csv file..
Done.

Sorting the dataframe by date..
Done..

In [22]:

```
df.describe()['rating']
```

Out[22]:

```
count    1.004805e+08
mean     3.604290e+00
std      1.085219e+00
min      1.000000e+00
25%      3.000000e+00
50%      4.000000e+00
75%      4.000000e+00
max      5.000000e+00
Name: rating, dtype: float64
```

3.1.2 Checking for NaN values

In [23]:

```
# just to make sure that all Nan containing rows are deleted..
print("No of Nan values in our dataframe : ", sum(df.isnull().any()))
```

No of Nan values in our dataframe : 0

3.1.3 Removing Duplicates

In [24]:

```
dup_bool = df.duplicated(['movie', 'user', 'rating'])
dups = sum(dup_bool) # by considering all columns..( including timestamp)
print("There are {} duplicate rating entries in the data..".format(dups))
```

There are 0 duplicate rating entries in the data..

3.1.4 Basic Statistics (#Ratings, #Users, and #Movies)

In [25]:

```
print("Total data ")
print("-"*50)
print("\nTotal no of ratings :", df.shape[0])
print("Total No of Users   :", len(np.unique(df.user)))
print("Total No of movies  :", len(np.unique(df.movie)))
```

Total data

```
-----
Total no of ratings : 100480507
Total No of Users   : 480189
Total No of movies  : 17770
```

3.2 Splitting data into Train and Test(80:20)

In [26]:

```

if not os.path.isfile('train.csv'):
    # create the dataframe and store it in the disk for offline purposes..
    df.iloc[:int(df.shape[0]*0.80)].to_csv("train.csv", index=False)

if not os.path.isfile('test.csv'):
    # create the dataframe and store it in the disk for offline purposes..
    df.iloc[int(df.shape[0]*0.80):].to_csv("test.csv", index=False)

train_df = pd.read_csv("train.csv", parse_dates=['date'])
test_df = pd.read_csv("test.csv")

```

In [27]:

```
train_df.shape
```

Out[27]:

(80384405, 4)

3.2.1 Basic Statistics in Train data (#Ratings, #Users, and #Movies)

In [28]:

```

# movies = train_df.movie.value_counts()
# users = train_df.user.value_counts()
print("Training data ")
print("-"*50)
print("\nTotal no of ratings :", train_df.shape[0])
print("Total No of Users    :", len(np.unique(train_df.user)))
print("Total No of movies   :", len(np.unique(train_df.movie)))

```

Training data

Total no of ratings : 80384405
 Total No of Users : 405041
 Total No of movies : 17424

3.2.2 Basic Statistics in Test data (#Ratings, #Users, and #Movies)

In [29]:

```

print("Test data ")
print("-"*50)
print("\nTotal no of ratings :", test_df.shape[0])
print("Total No of Users    :", len(np.unique(test_df.user)))
print("Total No of movies   :", len(np.unique(test_df.movie)))

```

Test data

Total no of ratings : 20096102
 Total No of Users : 349312
 Total No of movies : 17757

3.3 Exploratory Data Analysis on Train data

In [30]:

```
# method to make y-axis more readable
def human(num, units = 'M'):
    units = units.lower()
    num = float(num)
    if units == 'k':
        return str(num/10**3) + " K"
    elif units == 'm':
        return str(num/10**6) + " M"
    elif units == 'b':
        return str(num/10**9) + " B"
```

3.3.1 Distribution of ratings

In [31]:

```
fig, ax = plt.subplots()
plt.title('Distribution of ratings over Training dataset', fontsize=15)
sns.countplot(train_df.rating)
ax.set_yticklabels([human(item, 'M') for item in ax.get_yticks()])
ax.set_ylabel('No. of Ratings(Millions)')

plt.show()
```

Figure 1



Add new column (week day) to the data set for analysis.

In [32]:

```
# It is used to skip the warning ''SettingWithCopyWarning' '..  
pd.options.mode.chained_assignment = None # default='warn'  
  
train_df['day_of_week'] = train_df.date.dt.weekday_name  
  
train_df.tail()
```

Out[32]:


	movie	user	rating	date	day_of_week
80384400	12074	2033618	4	2005-08-08	Monday
80384401	862	1797061	3	2005-08-08	Monday
80384402	10986	1498715	5	2005-08-08	Monday
80384403	14861	500016	4	2005-08-08	Monday
80384404	5926	1044015	5	2005-08-08	Monday

3.3.2 Number of Ratings per a month

In [33]:

```
ax = train_df.resample('m', on='date')['rating'].count().plot()
ax.set_title('No of ratings per month (Training data)')
plt.xlabel('Month')
plt.ylabel('No of ratings(per month)')
ax.set_yticklabels([human(item, 'M') for item in ax.get_yticks()])
plt.show()
```

Figure 2



3.3.3 Analysis on the Ratings given by user

In [34]:

```
no_of_rated_movies_per_user = train_df.groupby(by='user')['rating'].count().sort_values(asc
no_of_rated_movies_per_user.head()
```

Out[34]:

```
user
305344    17112
2439493    15896
387418     15402
1639792     9767
1461435     9447
Name: rating, dtype: int64
```

In [35]:


```
fig = plt.figure(figsize=plt.figaspect(.5))

ax1 = plt.subplot(121)
sns.kdeplot(no_of Rated movies per user, shade=True, ax=ax1)
plt.xlabel('No of ratings by user')
plt.title("PDF")

ax2 = plt.subplot(122)
sns.kdeplot(no_of Rated movies per user, shade=True, cumulative=True, ax=ax2)
plt.xlabel('No of ratings by user')
plt.title('CDF')

plt.show()
```

Figure 3



In [36]:

```
no_of Rated_movies_per_user.describe()
```

Out[36]:

```
count      405041.000000
mean        198.459921
std         290.793238
min          1.000000
25%         34.000000
50%         89.000000
75%        245.000000
max       17112.000000
Name: rating, dtype: float64
```

There, is something interesting going on with the quantiles..

In [37]:

```
quantiles = no_of Rated_movies_per_user.quantile(np.arange(0,1.01,0.01), interpolation='hig
```

In [38]:

```


plt.title("Quantiles and their Values")
quantiles.plot()
# quantiles with 0.05 difference
plt.scatter(x=quantiles.index[::5], y=quantiles.values[::5], c='orange', label="quantiles w
# quantiles with 0.25 difference
plt.scatter(x=quantiles.index[::25], y=quantiles.values[::25], c='m', label = "quantiles wi
plt.ylabel('No of ratings by user')
plt.xlabel('Value at the quantile')
plt.legend(loc='best')

# annotate the 25th, 50th, 75th and 100th percentile values....
for x,y in zip(quantiles.index[::25], quantiles[::25]):
    plt.annotate(s="({} , {})".format(x,y), xy=(x,y), xytext=(x-0.05, y+500)
                ,fontweight='bold')

plt.show()

```

Figure 4



In [39]:

```
quantiles[::5]
```

Out[39]:

0.00	1
0.05	7
0.10	15
0.15	21
0.20	27
0.25	34
0.30	41
0.35	50
0.40	60
0.45	73
0.50	89
0.55	109
0.60	133
0.65	163
0.70	199
0.75	245
0.80	307
0.85	392
0.90	520
0.95	749
1.00	17112

Name: rating, dtype: int64

how many ratings at the last 5% of all ratings??

In [40]:

```
print('\n No of ratings at last 5 percentile : {}'.format(sum(no_of Rated movies per user
```

No of ratings at last 5 percentile : 20305

3.3.4 Analysis of ratings of a movie given by a user


In [41]:

```
no_of_ratings_per_movie = train_df.groupby(by='movie')['rating'].count().sort_values(ascending=True)

fig = plt.figure(figsize=plt.figaspect(.5))
ax = plt.gca()
plt.plot(no_of_ratings_per_movie.values)
plt.title('# RATINGS per Movie')
plt.xlabel('Movie')
plt.ylabel('No of Users who rated a movie')
ax.set_xticklabels([])

plt.show()
```

Figure 5



- It is very skewed.. just like number of ratings given per user.


- There are some movies (which are very popular) which are rated by huge number of users.
- But most of the movies(like 90%) got some hundreds of ratings.

3.3.5 Number of ratings on each day of the week

In [42]:

```
fig, ax = plt.subplots()
sns.countplot(x='day_of_week', data=train_df, ax=ax)
plt.title('No of ratings on each day...')
plt.ylabel('Total no of ratings')
plt.xlabel('')
ax.set_yticklabels([human(item, 'M') for item in ax.get_yticks()])
plt.show()
```

Figure 6




In [34]:

```

start = datetime.now()
fig = plt.figure(figsize=plt.figaspect(.45))
sns.boxplot(y='rating', x='day_of_week', data=train_df)
plt.show()
print(datetime.now() - start)

```



0:00:31.713358

In [43]:

```

avg_week_df = train_df.groupby(by=['day_of_week'])['rating'].mean()
print(" AVerage ratings")
print("-"*30)
print(avg_week_df)
print("\n")

```

```

Average ratings
-----
day_of_week
Friday      3.585274
Monday      3.577250
Saturday    3.591791
Sunday      3.594144
Thursday    3.582463
Tuesday     3.574438
Wednesday   3.583751
Name: rating, dtype: float64

```

3.3.6 Creating sparse matrix from data frame

MOVIE_ID	USER_ID	RATING
1	1	3
2	1	4
3	1	2
3	2	1
4	2	4
8	2	2
1	3	3
7	3	1
10	3	5



	1	2	3	4	5	6	7	8	9	10	(movie)
1	3	4	2	-	-	-	-	-	-	-	
2	-	-	1	4	-	-	-	-	-	-	
3	3	-	-	-	-	-	1	-	-	5	
(user)											

3.3.6.1 Creating sparse matrix from train data frame

In [44]:

```
start = datetime.now()
if os.path.isfile('train_sparse_matrix.npz'):
    print("It is present in your pwd, getting it from disk....")
    # just get it from the disk instead of computing it
    train_sparse_matrix = sparse.load_npz('train_sparse_matrix.npz')
    print("DONE..")
else:
    print("We are creating sparse_matrix from the dataframe..")
    # create sparse_matrix and store it for after usage.
    # csr_matrix(data_values, (row_index, col_index), shape_of_matrix)
    # It should be in such a way that, MATRIX[row, col] = data
    train_sparse_matrix = sparse.csr_matrix((train_df.rating.values, (train_df.user.values,
                                                                    train_df.movie.values)),)

    print('Done. It\'s shape is : (user, movie) : ', train_sparse_matrix.shape)
    print('Saving it into disk for furthur usage..')
    # save it into disk
    sparse.save_npz("train_sparse_matrix.npz", train_sparse_matrix)
    print('Done..\n')

print(datetime.now() - start)
```

It is present in your pwd, getting it from disk....
 DONE..
 0:00:03.263051

The Sparsity of Train Sparse Matrix

In [45]:

```
us, mv = train_sparse_matrix.shape
elem = train_sparse_matrix.count_nonzero()

print("Sparsity Of Train matrix : {} %".format( (1-(elem/(us*mv))) * 100) )
```

Sparsity Of Train matrix : 99.8292709259195 %

3.3.6.2 Creating sparse matrix from test data frame

In [46]:

```

start = datetime.now()
if os.path.isfile('test_sparse_matrix.npz'):
    print("It is present in your pwd, getting it from disk....")
    # just get it from the disk instead of computing it
    test_sparse_matrix = sparse.load_npz('test_sparse_matrix.npz')
    print("DONE..")
else:
    print("We are creating sparse_matrix from the dataframe..")
    # create sparse_matrix and store it for after usage.
    # csr_matrix(data_values, (row_index, col_index), shape_of_matrix)
    # It should be in such a way that, MATRIX[row, col] = data
    test_sparse_matrix = sparse.csr_matrix((test_df.rating.values, (test_df.user.values,
                                                                    test_df.movie.values)))

    print('Done. It\'s shape is : (user, movie) : ', test_sparse_matrix.shape)
    print('Saving it into disk for furthur usage..')
    # save it into disk
    sparse.save_npz("test_sparse_matrix.npz", test_sparse_matrix)
    print('Done..\n')

print(datetime.now() - start)

```

It is present in your pwd, getting it from disk....
 DONE..
 0:00:00.892044

The Sparsity of Test data Matrix

In [47]:

```

us,mv = test_sparse_matrix.shape
elem = test_sparse_matrix.count_nonzero()

print("Sparsity Of Test matrix : {} % ".format( (1-(elem/(us*mv))) * 100) )

```

Sparsity Of Test matrix : 99.95731772988694 %

3.3.7 Finding Global average of all movie ratings, Average rating per user, and Average rating per movie

In [48]:

```
# get the user averages in dictionary (key: user_id/movie_id, value: avg rating)

def get_average_ratings(sparse_matrix, of_users):

    # average ratings of user/axes
    ax = 1 if of_users else 0 # 1 - User axes, 0 - Movie axes

    # ".A1" is for converting Column_Matrix to 1-D numpy array
    sum_of_ratings = sparse_matrix.sum(axis=ax).A1
    # Boolean matrix of ratings ( whether a user rated that movie or not)
    isRated = sparse_matrix!=0
    # no of ratings that each user OR movie..
    no_of_ratings = isRated.sum(axis=ax).A1

    # max_user and max_movie ids in sparse matrix
    u,m = sparse_matrix.shape
    # create a dictionary of users and their average ratings..
    average_ratings = { i : sum_of_ratings[i]/no_of_ratings[i]
                        for i in range(u if of_users else m)
                        if no_of_ratings[i] !=0}

    # return that dictionary of average ratings
    return average_ratings
```

3.3.7.1 finding global average of all movie ratings

In [49]:

```
train_averages = dict()
# get the global average of ratings in our train set.
train_global_average = train_sparse_matrix.sum()/train_sparse_matrix.count_nonzero()
train_averages['global'] = train_global_average
train_averages
```

Out[49]:

```
{'global': 3.582890686321557}
```

3.3.7.2 finding average rating per user

In [50]:

```
train_averages['user'] = get_average_ratings(train_sparse_matrix, of_users=True)
print('\nAverage rating of user 10 : ',train_averages['user'][10])
```

```
Average rating of user 10 : 3.3781094527363185
```

3.3.7.3 finding average rating per movie

In [51]:

```
train_averages['movie'] = get_average_ratings(train_sparse_matrix, of_users=False)
print('\n AVerage rating of movie 15 :',train_averages['movie'][15])
```

Average rating of movie 15 : 3.3038461538461537

3.3.7.4 PDF's & CDF's of Avg.Ratings of Users & Movies (In Train Data)

In [52]:

```

start = datetime.now()
# draw pdfs for average rating per user and average
fig, (ax1, ax2) = plt.subplots(nrows=1, ncols=2, figsize=plt.figaspect(.5))
fig.suptitle('Avg Ratings per User and per Movie', fontsize=15)


ax1.set_title('Users-Avg-Ratings')
# get the list of average user ratings from the averages dictionary..
user_averages = [rat for rat in train_averages['user'].values()]
sns.distplot(user_averages, ax=ax1, hist=False,
              kde_kws=dict(cumulative=True), label='Cdf')
sns.distplot(user_averages, ax=ax1, hist=False, label='Pdf')


ax2.set_title('Movies-Avg-Rating')
# get the list of movie_average_ratings from the dictionary..
movie_averages = [rat for rat in train_averages['movie'].values()]
sns.distplot(movie_averages, ax=ax2, hist=False,
              kde_kws=dict(cumulative=True), label='Cdf')
sns.distplot(movie_averages, ax=ax2, hist=False, label='Pdf')

plt.show()
print(datetime.now() - start)

```

Figure 7





0:00:44.899368

3.3.8 Cold Start problem

3.3.8.1 Cold Start problem with Users

In [53]:

```
total_users = len(np.unique(df.user))
users_train = len(train_averages['user'])
new_users = total_users - users_train

print('\nTotal number of Users  :', total_users)
print('\nNumber of Users in Train data :', users_train)
print("\nNo of Users that didn't appear in train data: {}({} %) \n ".format(new_users,
                                                                              np.round((new_users
```

Total number of Users : 480189

Number of Users in Train data : 405041

No of Users that didn't appear in train data: 75148(15.65 %)

We might have to handle **new users (75148)** who didn't appear in train data.

3.3.8.2 Cold Start problem with Movies

In [54]:

```
total_movies = len(np.unique(df.movie))
movies_train = len(train_averages['movie'])
new_movies = total_movies - movies_train

print('\nTotal number of Movies  :', total_movies)
print('\nNumber of Users in Train data :', movies_train)
print("\nNo of Movies that didn't appear in train data: {}({} %) \n ".format(new_movies,
                                                                              np.round((new_movies / total_movies) * 100, 2)))
```

Total number of Movies : 17770

Number of Users in Train data : 17424

No of Movies that didn't appear in train data: 346(1.95 %)

We might have to handle **346 movies** (small comparatively) in test data

3.4 Computing Similarity matrices

3.4.1 Computing User-User Similarity matrix

3.4.1.1 Trying with all dimensions (17k dimensions per user)

In [55]:

```

from sklearn.metrics.pairwise import cosine_similarity

def compute_user_similarity(sparse_matrix, compute_for_few=False, top = 100, verbose=False,
                           draw_time_taken=True):
    no_of_users, _ = sparse_matrix.shape
    # get the indices of non zero rows(users) from our sparse matrix
    row_ind, col_ind = sparse_matrix.nonzero()
    row_ind = sorted(set(row_ind)) # we don't have to
    time_taken = list() # time taken for finding similar users for an user..

    # we create rows, cols, and data lists.., which can be used to create sparse matrices
    rows, cols, data = list(), list(), list()
    if verbose: print("Computing top",top,"similarities for each user..")

    start = datetime.now()
    temp = 0

    for row in row_ind[:top] if compute_for_few else row_ind:
        temp = temp+1
        prev = datetime.now()

        # get the similarity row for this user with all other users
        sim = cosine_similarity(sparse_matrix.getrow(row), sparse_matrix).ravel()
        # We will get only the top 'top' most similar users and ignore rest of them..
        top_sim_ind = sim.argsort()[-top:]
        top_sim_val = sim[top_sim_ind]

        # add them to our rows, cols and data
        rows.extend([row]*top)
        cols.extend(top_sim_ind)
        data.extend(top_sim_val)
        time_taken.append(datetime.now().timestamp() - prev.timestamp())
        if verbose:
            if temp%verb_for_n_rows == 0:
                print("computing done for {} users [ time elapsed : {} ]".format(temp, datetime.now()-start))

    # Lets create sparse matrix out of these and return it
    if verbose: print('Creating Sparse matrix from the computed similarities')
    #return rows, cols, data

    if draw_time_taken:
        plt.plot(time_taken, label = 'time taken for each user')
        plt.plot(np.cumsum(time_taken), label='Total time')
        plt.legend(loc='best')
        plt.xlabel('User')
        plt.ylabel('Time (seconds)')
        plt.show()

    return sparse.csr_matrix((data, (rows, cols)), shape=(no_of_users, no_of_users)), time_

```

In [56]:

```

start = datetime.now()
u_u_sim_sparse, _ = compute_user_similarity(train_sparse_matrix, compute_for_few=True, top
                                             verbose=True)
print("-"*100)
print("Time taken :",datetime.now()-start)

```


Computing top 100 similarities for each user..

```

computing done for 20 users [ time elapsed : 0:01:19.331064 ]
computing done for 40 users [ time elapsed : 0:02:35.099448 ]
computing done for 60 users [ time elapsed : 0:03:48.533472 ]
computing done for 80 users [ time elapsed : 0:05:04.426940 ]
computing done for 100 users [ time elapsed : 0:06:19.325751 ]
Creating Sparse matrix from the computed similarities

```

Figure 8



Time taken : 0:06:29.782266

3.4.1.2 Trying with reduced dimensions (Using TruncatedSVD for dimensionality reduction of user vector)

- We have **405,041 users** in our training set and computing similarities between them..(**17K dimensional**

vector..) is time consuming..

- From above plot, It took roughly **8.88 sec** for computing simlilar users for **one user**
- We have **405,041 users** with us in training set.
- $405041 \times 8.88 = 3596764.08 \text{ sec} = 59946.068 \text{ min} = 999.101133333 \text{ hours} = 41.629213889 \text{ days}$.
 - Even if we run on 4 cores parallely (a typical system now a days), It will still take almost **10 and 1/2** days.

IDEA: Instead, we will try to reduce the dimentsions using SVD, so that **it might** speed up the process...

In [27]:

```
from datetime import datetime
from sklearn.decomposition import TruncatedSVD

start = datetime.now()

# initilaize the algorithm with some parameters..
# All of them are default except n_components. n_itr is for Randomized SVD solver.
netflix_svd = TruncatedSVD(n_components=500, algorithm='randomized', random_state=15)
trunc_svd = netflix_svd.fit_transform(train_sparse_matrix)

print(datetime.now()-start)
```

Out[27]:

```
"from datetime import datetime\nfrom sklearn.decomposition import TruncatedSVD\n\nstart = datetime.now()\n\n# initilaize the algorithm with some parameters..\n# All of them are default except n_components. n_itr is for Randomized SVD solver.\nnetflix_svd = TruncatedSVD(n_components=500, algorithm='randomized', random_state=15)\ntrunc_svd = netflix_svd.fit_transform(train_sparse_matrix)\n\nprint(datetime.now()-start)"
```

Here,

- $\Sigma \leftarrow (\text{netflix_svd.singular_values_})$
- $V^T \leftarrow (\text{netflix_svd.components_})$
- U is not returned. instead **Projection_of_X** onto the new vectorspace is returned.
- It uses **randomized svd** internally, which returns **All 3 of them saperately**. Use that instead..

In [57]:

```
expl_var = np.cumsum(netflix_svd.explained_variance_ratio_)
```

Out[57]:

```
'expl_var = np.cumsum(netflix_svd.explained_variance_ratio_)'
```

In [58]:

```
fig, (ax1, ax2) = plt.subplots(nrows=1, ncols=2, figsize=plt.figaspect(.5))

ax1.set_ylabel("Variance Explained", fontsize=15)
ax1.set_xlabel("# Latent Facors", fontsize=15)
ax1.plot(expl_var)
# annotate some (latentfactors, expl_var) to make it clear
ind = [1, 2,4,8,20, 60, 100, 200, 300, 400, 500]
ax1.scatter(x = [i-1 for i in ind], y = expl_var[[i-1 for i in ind]], c='#ff3300')
for i in ind:
    ax1.annotate(s="{}, {}".format(i, np.round(expl_var[i-1], 2)), xy=(i-1, expl_var[i-1]),
                xytext = ( i+20, expl_var[i-1] - 0.01), fontweight='bold')

change_in_expl_var = [expl_var[i+1] - expl_var[i] for i in range(len(expl_var)-1)]
ax2.plot(change_in_expl_var)

ax2.set_ylabel("Gain in Var_Expl with One Additional LF", fontsize=10)
ax2.yaxis.set_label_position("right")
ax2.set_xlabel("# Latent Facors", fontsize=20)

plt.show()
```

Out[58]:

```
'fig, (ax1, ax2) = plt.subplots(nrows=1, ncols=2, figsize=plt.figaspect(.5))
\n\nax1.set_ylabel("Variance Explained", fontsize=15)\nax1.set_xlabel("# Lat
ent Facors", fontsize=15)\nax1.plot(expl_var)\n# annotate some (latentfactors,
expl_var) to make it clear\nind = [1, 2,4,8,20, 60, 100, 200, 300, 400, 500]
\nax1.scatter(x = [i-1 for i in ind], y = expl_var[[i-1 for i in ind]], c=
'\n\nfor i in ind:\n    ax1.annotate(s="{}, {}".format(i, np.ro
und(expl_var[i-1], 2)), xy=(i-1, expl_var[i-1]),\n                xytext = (
i+20, expl_var[i-1] - 0.01), fontweight='\n\nchange_in_expl_var = [e
xpl_var[i+1] - expl_var[i] for i in range(len(expl_var)-1)]\nax2.plot(change
_in_expl_var)\n\n\nax2.set_ylabel("Gain in Var_Expl with One Additional L
F", fontsize=10)\nax2.yaxis.set_label_position("right")\nax2.set_xlabel("# L
atent Facors", fontsize=20)\n\nplt.show()'
```

In [59]:

```
for i in ind:
    print("{}, {}".format(i, np.round(expl_var[i-1], 2)))
```

Out[59]:

```
'for i in ind:\n    print("{}, {}".format(i, np.round(expl_var[i-1], 2)))'
```

I think 500 dimensions is good enough

- By just taking **(20 to 30)** latent factors, explained variance that we could get is **20 %**.
- To take it to **60%**, we have to take **almost 400 latent factors**. It is not fare.
- It basically is the **gain of variance explained**, if we **add one additional latent factor to it**.

- By adding one by one latent factor too it, the **_gain in expained variance** with that addition is decreasing. (Obviously, because they are sorted that way).
- **LHS Graph:**
 - **x** --- (No of latent factos),
 - **y** --- (The variance explained by taking x latent factors)
- **__More decrease in the line (RHS graph) __:**
 - We are getting more expained variance than before.
- **Less decrease in that line (RHS graph) :**
 - We are not getting benifitted from adding latent factor furthur. This is what is shown in the plots.
- **RHS Graph:**
 - **x** --- (No of latent factors),
 - **y** --- (Gain n Expl_Var by taking one additional latent factor)

In []:

```
# Let's project our Original U_M matrix into into 500 Dimensional space...
start = datetime.now()
trunc_matrix = train_sparse_matrix.dot(netflix_svd.components_.T)
print(datetime.now() - start)
```

In []:

```
type(trunc_matrix), trunc_matrix.shape
```

- Let's convert this to actual sparse matrix and store it for future purposes

In [66]:

```
if not os.path.isfile('trunc_sparse_matrix.npz'):
    # create that sparse sparse matrix
    trunc_sparse_matrix = sparse.csr_matrix(trunc_matrix)
    # Save this truncated sparse matrix for later usage..
    sparse.save_npz('trunc_sparse_matrix', trunc_sparse_matrix)
else:
    trunc_sparse_matrix = sparse.load_npz('trunc_sparse_matrix.npz')
```

In [67]:

```
trunc_sparse_matrix.shape
```

Out[67]:

```
(2649430, 500)
```

In [69]:

```
start = datetime.now()
trunc_u_u_sim_matrix, _ = compute_user_similarity(trunc_sparse_matrix, compute_for_few=True,
                                                  verb_for_n_rows=10)
print("-"*50)
print("time:", datetime.now()-start)
```

Out[69]:

```
'start = datetime.now()\ntrunc_u_u_sim_matrix, _ = compute_user_similarity(t
runc_sparse_matrix, compute_for_few=True, top=50, verbose=True, \n
verb_for_n_rows=10)\nprint("-"*50)\nprint("time:", datetime.now()-start)'
```

: This is taking more time for each user than Original one.

- from above plot, It took almost **12.18** for computing similar users for **one user**
- We have **405041 users** with us in training set.
- $405041 \times 12.18 \text{ sec} = 4933399.38 \text{ sec} = 82223.323 \text{ min} = 1370.388716667 \text{ hours}$
 - Even we run on 4 cores parallely (a typical system now a days), It will still take almost 14 - 15 days.

• Why did this happen...??

- Just think about it. It's not that difficult.

-----_(sparse & dense.....get it ??)_-----

Is there any other way to compute user user similarity..??

-An alternative is to compute similar users for a particular user, whenever required (ie., **Run time**)

- We maintain a binary Vector for users, which tells us whether we already computed or not..
- ***If not*** :
 - Compute top (let's just say, 1000) most similar users for this given user, and add this to our datastructure, so that we can just access it(similar users) without recomputing it again.
-
- ***If It is already Computed***:
 - Just get it directly from our datastructure, which has that information.
 - In production time, We might have to recompute similarities, if it is computed a long time ago. Because user preferences changes over time. If we could maintain in some kind of Timer, which when expires, we have to update it (recompute it).
-
- ***Which datastructure to use***
 - It is purely implementation dependant.
 - One simple method is to maintain a **Dictionary Of Dictionaries**.
 -
 - **key** : **_userid_**
 - **__value__**: **_Again a dictionary_**
 - **__key__** : **_Similar User_**
 - **__value__**: **_Similarity Value_**

3.4.2 Computing Movie-Movie Similarity matrix

In [72]:

```

start = datetime.now()
if not os.path.isfile('m_m_sim_sparse.npz'):
    print("It seems you don't have that file. Computing movie_movie similarity...")
    start = datetime.now()
    m_m_sim_sparse = cosine_similarity(X=train_sparse_matrix.T, dense_output=False)
    print("Done..")
    # store this sparse matrix in disk before using it. For future purposes.
    print("Saving it to disk without the need of re-computing it again.. ")
    sparse.save_npz("m_m_sim_sparse.npz", m_m_sim_sparse)
    print("Done..")
else:
    print("It is there, We will get it.")
    m_m_sim_sparse = sparse.load_npz("m_m_sim_sparse.npz")
    print("Done ...")

print("It's a ",m_m_sim_sparse.shape," dimensional matrix")

print(datetime.now() - start)

```

Out[72]:

```

'start = datetime.now()\nif not os.path.isfile(\'m_m_sim_sparse.npz\'):\n
print("It seems you don\'t have that file. Computing movie_movie similarit
y...")\n    start = datetime.now()\n    m_m_sim_sparse = cosine_similarity(X
=train_sparse_matrix.T, dense_output=False)\n    print("Done..")\n    # stor
e this sparse matrix in disk before using it. For future purposes.\n    prin
t("Saving it to disk without the need of re-computing it again.. ")\n    spa
rse.save_npz("m_m_sim_sparse.npz", m_m_sim_sparse)\n    print("Done..")\nels
e:\n    print("It is there, We will get it.")\n    m_m_sim_sparse = sparse.l
oad_npz("m_m_sim_sparse.npz")\n    print("Done ...")\n\nprint("It\'s a ",m_m
_sim_sparse.shape," dimensional matrix")\n\nprint(datetime.now() - start)'

```

In []:

```
m_m_sim_sparse.shape
```

- Even though we have similarity measure of each movie, with all other movies, We generally don't care much about least similar movies.
- Most of the times, only top_xxx similar items matters. It may be 10 or 100.
- We take only those top similar movie ratings and store them in a saperate dictionary.

In []:

```
movie_ids = np.unique(m_m_sim_sparse.nonzero()[1])
```


In []:

```

start = datetime.now()
similar_movies = dict()
for movie in movie_ids:
    # get the top similar movies and store them in the dictionary
    sim_movies = m_m_sim_sparse[movie].toarray().ravel().argsort()[::-1][1:]
    similar_movies[movie] = sim_movies[:100]
print(datetime.now() - start)

# just testing similar movies for movie_15
similar_movies[15]

```

In [74]:

```
m_m_sim_sparse = np.load("m_m_sim_sparse.npz")
```

3.4.3 Finding most similar movies using similarity matrix

_ Does Similarity really works as the way we expected...? _

_Let's pick some random movie and check for its similar movies....

In [75]:

```

# First Let's load the movie details into soe dataframe..
# movie details are in 'netflix/movie_titles.csv'

movie_titles = pd.read_csv("movie_titles.csv", sep=',', header = None,
                           names=['movie_id', 'year_of_release', 'title'], verbose=True,
                           index_col = 'movie_id', encoding = "ISO-8859-1")

movie_titles.head()

```

Tokenization took: 7.98 ms

Type conversion took: 27.92 ms

Parser memory cleanup took: 0.00 ms

Out[75]:

	year_of_release	title
movie_id		
1	2003.0	Dinosaur Planet
2	2004.0	Isle of Man TT 2004 Review
3	1997.0	Character
4	1994.0	Paula Abdul's Get Up & Dance
5	2004.0	The Rise and Fall of ECW

Similar Movies for 'Vampire Journals'

In [77]:

```
mv_id = 67

print("\nMovie ----->",movie_titles.loc[mv_id].values[1])

print("\nIt has {} Ratings from users.".format(train_sparse_matrix[:,mv_id].getnnz()))

print("\nWe have {} movies which are similar to this and we will get only top most..".format
```

Out[77]:

```
'mv_id = 67\n\nprint("\nMovie ----->",movie_titles.loc[mv_id].values[1])\n\nprint("\nIt has {} Ratings from users.".format(train_sparse_matrix[:,mv_id].getnnz()))\n\nprint("\nWe have {} movies which are similar to this and we will get only top most..".format(m_m_sim_sparse[:,mv_id].getnnz()))'
```

In [78]:

```
similarities = m_m_sim_sparse[mv_id].toarray().ravel()

similar_indices = similarities.argsort()[::-1][1:]

similarities[similar_indices]

sim_indices = similarities.argsort()[::-1][1:] # It will sort and reverse the array and ignore its similarity (ie.,1)\n\n# and return its indices(movie_ids)
```

Out[78]:

```
'similarities = m_m_sim_sparse[mv_id].toarray().ravel()\n\nsimilar_indices = similarities.argsort()[::-1][1:]\n\nsimilarities[similar_indices]\n\nsim_indices = similarities.argsort()[::-1][1:] # It will sort and reverse the array and ignore its similarity (ie.,1)\n\n# and return its indices(movie_ids)'
```

In []:

```
plt.plot(similarities[sim_indices], label='All the ratings')
plt.plot(similarities[sim_indices[:100]], label='top 100 similar movies')
plt.title("Similar Movies of {}(movie_id)".format(mv_id), fontsize=20)
plt.xlabel("Movies (Not Movie_Ids)", fontsize=15)
plt.ylabel("Cosine Similarity", fontsize=15)
plt.legend()
plt.show()
```

Top 10 similar movies

In []:

```
movie_titles.loc[sim_indices[:10]]
```

Similarly, we can **find similar users** and compare how similar they are.

4. Machine Learning Models

In [79]:

```

def get_sample_sparse_matrix(sparse_matrix, no_users, no_movies, path, verbose = True):
    """
        It will get it from the 'path' if it is present or It will create
        and store the sampled sparse matrix in the path specified.
    """

    # get (row, col) and (rating) tuple from sparse_matrix...
    row_ind, col_ind, ratings = sparse.find(sparse_matrix)
    users = np.unique(row_ind)
    movies = np.unique(col_ind)

    print("Original Matrix : (users, movies) -- ({ } { })".format(len(users), len(movies)))
    print("Original Matrix : Ratings -- { }\n".format(len(ratings)))

    # It just to make sure to get same sample everytime we run this program..
    # and pick without replacement....
    np.random.seed(15)
    sample_users = np.random.choice(users, no_users, replace=False)
    sample_movies = np.random.choice(movies, no_movies, replace=False)
    # get the boolean mask or these sampled_items in originl row/col_inds..
    mask = np.logical_and( np.isin(row_ind, sample_users),
                           np.isin(col_ind, sample_movies) )

    sample_sparse_matrix = sparse.csr_matrix((ratings[mask], (row_ind[mask], col_ind[mask])
                                              shape=(max(sample_users)+1, max(sample_movies)

    if verbose:
        print("Sampled Matrix : (users, movies) -- ({ } { })".format(len(sample_users), len(s
        print("Sampled Matrix : Ratings --", format(ratings[mask].shape[0]))

    print('Saving it into disk for furthur usage..')
    # save it into disk
    sparse.save_npz(path, sample_sparse_matrix)
    if verbose:
        print('Done..\n')

    return sample_sparse_matrix

```

4.1 Sampling Data

4.1.1 Build sample train data from the train data

In [80]:

```
start = datetime.now()
path = "sample_train_sparse_matrix.npz"
if os.path.isfile(path):
    print("It is present in your pwd, getting it from disk....")
    # just get it from the disk instead of computing it
    sample_train_sparse_matrix = sparse.load_npz(path)
    print("DONE..")
else:
    # get 25k users and 3k movies from available data
    sample_train_sparse_matrix = get_sample_sparse_matrix(train_sparse_matrix, no_users=120,
                                                         path = path)

print(datetime.now() - start)
```

It is present in your pwd, getting it from disk....
DONE..
0:00:00.093744

4.1.2 Build sample test data from the test data

In [81]:

```
start = datetime.now()

path = "sample_test_sparse_matrix.npz"
if os.path.isfile(path):
    print("It is present in your pwd, getting it from disk....")
    # just get it from the disk instead of computing it
    sample_test_sparse_matrix = sparse.load_npz(path)
    print("DONE..")
else:
    # get 5k users and 500 movies from available data
    sample_test_sparse_matrix = get_sample_sparse_matrix(test_sparse_matrix, no_users=5000,
                                                         path = "sample_test_sparse_matrix.npz")

print(datetime.now() - start)
```

It is present in your pwd, getting it from disk....
DONE..
0:00:00.081780

4.2 Finding Global Average of all movie ratings, Average rating per User, and Average rating per Movie (from sampled train)

In [82]:

```
sample_train_averages = dict()
```

4.2.1 Finding Global Average of all movie ratings

In [83]:

```
# get the global average of ratings in our train set.
global_average = sample_train_sparse_matrix.sum()/sample_train_sparse_matrix.count_nonzero()
sample_train_averages['global'] = global_average
sample_train_averages
```

Out[83]:

```
{'global': 3.5744906003423225}
```

4.2.2 Finding Average rating per User

In [84]:

```
sample_train_averages['user'] = get_average_ratings(sample_train_sparse_matrix, of_users=True)
print('\nAverage rating of user 1515220 :', sample_train_averages['user'][1515220])
```

Average rating of user 1515220 : 3.9555555555555557

4.2.3 Finding Average rating per Movie

In [85]:

```
sample_train_averages['movie'] = get_average_ratings(sample_train_sparse_matrix, of_movies=True)
print('\n Average rating of movie 15153 :', sample_train_averages['movie'][15153])
```

Average rating of movie 15153 : 2.603448275862069

4.3 Featurizing data

In [86]:

```
print('\n No of ratings in Our Sampled train matrix is : {}'.format(sample_train_sparse_matrix.count_nonzero()))
print('\n No of ratings in Our Sampled test matrix is : {}'.format(sample_test_sparse_matrix.count_nonzero()))
```

No of ratings in Our Sampled train matrix is : 281606

No of ratings in Our Sampled test matrix is : 7333

4.3.1 Featurizing data for regression problem

4.3.1.1 Featurizing train data

In [87]:

```
# get users, movies and ratings from our samples train sparse matrix  
sample_train_users, sample_train_movies, sample_train_ratings = sparse.find(sample_train_sp
```

In [88]:

```
#####
# It took me almost 10 hours to prepare this train dataset. #
#####
start = datetime.now()
if os.path.isfile('reg_train.csv'):
    print("File already exists you don't have to prepare again..." )
else:
    print('preparing {} tuples for the dataset.. \n'.format(len(sample_train_ratings)))
    with open('reg_train.csv', mode='w') as reg_data_file:
        count = 0
        for (user, movie, rating) in zip(sample_train_users, sample_train_movies, sample_train_ratings):
            st = datetime.now()
            # print(user, movie)
            #----- Ratings of "movie" by similar users of "user" -----
            # compute the similar Users of the "user"
            user_sim = cosine_similarity(sample_train_sparse_matrix[user], sample_train_sparse_matrix[movie])
            top_sim_users = user_sim.argsort()[::-1][1:] # we are ignoring 'The User' from
            # get the ratings of most similar users for this movie
            top_ratings = sample_train_sparse_matrix[top_sim_users, movie].toarray().ravel()
            # we will make it's length "5" by adding movie averages to .
            top_sim_users_ratings = list(top_ratings[top_ratings != 0][:5])
            top_sim_users_ratings.extend([sample_train_averages['movie'][movie]]*(5 - len(top_sim_users_ratings)))
            # print(top_sim_users_ratings, end=" ")

            #----- Ratings by "user" to similar movies of "movie" -----
            # compute the similar movies of the "movie"
            movie_sim = cosine_similarity(sample_train_sparse_matrix[:,movie].T, sample_train_sparse_matrix[:,user])
            top_sim_movies = movie_sim.argsort()[::-1][1:] # we are ignoring 'The User' from
            # get the ratings of most similar movie rated by this user..
            top_ratings = sample_train_sparse_matrix[user, top_sim_movies].toarray().ravel()
            # we will make it's length "5" by adding user averages to.
            top_sim_movies_ratings = list(top_ratings[top_ratings != 0][:5])
            top_sim_movies_ratings.extend([sample_train_averages['user'][user]]*(5-len(top_sim_movies_ratings)))
            # print(top_sim_movies_ratings, end=" : -- ")

            #-----prepare the row to be stores in a file-----#
            row = list()
            row.append(user)
            row.append(movie)
            # Now add the other features to this data...
            row.append(sample_train_averages['global']) # first feature
            # next 5 features are similar_users "movie" ratings
            row.extend(top_sim_users_ratings)
            # next 5 features are "user" ratings for similar_movies
            row.extend(top_sim_movies_ratings)
            # Avg_user rating
            row.append(sample_train_averages['user'][user])
            # Avg_movie rating
            row.append(sample_train_averages['movie'][movie])

            # finalley, The actual Rating of this user-movie pair...
            row.append(rating)
            count = count + 1

            # add rows to the file opened..
            reg_data_file.write(','.join(map(str, row)))
            reg_data_file.write('\n')
            if (count)%10000 == 0:
```

```
# print(', '.join(map(str, row)))
print("Done for {} rows----- {}".format(count, datetime.now() - start))
```

```
print(datetime.now() - start)
```

File already exists you don't have to prepare again...
0:00:00.000998

Reading from the file to make a Train_dataframe

In [89]:

```
reg_train = pd.read_csv('reg_train.csv', names = ['user', 'movie', 'GAvg', 'sur1', 'sur2',
reg_train.head()
```

Out[89]:

	user	movie	GAvg	sur1	sur2	sur3	sur4	sur5	smr1	smr2	smr3	smr4	smr5
0	53406	33	3.581679	4.0	5.0	5.0	4.0	1.0	5.0	2.0	5.0	3.0	1.0
1	99540	33	3.581679	5.0	5.0	5.0	4.0	5.0	3.0	4.0	4.0	3.0	5.0
2	99865	33	3.581679	5.0	5.0	4.0	5.0	3.0	5.0	4.0	4.0	5.0	4.0
3	101620	33	3.581679	2.0	3.0	5.0	5.0	4.0	4.0	3.0	3.0	4.0	5.0
4	112974	33	3.581679	5.0	5.0	5.0	5.0	5.0	3.0	5.0	5.0	5.0	3.0

- **GAvg** : Average rating of all the ratings
- **Similar users rating of this movie:**
 - sur1, sur2, sur3, sur4, sur5 (top 5 similar users who rated that movie..)
- **Similar movies rated by this user:**
 - smr1, smr2, smr3, smr4, smr5 (top 5 similar movies rated by this movie..)
- **UAvg** : User's Average rating
- **MAvg** : Average rating of this movie
- **rating** : Rating of this movie by this user.

4.3.1.2 Featurizing test data

In [90]:

```
# get users, movies and ratings from the Sampled Test
sample_test_users, sample_test_movies, sample_test_ratings = sparse.find(sample_test_sparse
```


In [91]:

```
sample_train_averages['global']
```

Out[91]:

3.5744906003423225

In [92]:

```

start = datetime.now()

if os.path.isfile('reg_test.csv'):
    print("It is already created...")
else:

    print('preparing {} tuples for the dataset..\\n'.format(len(sample_test_ratings)))
    with open('reg_test.csv', mode='w') as reg_data_file:
        count = 0
        for (user, movie, rating) in zip(sample_test_users, sample_test_movies, sample_test_ratings):
            st = datetime.now()

            #----- Ratings of "movie" by similar users of "user" -----
            #print(user, movie)
            try:
                # compute the similar Users of the "user"
                user_sim = cosine_similarity(sample_train_sparse_matrix[user], sample_train_sparse_matrix)
                top_sim_users = user_sim.argsort()[::-1][1:] # we are ignoring 'The User' first
                # get the ratings of most similar users for this movie
                top_ratings = sample_train_sparse_matrix[top_sim_users, movie].toarray().ravel()
                # we will make it's length "5" by adding movie averages to .
                top_sim_users_ratings = list(top_ratings[top_ratings != 0][:5])
                top_sim_users_ratings.extend([sample_train_averages['movie'][movie]]*(5 - len(top_sim_users_ratings)))
                # print(top_sim_users_ratings, end="--")

            except (IndexError, KeyError):
                # It is a new User or new Movie or there are no ratings for given user for
                ##### Cold Start Problem #####
                top_sim_users_ratings.extend([sample_train_averages['global']]*(5 - len(top_sim_users_ratings)))
                #print(top_sim_users_ratings)
            except:
                print(user, movie)
                # we just want KeyErrors to be resolved. Not every Exception...
                raise

            #----- Ratings by "user" to similar movies of "movie" -----
            try:
                # compute the similar movies of the "movie"
                movie_sim = cosine_similarity(sample_train_sparse_matrix[:,movie].T, sample_train_sparse_matrix)
                top_sim_movies = movie_sim.argsort()[::-1][1:] # we are ignoring 'The User' first
                # get the ratings of most similar movie rated by this user..
                top_ratings = sample_train_sparse_matrix[user, top_sim_movies].toarray().ravel()
                # we will make it's length "5" by adding user averages to .
                top_sim_movies_ratings = list(top_ratings[top_ratings != 0][:5])
                top_sim_movies_ratings.extend([sample_train_averages['user'][user]]*(5-len(top_sim_movies_ratings)))
                #print(top_sim_movies_ratings)
            except (IndexError, KeyError):
                #print(top_sim_movies_ratings, end=" : -- ")
                top_sim_movies_ratings.extend([sample_train_averages['global']]*(5-len(top_sim_movies_ratings)))
                #print(top_sim_movies_ratings)
            except :
                raise

            #-----prepare the row to be stores in a file-----#
            row = list()
            # add usser and movie name first
            row.append(user)

```

```

row.append(movie)
row.append(sample_train_averages['global']) # first feature
#print(row)
# next 5 features are similar_users "movie" ratings
row.extend(top_sim_users_ratings)
#print(row)
# next 5 features are "user" ratings for similar_movies
row.extend(top_sim_movies_ratings)
#print(row)
# Avg_user rating
try:
    row.append(sample_train_averages['user'][user])
except KeyError:
    row.append(sample_train_averages['global'])
except:
    raise
#print(row)
# Avg_movie rating
try:
    row.append(sample_train_averages['movie'][movie])
except KeyError:
    row.append(sample_train_averages['global'])
except:
    raise
#print(row)
# finalley, The actual Rating of this user-movie pair...
row.append(rating)
#print(row)
count = count + 1

# add rows to the file opened..
reg_data_file.write(','.join(map(str, row)))
#print(','.join(map(str, row)))
reg_data_file.write('\n')
if (count)%1000 == 0:
    #print(','.join(map(str, row)))
    print("Done for {} rows----- {}".format(count, datetime.now() - start))
print("",datetime.now() - start)

```

It is already created...

__Reading from the file to make a test dataframe __

In [93]:

```
reg_test_df = pd.read_csv('reg_test.csv', names = ['user', 'movie', 'GAvg', 'sur1', 'sur2',
                                                    'smr1', 'smr2', 'smr3', 'smr4', 'UAvg', 'MAvg', 'rating'], header=0)
reg_test_df.head(4)
```

Out[93]:

	user	movie	GAvg	sur1	sur2	sur3	sur4	sur5	smr1	smr2	smr3	smr4	rating
0	808635	71	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.58
1	941866	71	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.58
2	1737912	71	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.58
3	1849204	71	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.58

- **GAvg** : Average rating of all the ratings
- **Similar users rating of this movie:**
 - sur1, sur2, sur3, sur4, sur5 (top 5 similar users who rated that movie..)
- **Similar movies rated by this user:**
 - smr1, smr2, smr3, smr4, smr5 (top 5 similar movies rated by this movie..)
- **UAvg** : User Average rating
- **MAvg** : Average rating of this movie
- **rating** : Rating of this movie by this user.

4.3.2 Transforming data for Surprise models

In [94]:

```
from surprise import Reader, Dataset
```

4.3.2.1 Transforming train data

- We can't give raw data (movie, user, rating) to train the model in Surprise library.
- They have a separate format for TRAIN and TEST data, which will be useful for training the models like SVD, KNNBaseLineOnly....etc., in Surprise.
- We can form the trainset from a file, or from a Pandas DataFrame.
http://surprise.readthedocs.io/en/stable/getting_started.html#load-dom-dataframe-py
[\(http://surprise.readthedocs.io/en/stable/getting_started.html#load-dom-dataframe-py\)](http://surprise.readthedocs.io/en/stable/getting_started.html#load-dom-dataframe-py)

In [95]:

```
# It is to specify how to read the dataframe.
# for our dataframe, we don't have to specify anything extra..
reader = Reader(rating_scale=(1,5))

# create the traindata from the dataframe...
train_data = Dataset.load_from_df(reg_train[['user', 'movie', 'rating']], reader)

# build the trainset from traindata.. It is of dataset format from surprise library..
trainset = train_data.build_full_trainset()
```

4.3.2.2 Transforming test data

- Testset is just a list of (user, movie, rating) tuples. (Order in the tuple is important)

In [96]:

```
testset = list(zip(reg_test_df.user.values, reg_test_df.movie.values, reg_test_df.rating.values))
testset[:3]
```

Out[96]:

```
[(808635, 71, 5), (941866, 71, 4), (1737912, 71, 3)]
```

4.4 Applying Machine Learning models

- Global dictionary that stores rmse and mape for all the models....
 - It stores the metrics in a dictionary of dictionaries

keys : model names(string)

value: dict(**key** : metric, **value** : value)

In [97]:

```
models_evaluation_train = dict()
models_evaluation_test = dict()

models_evaluation_train, models_evaluation_test
```

Out[97]:

```
({}, {})
```

Utility functions for running regression models

In [98]:

```

# to get rmse and mape given actual and predicted ratings..
def get_error_metrics(y_true, y_pred):
    rmse = np.sqrt(np.mean([ (y_true[i] - y_pred[i])**2 for i in range(len(y_pred)) ]))
    mape = np.mean(np.abs( (y_true - y_pred)/y_true )) * 100
    return rmse, mape

#####
#####
def run_xgboost(algo, x_train, y_train, x_test, y_test, verbose=True):
    """
    It will return train_results and test_results
    """

    # dictionaries for storing train and test results
    train_results = dict()
    test_results = dict()

    # fit the model
    print('Training the model..')
    start = datetime.now()
    algo.fit(x_train, y_train, eval_metric = 'rmse')
    print('Done. Time taken : {}'.format(datetime.now()-start))
    print('Done \n')

    # from the trained model, get the predictions....
    print('Evaluating the model with TRAIN data...')
    start = datetime.now()
    y_train_pred = algo.predict(x_train)
    # get the rmse and mape of train data...
    rmse_train, mape_train = get_error_metrics(y_train.values, y_train_pred)

    # store the results in train_results dictionary..
    train_results = {'rmse': rmse_train,
                    'mape' : mape_train,
                    'predictions' : y_train_pred}

    #####
    # get the test data predictions and compute rmse and mape
    print('Evaluating Test data')
    y_test_pred = algo.predict(x_test)
    rmse_test, mape_test = get_error_metrics(y_true=y_test.values, y_pred=y_test_pred)
    # store them in our test results dictionary.
    test_results = {'rmse': rmse_test,
                   'mape' : mape_test,
                   'predictions':y_test_pred}

    if verbose:
        print('\nTEST DATA')
        print('-'*30)
        print('RMSE : ', rmse_test)
        print('MAPE : ', mape_test)

    # return these train and test results...
    return train_results, test_results

```


In [99]:

```

# it is just to make sure that all of our algorithms should produce same results
# everytime they run...

my_seed = 15
random.seed(my_seed)
np.random.seed(my_seed)

#####
# get (actual_list , predicted_list) ratings given list
# of predictions (prediction is a class in Surprise).
#####
def get_ratings(predictions):
    actual = np.array([pred.r_ui for pred in predictions])
    pred = np.array([pred.est for pred in predictions])

    return actual, pred

#####
# get 'rmse' and 'mape' , given list of prediction objects
#####
def get_errors(predictions, print_them=False):

    actual, pred = get_ratings(predictions)
    rmse = np.sqrt(np.mean((pred - actual)**2))
    mape = np.mean(np.abs(pred - actual)/actual)

    return rmse, mape*100

#####
# It will return predicted ratings, rmse and mape of both train and test data #
#####
def run_surprise(algo, trainset, testset, verbose=True):
    '''
        return train_dict, test_dict

        It returns two dictionaries, one for train and the other is for test
        Each of them have 3 key-value pairs, which specify 'rmse', 'mape', and 'predic
    '''
    start = datetime.now()
    # dictionaries that stores metrics for train and test..
    train = dict()
    test = dict()

    # train the algorithm with the trainset
    st = datetime.now()
    print('Training the model...')
    algo.fit(trainset)
    print('Done. time taken : {} \n'.format(datetime.now()-st))

    # ----- Evaluating train data-----#
    st = datetime.now()
    print('Evaluating the model with train data..')
    # get the train predictions (list of prediction class inside Surprise)
    train_preds = algo.test(trainset.build_testset())
    # get predicted ratings from the train predictions..
    train_actual_ratings, train_pred_ratings = get_ratings(train_preds)
    # get 'rmse' and 'mape' from the train predictions.
    train_rmse, train_mape = get_errors(train_preds)
    print('time taken : {}'.format(datetime.now()-st))

```



```

if verbose:
    print('-'*15)
    print('Train Data')
    print('-'*15)
    print("RMSE : {}\nMAPE : {}".format(train_rmse, train_mape))

#store them in the train dictionary
if verbose:
    print('adding train results in the dictionary..')
train['rmse'] = train_rmse
train['mape'] = train_mape
train['predictions'] = train_pred_ratings

#----- Evaluating Test data-----#
st = datetime.now()
print('\nEvaluating for test data...')
# get the predictions( list of prediction classes) of test data
test_preds = algo.test(testset)
# get the predicted ratings from the list of predictions
test_actual_ratings, test_pred_ratings = get_ratings(test_preds)
# get error metrics from the predicted and actual ratings
test_rmse, test_mape = get_errors(test_preds)
print('time taken : {}'.format(datetime.now()-st))

if verbose:
    print('-'*15)
    print('Test Data')
    print('-'*15)
    print("RMSE : {}\nMAPE : {}".format(test_rmse, test_mape))
# store them in test dictionary
if verbose:
    print('storing the test results in test dictionary...')
test['rmse'] = test_rmse
test['mape'] = test_mape
test['predictions'] = test_pred_ratings

print('\n'+ '-'*45)
print('Total time taken to run this algorithm :', datetime.now() - start)

# return two dictionaries train and test
return train, test

```

4.4.1 XGBoost with initial 13 features **With Hyperparameter tuning**

In [100]:

```

# prepare Train data
x_train = reg_train.drop(['user', 'movie', 'rating'], axis=1)
y_train = reg_train['rating']

# Prepare Test data
x_test = reg_test_df.drop(['user', 'movie', 'rating'], axis=1)
y_test = reg_test_df['rating']

```

In [83]:

```

from sklearn.model_selection import GridSearchCV
import xgboost as xgb

xgb1 = xgb.XGBRegressor()
parameters = {'nthread':[1,2,3,4], #when use hyperthread, xgboost may become slower
              'learning_rate': [0.01,0.001,0.0001,0.03, 0.05, 0.07], #so called `eta` value
              'max_depth': [2, 3, 5, 6, 7,10, 12],
              'n_estimators': [20, 50, 100, 200, 300,500, 700, 1000]}

xgb_grid = GridSearchCV(xgb1,parameters,cv = 2,n_jobs = 5,verbose=True)

xgb_grid.fit(x_train,y_train)

print(xgb_grid.best_score_)
print(xgb_grid.best_params_)

```

Fitting 2 folds for each of 1344 candidates, totalling 2688 fits

```

[Parallel(n_jobs=5)]: Using backend LokyBackend with 5 concurrent workers.
[Parallel(n_jobs=5)]: Done 40 tasks      | elapsed: 36.8s
[Parallel(n_jobs=5)]: Done 190 tasks    | elapsed: 9.6min
[Parallel(n_jobs=5)]: Done 440 tasks    | elapsed: 41.1min
[Parallel(n_jobs=5)]: Done 790 tasks    | elapsed: 65.9min
[Parallel(n_jobs=5)]: Done 1240 tasks   | elapsed: 107.4min
[Parallel(n_jobs=5)]: Done 1790 tasks   | elapsed: 168.0min
[Parallel(n_jobs=5)]: Done 2440 tasks   | elapsed: 219.1min
[Parallel(n_jobs=5)]: Done 2688 out of 2688 | elapsed: 250.9min finished
C:\Users\vansh\Anaconda3\envs\env\lib\site-packages\xgboost\core.py:587: FutureWarning: Series.base is deprecated and will be removed in a future version
    if getattr(data, 'base', None) is not None and \
C:\Users\vansh\Anaconda3\envs\env\lib\site-packages\xgboost\core.py:588: FutureWarning: Series.base is deprecated and will be removed in a future version
    data.base is not None and isinstance(data, np.ndarray) \

[03:10:58] WARNING: C:/Jenkins/workspace/xgboost-win64_release_0.90/src/objective/regression_obj.cu:152: reg:linear is now deprecated in favor of reg:squarederror.
0.37613024354760494
{'learning_rate': 0.07, 'max_depth': 5, 'n_estimators': 500, 'nthread': 1}

```

In [102]:

```
# initialize Our first XGBoost model...
import xgboost as xgb
first_xgb = xgb.XGBRegressor(silent=False, learning_rate=0.07, max_depth= 5, n_estimators=500)
train_results, test_results = run_xgboost(first_xgb, x_train, y_train, x_test, y_test)

# store the results in models_evaluations dictionaries
models_evaluation_train['first_algo'] = train_results
models_evaluation_test['first_algo'] = test_results

xgb.plot_importance(first_xgb)
plt.show()
```

Training the model..

[22:53:05] WARNING: C:/Jenkins/workspace/xgboost-win64_release_0.90/src/objective/regression_obj.cu:152: reg:linear is now deprecated in favor of reg:squarederror.

C:\Users\vansh\Anaconda3\envs\env\lib\site-packages\xgboost\core.py:587: FutureWarning: Series.base is deprecated and will be removed in a future version

if getattr(data, 'base', None) is not None and \n C:\Users\vansh\Anaconda3\envs\env\lib\site-packages\xgboost\core.py:588: FutureWarning: Series.base is deprecated and will be removed in a future version

data.base is not None and isinstance(data, np.ndarray) \

Done. Time taken : 0:00:48.943570

Done

Evaluating the model with TRAIN data...


Evaluating Test data

TEST DATA

RMSE : 1.0846515217011976
MAPE : 33.98071389809766



Figure 9



4.4.2 Surprise BaselineModel

In [103]:

```
from surprise import BaselineOnly
```

__Predicted_rating : (baseline prediction) __

- http://surprise.readthedocs.io/en/stable/basic_algorithms.html#surprise.prediction_algorithms.baseline_only.BaselineOnly

$$\hat{r}_{ui} = b_{ui} = \mu + b_u + b_i$$

- μ : Average of all trainings in training data.
- b_u : User bias
- b_i : Item bias (movie biases)

__Optimization function (Least Squares Problem) __

- http://surprise.readthedocs.io/en/stable/prediction_algorithms.html#baselines-estimates-configuration

$$\sum_{r_{ui} \in R_{train}} (r_{ui} - (\mu + b_u + b_i))^2 + \lambda (b_u^2 + b_i^2) . \text{ [mimimize } b_u, b_i]$$

In [104]:

```
# options are to specify.., how to compute those user and item biases
bsl_options = {'method': 'sgd',
               'learning_rate': .001
              }
bsl_algo = BaselineOnly(bsl_options=bsl_options)
# run this algorithm.., It will return the train and test results..
bsl_train_results, bsl_test_results = run_surprise(bsl_algo, trainset, testset, verbose=True)

# Just store these error metrics in our models_evaluation datastructure
models_evaluation_train['bsl_algo'] = bsl_train_results
models_evaluation_test['bsl_algo'] = bsl_test_results
```

Training the model...

Estimating biases using sgd...

Done. time taken : 0:00:00.665440

Evaluating the model with train data..

time taken : 0:00:00.784440

Train Data

RMSE : 0.9347153928678286

MAPE : 29.389572652358183

adding train results in the dictionary..

Evaluating for test data...

time taken : 0:00:00.054853

Test Data

RMSE : 1.0730330260516174

MAPE : 35.04995544572911

storing the test results in test dictionary...

Total time taken to run this algorithm : 0:00:01.505730

4.4.3 XGBoost with initial 13 features + Surprise Baseline predictor with hyperparameter tuning

Updating Train Data

In [105]:

```
# add our baseline_predicted value as our feature..
reg_train['bslpr'] = models_evaluation_train['bsl_algo']['predictions']
reg_train.head(2)
```

Out[105]:

	user	movie	GAvg	sur1	sur2	sur3	sur4	sur5	smr1	smr2	smr3	smr4	smr5	
0	53406	33	3.581679	4.0	5.0	5.0	4.0	1.0	5.0	2.0	5.0	3.0	1.0	3.37
1	99540	33	3.581679	5.0	5.0	5.0	4.0	5.0	3.0	4.0	4.0	3.0	5.0	3.55

Updating Test Data

In [106]:

```
# add that baseline predicted ratings with Surprise to the test data as well
reg_test_df['bslpr'] = models_evaluation_test['bsl_algo']['predictions']
reg_test_df.head(2)
```

Out[106]:

	user	movie	GAvg	sur1	sur2	sur3	sur4	sur5	smr1	smr2	smr3	smr4	smr5	si
0	808635	71	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679
1	941866	71	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679

In [108]:

```

from sklearn.model_selection import GridSearchCV
# prepare train data
x_train = reg_train.drop(['user', 'movie', 'rating'], axis=1)
y_train = reg_train['rating']

# Prepare Test data
x_test = reg_test_df.drop(['user', 'movie', 'rating'], axis=1)
y_test = reg_test_df['rating']

xgb1 = xgb.XGBRegressor()
parameters = {'nthread': [1, 2, 3, 4], #when use hyperthread, xgboost may become slower
              'learning_rate': [0.01, 0.001, 0.0001, 0.03, 0.05, 0.07], #so called `eta` value
              'max_depth': [2, 3, 5, 6, 7, 10, 12],
              'n_estimators': [20, 50, 100, 200, 300, 500, 700, 1000]}

xgb_grid = GridSearchCV(xgb1, parameters, cv = 2, n_jobs = 5, verbose=True)

xgb_grid.fit(x_train, y_train)

print(xgb_grid.best_score_)
print(xgb_grid.best_params_)

```

Fitting 2 folds for each of 1344 candidates, totalling 2688 fits

```

[Parallel(n_jobs=5)]: Using backend LokyBackend with 5 concurrent workers.
[Parallel(n_jobs=5)]: Done 40 tasks      | elapsed: 53.4s
[Parallel(n_jobs=5)]: Done 190 tasks    | elapsed: 12.4min
[Parallel(n_jobs=5)]: Done 440 tasks    | elapsed: 58.7min
[Parallel(n_jobs=5)]: Done 790 tasks    | elapsed: 97.2min
[Parallel(n_jobs=5)]: Done 1240 tasks   | elapsed: 159.0min
[Parallel(n_jobs=5)]: Done 1790 tasks   | elapsed: 601.4min
[Parallel(n_jobs=5)]: Done 2440 tasks   | elapsed: 696.6min
[Parallel(n_jobs=5)]: Done 2688 out of 2688 | elapsed: 738.7min finished

```

```

[11:13:26] WARNING: C:/Jenkins/workspace/xgboost-win64_release_0.90/src/objective/regression_obj.cu:152: reg:linear is now deprecated in favor of reg:squarederror.
0.3745085331359887
{'learning_rate': 0.05, 'max_depth': 5, 'n_estimators': 200, 'nthread': 1}

```

In [109]:

```
# initialize Our first XGBoost model...
xgb_bsl = xgb.XGBRegressor(silent=False, learning_rate= 0.05, max_depth= 5, n_estimators=200)
train_results, test_results = run_xgboost(xgb_bsl, x_train, y_train, x_test, y_test)

# store the results in models_evaluations dictionaries
models_evaluation_train['xgb_bsl'] = train_results
models_evaluation_test['xgb_bsl'] = test_results

xgb.plot_importance(xgb_bsl)
plt.show()
```

Training the model..

[11:57:36] WARNING: C:/Jenkins/workspace/xgboost-win64_release_0.90/src/objective/regression_obj.cu:152: reg:linear is now deprecated in favor of reg:squarederror.

Done. Time taken : 0:00:55.811345

Done

Evaluating the model with TRAIN data...


Evaluating Test data

TEST DATA

RMSE : 1.0775409822076403

MAPE : 34.39214865567116

Figure 10



4.4.4 Surprise KNNBaseline predictor

In [111]:

```
from surprise import KNNBaseline
```

- KNN BASELINE

- http://surprise.readthedocs.io/en/stable/knn_inspired.html#surprise.prediction_algorithms.knns.KNNBaseline
(http://surprise.readthedocs.io/en/stable/knn_inspired.html#surprise.prediction_algorithms.knns.KNNBaseline)

- PEARSON_BASELINE SIMILARITY

- http://surprise.readthedocs.io/en/stable/similarities.html#surprise.similarities.pearson_baseline
(http://surprise.readthedocs.io/en/stable/similarities.html#surprise.similarities.pearson_baseline)

- SHRINKAGE

- 2.2 Neighborhood Models in <http://courses.ischool.berkeley.edu/i290-dm/s11/SECURE/a1-koren.pdf>
(<http://courses.ischool.berkeley.edu/i290-dm/s11/SECURE/a1-koren.pdf>)

- predicted Rating : (_ based on User-User similarity _)

$$\hat{r}_{ui} = b_{ui} + \frac{\sum_{v \in N_i^k(u)} \text{sim}(u, v) \cdot (r_{vi} - b_{vi})}{\sum_{v \in N_i^k(u)} \text{sim}(u, v)}$$

- b_{ui} - Baseline prediction of (user, movie) rating
- $N_i^k(u)$ - Set of **K similar** users (neighbours) of **user (u)** who rated **movie(i)**
- $\text{sim}(u, v)$ - **Similarity** between users **u** and **v**
 - Generally, it will be cosine similarity or Pearson correlation coefficient.
 - But we use **shrunk Pearson-baseline correlation coefficient**, which is based on the pearsonBaseline similarity (we take base line predictions instead of mean rating of user/item)

- __ Predicted rating __ (based on Item Item similarity):

$$\hat{r}_{ui} = b_{ui} + \frac{\sum_{j \in N_u^k(i)} \text{sim}(i, j) \cdot (r_{uj} - b_{uj})}{\sum_{j \in N_u^k(i)} \text{sim}(i, j)}$$

- _Notations follows same as above (user user based predicted rating) _

4.4.4.1 Surprise KNNBaseline with user user similarities

In [112]:

```
# we specify , how to compute similarities and what to consider with sim_options to our alg
sim_options = {'user_based' : True,
               'name': 'pearson_baseline',
               'shrinkage': 100,
               'min_support': 2
              }
# we keep other parameters like regularization parameter and learning_rate as default value
bsl_options = {'method': 'sgd'}

knn_bsl_u = KNNBaseline(k=40, sim_options = sim_options, bsl_options = bsl_options)
knn_bsl_u_train_results, knn_bsl_u_test_results = run_surprise(knn_bsl_u, trainset, testset)

# Just store these error metrics in our models_evaluation datastructure
models_evaluation_train['knn_bsl_u'] = knn_bsl_u_train_results
models_evaluation_test['knn_bsl_u'] = knn_bsl_u_test_results
```

Training the model...
 Estimating biases using sgd...
 Computing the pearson_baseline similarity matrix...
 Done computing similarity matrix.
 Done. time taken : 0:01:20.847121

Evaluating the model with train data..
 time taken : 0:14:45.341548

 Train Data

 RMSE : 0.33642097416508826

MAPE : 9.145093375416348

adding train results in the dictionary..

Evaluating for test data...
 time taken : 0:00:00.057123

 Test Data

 RMSE : 1.0726493739667242

MAPE : 35.02094499698424

storing the test results in test dictionary...

 Total time taken to run this algorithm : 0:16:06.249791

4.4.4.2 Surprise KNNBaseline with movie movie similarities

In [113]:

```
# we specify , how to compute similarities and what to consider with sim_options to our alg
# 'user_based' : Fals => this considers the similarities of movies instead of users

sim_options = {'user_based' : False,
               'name': 'pearson_baseline',
               'shrinkage': 100,
               'min_support': 2
              }
# we keep other parameters like regularization parameter and learning_rate as default value
bsl_options = {'method': 'sgd'}

knn_bsl_m = KNNBaseline(k=40, sim_options = sim_options, bsl_options = bsl_options)

knn_bsl_m_train_results, knn_bsl_m_test_results = run_surprise(knn_bsl_m, trainset, testset)

# Just store these error metrics in our models_evaluation datastructure
models_evaluation_train['knn_bsl_m'] = knn_bsl_m_train_results
models_evaluation_test['knn_bsl_m'] = knn_bsl_m_test_results
```

Training the model...
 Estimating biases using sgd...
 Computing the pearson_baseline similarity matrix...
 Done computing similarity matrix.
 Done. time taken : 0:00:01.000853

Evaluating the model with train data..
 time taken : 0:00:07.087081

 Train Data

 RMSE : 0.32584796251610554

MAPE : 8.447062581998374

adding train results in the dictionary..

Evaluating for test data...
 time taken : 0:00:00.053972

 Test Data

 RMSE : 1.072758832653683

MAPE : 35.02269653015042

storing the test results in test dictionary...

 Total time taken to run this algorithm : 0:00:08.141906

4.4.5 XGBoost with initial 13 features + Surprise Baseline predictor + KNNBaseline predictor

- First we will run XGBoost with predictions from both KNN's (that uses User_User and Item_Item similarities along with our previous features.
- Then we will run XGBoost with just predictions form both knn models and preditions from our baseline model.

__Preparing Train data __

In [114]:

```
# add the predicted values from both knns to this dataframe
reg_train['knn_bsl_u'] = models_evaluation_train['knn_bsl_u']['predictions']
reg_train['knn_bsl_m'] = models_evaluation_train['knn_bsl_m']['predictions']

reg_train.head(2)
```

Out[114]:

	user	movie	GAvg	sur1	sur2	sur3	sur4	sur5	smr1	smr2	smr3	smr4	smr5	
0	53406	33	3.581679	4.0	5.0	5.0	4.0	1.0	5.0	2.0	5.0	3.0	1.0	3.37
1	99540	33	3.581679	5.0	5.0	5.0	4.0	5.0	3.0	4.0	4.0	3.0	5.0	3.55

__Preparing Test data __

In [115]:

```
reg_test_df['knn_bsl_u'] = models_evaluation_test['knn_bsl_u']['predictions']
reg_test_df['knn_bsl_m'] = models_evaluation_test['knn_bsl_m']['predictions']

reg_test_df.head(2)
```

Out[115]:

	user	movie	GAvg	sur1	sur2	sur3	sur4	sur5	smr1	smr2	smr3	smr4	smr5	
0	808635	71	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	
1	941866	71	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	

Hyperparameter tuning

In [116]:

```

# prepare the train data....
x_train = reg_train.drop(['user', 'movie', 'rating'], axis=1)
y_train = reg_train['rating']

# prepare the train data....
x_test = reg_test_df.drop(['user', 'movie', 'rating'], axis=1)
y_test = reg_test_df['rating']

xgb1 = xgb.XGBRegressor()
parameters = {'nthread': [1, 2, 3, 4], #when use hyperthread, xgboost may become slower
              'learning_rate': [0.01, 0.001, 0.0001, 0.03, 0.05, 0.07], #so called `eta` value
              'max_depth': [2, 3, 5, 6, 7, 10, 12],
              'n_estimators': [20, 50, 100, 200, 300, 500, 700, 1000]}

xgb_grid = GridSearchCV(xgb1, parameters, cv = 2, n_jobs = 5, verbose=True)

xgb_grid.fit(x_train, y_train)

print(xgb_grid.best_score_)
print(xgb_grid.best_params_)

```

Fitting 2 folds for each of 1344 candidates, totalling 2688 fits

```

[Parallel(n_jobs=5)]: Using backend LokyBackend with 5 concurrent workers.
[Parallel(n_jobs=5)]: Done 40 tasks      | elapsed: 1.7min
[Parallel(n_jobs=5)]: Done 190 tasks    | elapsed: 31.4min
[Parallel(n_jobs=5)]: Done 440 tasks    | elapsed: 97.9min
[Parallel(n_jobs=5)]: Done 790 tasks    | elapsed: 137.9min
[Parallel(n_jobs=5)]: Done 1240 tasks   | elapsed: 206.5min
[Parallel(n_jobs=5)]: Done 1790 tasks   | elapsed: 305.3min
[Parallel(n_jobs=5)]: Done 2440 tasks   | elapsed: 389.0min
[Parallel(n_jobs=5)]: Done 2688 out of 2688 | elapsed: 446.9min finished

[19:43:26] WARNING: C:/Jenkins/workspace/xgboost-win64_release_0.90/src/objective/regression_obj.cu:152: reg:linear is now deprecated in favor of reg:squarederror.
0.3738852164946025
{'learning_rate': 0.03, 'max_depth': 5, 'n_estimators': 300, 'nthread': 1}

```

In [117]:

```
# declare the model
xgb_knn_bsl = xgb.XGBRegressor(silent=False, learning_rate= 0.03, max_depth= 5, n_estimators= 100)
train_results, test_results = run_xgboost(xgb_knn_bsl, x_train, y_train, x_test, y_test)

# store the results in models_evaluations dictionaries
models_evaluation_train['xgb_knn_bsl'] = train_results
models_evaluation_test['xgb_knn_bsl'] = test_results

xgb.plot_importance(xgb_knn_bsl)
plt.show()
```

Training the model..

[19:47:48] WARNING: C:/Jenkins/workspace/xgboost-win64_release_0.90/src/objective/regression_obj.cu:152: reg:linear is now deprecated in favor of reg:squarederror.

Done. Time taken : 0:01:59.718489

Done

Evaluating the model with TRAIN data...

Evaluating Test data


TEST DATA

RMSE : 1.0773481404411078

MAPE : 34.40629706111159



Figure 11



4.4.6 Matrix Factorization Techniques

4.4.6.1 SVD Matrix Factorization User Movie interactions

In [118]:

```
from surprise import SVD
```

http://surprise.readthedocs.io/en/stable/matrix_factorization.html#surprise.prediction_algorithms.matrix_factorization
http://surprise.readthedocs.io/en/stable/matrix_factorization.html#surprise.prediction_algorithms.matrix_factorization

- ___ Predicted Rating : ___

-

- $\hat{r}_{ui} = \mu + b_u + b_i + q_i^T p_u$

- q_i - Representation of item(movie) in latent factor space
- p_u - Representation of user in new latent factor space

- A BASIC MATRIX FACTORIZATION MODEL in [https://datajobs.com/data-science-repo/Recommender-Systems-\[Netflix\].pdf](https://datajobs.com/data-science-repo/Recommender-Systems-[Netflix].pdf) (<https://datajobs.com/data-science-repo/Recommender-Systems-%5BNetflix%5D.pdf>)

- **Optimization problem with user item interactions and regularization (to avoid overfitting)**

- - $$\sum_{r_{ui} \in R_{train}} (r_{ui} - \hat{r}_{ui})^2 + \lambda (b_i^2 + b_u^2 + ||q_i||^2 + ||p_u||^2)$$

In [119]:

```
# initialize the model
svd = SVD(n_factors=100, biased=True, random_state=15, verbose=True)
svd_train_results, svd_test_results = run_surprise(svd, trainset, testset, verbose=True)

# Just store these error metrics in our models_evaluation datastructure
models_evaluation_train['svd'] = svd_train_results
models_evaluation_test['svd'] = svd_test_results
```

Training the model...

Processing epoch 0
 Processing epoch 1
 Processing epoch 2
 Processing epoch 3
 Processing epoch 4
 Processing epoch 5
 Processing epoch 6
 Processing epoch 7
 Processing epoch 8
 Processing epoch 9
 Processing epoch 10
 Processing epoch 11
 Processing epoch 12
 Processing epoch 13
 Processing epoch 14
 Processing epoch 15
 Processing epoch 16
 Processing epoch 17
 Processing epoch 18
 Processing epoch 19

Done. time taken : 0:00:17.353360

Evaluating the model with train data..

time taken : 0:00:03.231120

Train Data

RMSE : 0.6574721240954099

MAPE : 19.704901088660478

adding train results in the dictionary..

Evaluating for test data...

time taken : 0:00:00.140223

Test Data

RMSE : 1.0726046873826458

MAPE : 35.01953535988152

storing the test results in test dictionary...

Total time taken to run this algorithm : 0:00:20.729877

4.4.6.2 SVD Matrix Factorization with implicit feedback from user (user rated movies)

In [120]:

```
from surprise import SVDpp
```

- -----> 2.5 Implicit Feedback in <http://courses.ischool.berkeley.edu/i290-dm/s11/SECURE/a1-koren.pdf>
(<http://courses.ischool.berkeley.edu/i290-dm/s11/SECURE/a1-koren.pdf>)
- __ Predicted Rating : __
 - $$\hat{r}_{ui} = \mu + b_u + b_i + q_i^T \left(p_u + |I_u|^{-\frac{1}{2}} \sum_{j \in I_u} y_j \right)$$
- I_u --- the set of all items rated by user u
- y_j --- Our new set of item factors that capture implicit ratings.
- **Optimization problem with user item interactions and regularization (to avoid overfitting)**
 - $$\sum_{r_{ui} \in R_{train}} (r_{ui} - \hat{r}_{ui})^2 + \lambda (b_i^2 + b_u^2 + ||q_i||^2 + ||p_u||^2 + ||y_j||^2)$$

In [121]:

```
# initialize the model
svdpp = SVDpp(n_factors=50, random_state=15, verbose=True)
svdpp_train_results, svdpp_test_results = run_surprise(svdpp, trainset, testset, verbose=True)

# Just store these error metrics in our models_evaluation datastructure
models_evaluation_train['svdpp'] = svdpp_train_results
models_evaluation_test['svdpp'] = svdpp_test_results
```

Training the model...

```
processing epoch 0
processing epoch 1
processing epoch 2
processing epoch 3
processing epoch 4
processing epoch 5
processing epoch 6
processing epoch 7
processing epoch 8
processing epoch 9
processing epoch 10
processing epoch 11
processing epoch 12
processing epoch 13
processing epoch 14
processing epoch 15
processing epoch 16
processing epoch 17
processing epoch 18
processing epoch 19
```

Done. time taken : 0:04:28.700122

Evaluating the model with train data..

time taken : 0:00:12.766503

Train Data

RMSE : 0.6032438403305899

MAPE : 17.49285063490268

adding train results in the dictionary..

Evaluating for test data...

time taken : 0:00:00.252428

Test Data

RMSE : 1.0728491944183447

MAPE : 35.03817913919887

storing the test results in test dictionary...

Total time taken to run this algorithm : 0:04:41.722241

4.4.7 XgBoost with 13 features + Surprise Baseline + Surprise KNNbaseline + MF Techniques

Preparing Train data

In [122]:

```
# add the predicted values from both knns to this dataframe
reg_train['svd'] = models_evaluation_train['svd']['predictions']
reg_train['svdpp'] = models_evaluation_train['svdpp']['predictions']

reg_train.head(2)
```

Out[122]:

	user	movie	GAvg	sur1	sur2	sur3	sur4	sur5	smr1	smr2	...	smr4	smr5	UAV
0	53406	33	3.581679	4.0	5.0	5.0	4.0	1.0	5.0	2.0	...	3.0	1.0	3.3703
1	99540	33	3.581679	5.0	5.0	5.0	4.0	5.0	3.0	4.0	...	3.0	5.0	3.5555

2 rows × 21 columns

__Preparing Test data __

In [123]:

```
reg_test_df['svd'] = models_evaluation_test['svd']['predictions']
reg_test_df['svdpp'] = models_evaluation_test['svdpp']['predictions']

reg_test_df.head(2)
```

Out[123]:

	user	movie	GAvg	sur1	sur2	sur3	sur4	sur5	smr1	si
0	808635	71	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581
1	941866	71	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581

2 rows × 21 columns

Hyperparameter tuning

In [124]:

```

# prepare x_train and y_train
x_train = reg_train.drop(['user', 'movie', 'rating'], axis=1)
y_train = reg_train['rating']

# prepare test data
x_test = reg_test_df.drop(['user', 'movie', 'rating'], axis=1)
y_test = reg_test_df['rating']

xgb1 = xgb.XGBRegressor()
parameters = {'nthread': [1,2,3,4], #when use hyperthread, xgboost may become slower
              'learning_rate': [0.01,0.001,0.0001,0.03, 0.05, 0.07], #so called `eta` value
              'max_depth': [2, 3, 5, 6, 7,10, 12],
              'n_estimators': [20, 50, 100, 200, 300,500, 700, 1000]}

xgb_grid = GridSearchCV(xgb1,parameters,cv = 2,n_jobs = 5,verbose=True)

xgb_grid.fit(x_train,y_train)

print(xgb_grid.best_score_)
print(xgb_grid.best_params_)

```

Fitting 2 folds for each of 1344 candidates, totalling 2688 fits

```

[Parallel(n_jobs=5)]: Using backend LokyBackend with 5 concurrent workers.
[Parallel(n_jobs=5)]: Done 40 tasks      | elapsed: 2.2min
[Parallel(n_jobs=5)]: Done 190 tasks    | elapsed: 35.6min
[Parallel(n_jobs=5)]: Done 440 tasks    | elapsed: 197.3min
[Parallel(n_jobs=5)]: Done 790 tasks    | elapsed: 246.0min
[Parallel(n_jobs=5)]: Done 1240 tasks   | elapsed: 327.5min
[Parallel(n_jobs=5)]: Done 1790 tasks   | elapsed: 451.3min
[Parallel(n_jobs=5)]: Done 2440 tasks   | elapsed: 552.9min
[Parallel(n_jobs=5)]: Done 2688 out of 2688 | elapsed: 615.2min finished

```

```

[06:12:20] WARNING: C:/Jenkins/workspace/xgboost-win64_release_0.90/src/objective/regression_obj.cu:152: reg:linear is now deprecated in favor of reg:squarederror.

```

```
0.37353399234210544
```

```
{'learning_rate': 0.05, 'max_depth': 5, 'n_estimators': 200, 'nthread': 1}
```

In [126]:

```
xgb_final = xgb.XGBRegressor(silent=False, learning_rate= 0.05, max_depth= 5, n_estimators=
train_results, test_results = run_xgboost(xgb_final, x_train, y_train, x_test, y_test)
```

```
# store the results in models_evaluations dictionaries
```

```
models_evaluation_train['xgb_final'] = train_results
```

```
models_evaluation_test['xgb_final'] = test_results
```

```
xgb.plot_importance(xgb_final)
plt.show()
```

Training the model..

[06:51:08] WARNING: C:/Jenkins/workspace/xgboost-win64_release_0.90/src/objective/regression_obj.cu:152: reg:linear is now deprecated in favor of reg:squarederror.

Done. Time taken : 0:00:34.767999

Done

Evaluating the model with TRAIN data...

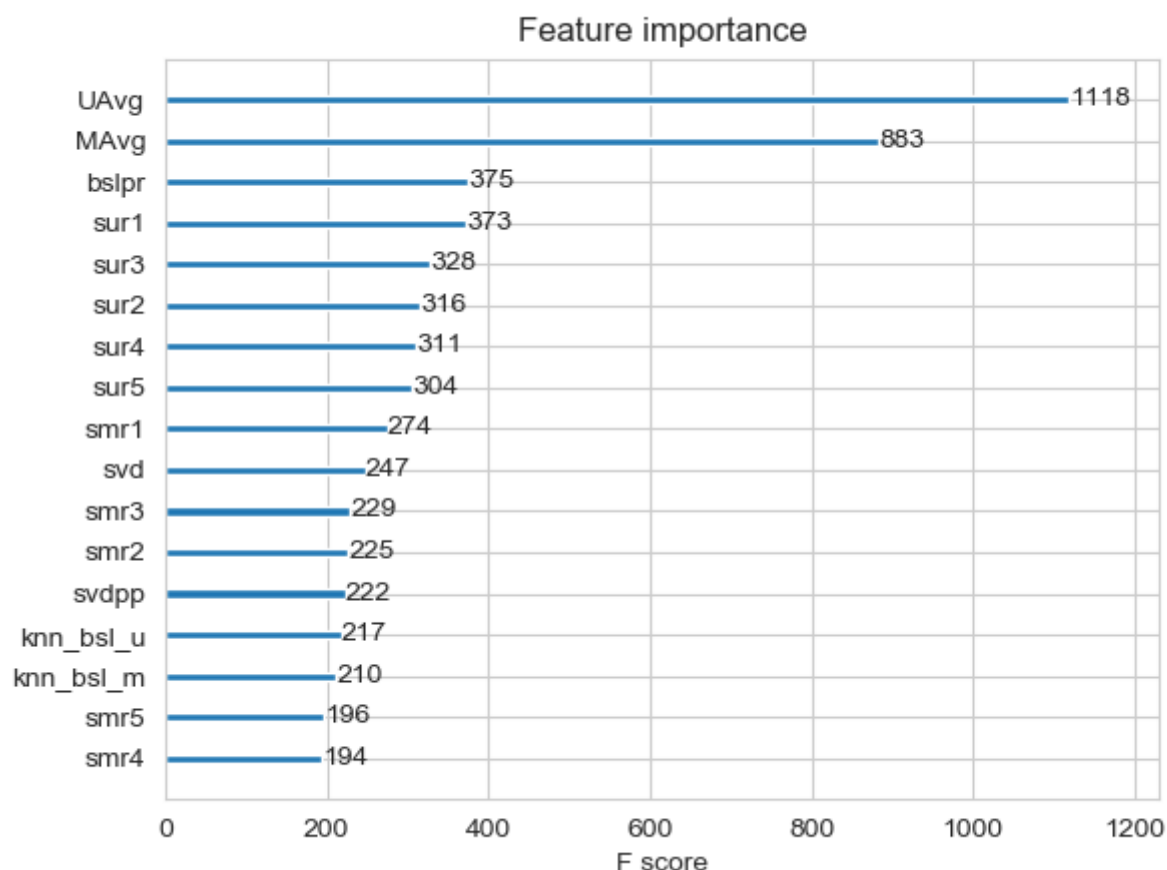
Evaluating Test data

TEST DATA

RMSE : 1.0777761952448184

MAPE : 34.36773182836531

Figure 12



4.4.8 XgBoost with Surprise Baseline + Surprise KNNbaseline + MF Techniques

Hyperparameter tuning

In [127]:

```
# prepare train data
x_train = reg_train[['knn_bsl_u', 'knn_bsl_m', 'svd', 'svdpp']]
y_train = reg_train['rating']

# test data
x_test = reg_test_df[['knn_bsl_u', 'knn_bsl_m', 'svd', 'svdpp']]
y_test = reg_test_df['rating']

xgb1 = xgb.XGBRegressor()
parameters = {'nthread': [1,2,3,4], #when use hyperthread, xgboost may become slower
              'learning_rate': [0.01,0.001,0.0001,0.03, 0.05, 0.07], #so called `eta` value
              'max_depth': [2, 3, 5, 6, 7,10, 12],
              'n_estimators': [20, 50, 100, 200, 300,500, 700, 1000]}

xgb_grid = GridSearchCV(xgb1,parameters,cv = 2,n_jobs = 5,verbose=True)

xgb_grid.fit(x_train,y_train)

print(xgb_grid.best_score_)
print(xgb_grid.best_params_)
```

Fitting 2 folds for each of 1344 candidates, totalling 2688 fits

```
[Parallel(n_jobs=5)]: Using backend LokyBackend with 5 concurrent workers.
[Parallel(n_jobs=5)]: Done 40 tasks      | elapsed: 36.3s
[Parallel(n_jobs=5)]: Done 190 tasks    | elapsed: 8.3min
[Parallel(n_jobs=5)]: Done 440 tasks    | elapsed: 33.6min
[Parallel(n_jobs=5)]: Done 790 tasks    | elapsed: 54.2min
[Parallel(n_jobs=5)]: Done 1240 tasks   | elapsed: 78.8min
[Parallel(n_jobs=5)]: Done 1790 tasks   | elapsed: 120.2min
[Parallel(n_jobs=5)]: Done 2440 tasks   | elapsed: 167.0min
[Parallel(n_jobs=5)]: Done 2688 out of 2688 | elapsed: 195.8min finished
```

```
[10:07:30] WARNING: C:/Jenkins/workspace/xgboost-win64_release_0.90/src/objective/regression_obj.cu:152: reg:linear is now deprecated in favor of reg:squarederror.
```

```
-0.002712634349758103
```

```
{'learning_rate': 0.03, 'max_depth': 2, 'n_estimators': 200, 'nthread': 1}
```

In [128]:

```

xgb_all_models = xgb.XGBRegressor(silent=False, learning_rate= 0.03, max_depth= 2, n_estimators= 1000)
train_results, test_results = run_xgboost(xgb_all_models, x_train, y_train, x_test, y_test)

# store the results in models_evaluations dictionaries
models_evaluation_train['xgb_all_models'] = train_results
models_evaluation_test['xgb_all_models'] = test_results

xgb.plot_importance(xgb_all_models)
plt.show()

```

Training the model..

[10:14:27] WARNING: C:/Jenkins/workspace/xgboost-win64_release_0.90/src/objective/regression_obj.cu:152: reg:linear is now deprecated in favor of reg:squarederror.

Done. Time taken : 0:00:07.768545

Done

Evaluating the model with TRAIN data...


Evaluating Test data

TEST DATA

RMSE : 1.075471898858508

MAPE : 35.012957238301325

Figure 13



4.5 Comparision between all models

In [130]:

```
models_evaluation_train
```

Out[130]:

```
{'first_algo': {'rmse': 0.8136983954128615,
  'mape': 23.96045466382881,
  'predictions': array([3.822063 , 4.10283 , 4.263296 , ..., 3.170802 , 3.6
471853,
  3.9148004], dtype=float32)},
'bsl_algo': {'rmse': 0.9347153928678286,
  'mape': 29.389572652358183,
  'predictions': array([3.89898205, 3.37140265, 3.3237082 , ..., 3.40556302,
3.54788112,
  3.6428903 ])},
'xgb_bsl': {'rmse': 0.8344119653430891,
  'mape': 24.75946475777376,
  'predictions': array([3.8615427, 4.160123 , 4.313834 , ..., 3.1586735, 3.6
19939 ,
  3.9573 ], dtype=float32)},
'knn_bsl_u': {'rmse': 0.33642097416508826,
  'mape': 9.145093375416348,
  'predictions': array([3.93001981, 3.17732964, 3.57250332, ..., 3.
4.
,
  4.
])},
'knn_bsl_m': {'rmse': 0.32584796251610554,
  'mape': 8.447062581998374,
  'predictions': array([3.86795834, 3.07630234, 3.40117708, ..., 3.
4.
,
  4.
])},
'xgb_knn_bsl': {'rmse': 0.8346283607932142,
  'mape': 24.785808523860204,
  'predictions': array([3.8881748, 4.1758423, 4.3292675, ..., 3.1644182, 3.6
282504,
  3.9577794], dtype=float32)},
'svd': {'rmse': 0.6574721240954099,
  'mape': 19.704901088660478,
  'predictions': array([4.30376383, 3.27896711, 3.38572269, ..., 3.25142545,
3.59289961,
  3.83325311])},
'svdp': {'rmse': 0.6032438403305899,
  'mape': 17.49285063490268,
  'predictions': array([3.62166401, 3.53824703, 2.87936055, ..., 2.87747167,
3.5273425 ,
  3.8486415 ])},
'xgb_final': {'rmse': 0.833468109564551,
  'mape': 24.74033115377659,
  'predictions': array([3.9052556, 4.1504154, 4.3200264, ..., 3.1457024, 3.6
164975,
  3.9588625], dtype=float32)},
'xgb_all_models': {'rmse': 1.0754320861991955,
  'mape': 35.2830757280161,
  'predictions': array([3.5706525, 3.5712123, 3.5626297, ..., 3.5712123, 3.5
631638,
  3.5774543], dtype=float32)}}}
```

In [131]:

```
models_evaluation_test
```

Out[131]:

```
{'first_algo': {'rmse': 1.0846515217011976,
  'mape': 33.98071389809766,
  'predictions': array([3.4303806, 3.4303806, 3.4303806, ..., 3.4303806, 3.4
303806,
    3.4303806], dtype=float32)},
'bsl_algo': {'rmse': 1.0730330260516174,
  'mape': 35.04995544572911,
  'predictions': array([3.58167938, 3.58167938, 3.58167938, ..., 3.58167938,
3.58167938,
    3.58167938])},
'xgb_bsl': {'rmse': 1.0775409822076403,
  'mape': 34.39214865567116,
  'predictions': array([3.490574, 3.490574, 3.490574, ..., 3.490574, 3.49057
4, 3.490574],
    dtype=float32)},
'knn_bsl_u': {'rmse': 1.0726493739667242,
  'mape': 35.02094499698424,
  'predictions': array([3.58167938, 3.58167938, 3.58167938, ..., 3.58167938,
3.58167938,
    3.58167938])},
'knn_bsl_m': {'rmse': 1.072758832653683,
  'mape': 35.02269653015042,
  'predictions': array([3.58167938, 3.58167938, 3.58167938, ..., 3.58167938,
3.58167938,
    3.58167938])},
'xgb_knn_bsl': {'rmse': 1.0773481404411078,
  'mape': 34.40629706111159,
  'predictions': array([3.4929602, 3.4929602, 3.4929602, ..., 3.4929602, 3.4
929602,
    3.4929602], dtype=float32)},
'svd': {'rmse': 1.0726046873826458,
  'mape': 35.01953535988152,
  'predictions': array([3.58167938, 3.58167938, 3.58167938, ..., 3.58167938,
3.58167938,
    3.58167938])},
'svdpp': {'rmse': 1.0728491944183447,
  'mape': 35.03817913919887,
  'predictions': array([3.58167938, 3.58167938, 3.58167938, ..., 3.58167938,
3.58167938,
    3.58167938])},
'xgb_final': {'rmse': 1.0777761952448184,
  'mape': 34.36773182836531,
  'predictions': array([3.4871788, 3.4871788, 3.4871788, ..., 3.4871788, 3.4
871788,
    3.4871788], dtype=float32)},
'xgb_all_models': {'rmse': 1.075471898858508,
  'mape': 35.012957238301325,
  'predictions': array([3.5596108, 3.5596108, 3.5596108, ..., 3.5596108, 3.5
596108,
    3.5596108], dtype=float32)}}
```

