

## In [1]:

1/49

Stack Overflow has over 4,000,000 registered users, and it exceeded 10,000,000 questions in late August 2015. Based on the type of tags assigned to questions, the top eight most discussed topics on the site are: Java, JavaScript, C#, PHP, Android, jQuery, Python and HTML.

## Problem Statement

Suggest the tags based on the content that was there in the question posted on Stackoverflow.

**Source:** <https://www.kaggle.com/c/facebook-recruiting-iii-keyword-extraction/>

## 1.2 Source / useful links

Data Source : <https://www.kaggle.com/c/facebook-recruiting-iii-keyword-extraction/data>  
(<https://www.kaggle.com/c/facebook-recruiting-iii-keyword-extraction/data>)

Youtube : <https://youtu.be/nNDqbUhtIRg> (<https://youtu.be/nNDqbUhtIRg>)

Research paper : <https://www.microsoft.com/en-us/research/wp-content/uploads/2016/02/tagging-1.pdf>  
(<https://www.microsoft.com/en-us/research/wp-content/uploads/2016/02/tagging-1.pdf>)

Research paper : <https://dl.acm.org/citation.cfm?id=2660970&dl=ACM&coll=DL> (<https://dl.acm.org/citation.cfm?id=2660970&dl=ACM&coll=DL>)

## 1.3 Real World / Business Objectives and Constraints

1. Predict as many tags as possible with high precision and recall.
2. Incorrect tags could impact customer experience on StackOverflow.
3. No strict latency constraints.

## 2. Machine Learning problem

### 2.1 Data

#### 2.1.1 Data Overview

Refer: <https://www.kaggle.com/c/facebook-recruiting-iii-keyword-extraction/data>  
(<https://www.kaggle.com/c/facebook-recruiting-iii-keyword-extraction/data>)

All of the data is in 2 files: Train and Test.

**Train.csv** contains 4 columns: Id,Title,Body,Tags.

**Test.csv** contains the same columns but without the Tags, which you are to predict.

**Size of Train.csv** - 6.75GB

**Size of Test.csv** - 2GB

**Number of rows in Train.csv** = 6034195

The questions are randomized and contains a mix of verbose text sites as well as sites related to math and programming. The number of questions from each site may vary, and no filtering has been performed on the questions (such as closed questions).

### Data Field Explanation

Dataset contains 6,034,195 rows. The columns in the table are:

**Id** - Unique identifier for each question

**Title** - The question's title

**Body** - The body of the question

**Tags** - The tags associated with the question in a space-separated format (all lowercase, should not contain tabs '\t' or ampersands '&')

### 2.1.2 Example Data point

**Title:** Implementing Boundary Value Analysis of Software Testing in a C++ program?

**Body :**

```

#include<
iostream>\n
#include<
stdlib.h>\n\n
using namespace std;\n\n
int main()\n
{\n
    int n,a[n],x,c,u[n],m[n],e[n][4];\n
    cout<<"Enter the number of variables";\n          cin>>n;\n
\n
    cout<<"Enter the Lower, and Upper Limits of the variable
s";\n

    for(int y=1; y<n+1; y++)\n
    {\n
        cin>>m[y];\n
        cin>>u[y];\n
    }\n
    for(x=1; x<n+1; x++)\n
    {\n
        a[x] = (m[x] + u[x])/2;\n
    }\n
    c=(n*4)-4;\n
    for(int a1=1; a1<n+1; a1++)\n
    {\n\n
        e[a1][0] = m[a1];\n
        e[a1][1] = m[a1]+1;\n
        e[a1][2] = u[a1]-1;\n
        e[a1][3] = u[a1];\n
    }\n
    for(int i=1; i<n+1; i++)\n
    {\n
        for(int l=1; l<=i; l++)\n
        {\n
            if(l!=1)\n
            {\n
                cout<<a[l]<<"\\t";\n
            }\n
        }\n
        for(int j=0; j<4; j++)\n
        {\n
            cout<<e[i][j];\n
            for(int k=0; k<n-(i+1); k++)\n
            {\n
                cout<<a[k]<<"\\t";\n
            }\n
            cout<<"\\n";\n
        }\n
    }\n
    }\n\n
    system("PAUSE");\n
    return 0;    \n

```

```
} \n
```

```
\n \n
```

<p>The answer should come in the form of a table like</p>\n\n

```
<pre><code>
```

```
1          50          50\n
2          50          50\n
99         50          50\n
100        50          50\n
50         1           50\n
50         2           50\n
50         99          50\n
50        100          50\n
50         50          1\n
50         50          2\n
50         50          99\n
50         50         100\n
```

```
</code></pre>\n\n
```

<p>if the no of inputs is 3 and their ranges are\n

```
1,100\n
```

```
1,100\n
```

```
1,100\n
```

(could be varied too)</p>\n\n

<p>The output is not coming,can anyone correct the code or tell me what  
's wrong?</p>\n'

Tags : 'c++ c'

## 2.2 Mapping the real-world problem to a Machine Learning Problem

### 2.2.1 Type of Machine Learning Problem

It is a multi-label classification problem

**Multi-label Classification:** Multilabel classification assigns to each sample a set of target labels. This can be thought as predicting properties of a data-point that are not mutually exclusive, such as topics that are relevant for a document. A question on Stackoverflow might be about any of C, Pointers, FileIO and/or memory-management at the same time or none of these.

\_\_\_Credit\_\_\_: <http://scikit-learn.org/stable/modules/multiclass.html>

### 2.2.2 Performance metric

**Micro-Averaged F1-Score (Mean F Score) :** The F1 score can be interpreted as a weighted average of the precision and recall, where an F1 score reaches its best value at 1 and worst score at 0. The relative contribution of precision and recall to the F1 score are equal. The formula for the F1 score is:

$$F1 = 2 * (precision * recall) / (precision + recall)$$

In the multi-class and multi-label case, this is the weighted average of the F1 score of each class.

#### 'Micro f1 score':

Calculate metrics globally by counting the total true positives, false negatives and false positives. This is a better metric when we have class imbalance.

#### 'Macro f1 score':

Calculate metrics for each label, and find their unweighted mean. This does not take label imbalance into account.

<https://www.kaggle.com/wiki/MeanFScore> (<https://www.kaggle.com/wiki/MeanFScore>)  
[http://scikit-learn.org/stable/modules/generated/sklearn.metrics.f1\\_score.html](http://scikit-learn.org/stable/modules/generated/sklearn.metrics.f1_score.html) ([http://scikit-learn.org/stable/modules/generated/sklearn.metrics.f1\\_score.html](http://scikit-learn.org/stable/modules/generated/sklearn.metrics.f1_score.html))

**Hamming loss** : The Hamming loss is the fraction of labels that are incorrectly predicted.

<https://www.kaggle.com/wiki/HammingLoss> (<https://www.kaggle.com/wiki/HammingLoss>)

## 3. Exploratory Data Analysis

### 3.1 Data Loading and Cleaning

#### 3.1.1 Using Pandas with SQLite to Load the data

In [2]:

```
#Creating db file from csv
#Learn SQL: https://www.w3schools.com/sql/default.asp
if not os.path.isfile('train.db'):
    start = datetime.now()
    disk_engine = create_engine('sqlite:///train.db')
    start = dt.datetime.now()
    chunksize = 180000
    j = 0
    index_start = 1
    for df in pd.read_csv('Train.csv', names=['Id', 'Title', 'Body', 'Tags'], chunksize=chunksize):
        df.index += index_start
        j+=1
        print('{} rows'.format(j*chunksize))
        df.to_sql('data', disk_engine, if_exists='append')
        index_start = df.index[-1] + 1
    print("Time taken to run this cell :", datetime.now() - start)
```

#### 3.1.2 Counting the number of rows

In [3]:

```

if os.path.isfile('train.db'):
    start = datetime.now()
    con = sqlite3.connect('train.db')
    num_rows = pd.read_sql_query("""SELECT count(*) FROM data""", con)
    #Always remember to close the database
    print("Number of rows in the database :", "\n", num_rows['count(*)'].values[0])
    con.close()
    print("Time taken to count the number of rows :", datetime.now() - start)
else:
    print("Please download the train.db file from drive or run the above cell to generate t

```

Number of rows in the database :

6034196

Time taken to count the number of rows : 0:00:04.153227

In [4]:

```
num_rows
```

Out[4]:

	count(*)
0	6034196

### 3.1.3 Checking for duplicates

In [5]:

```

#Learn SQL: https://www.w3schools.com/sql/default.asp
if os.path.isfile('train.db'):
    start = datetime.now()
    con = sqlite3.connect('train.db')
    df_no_dup = pd.read_sql_query('SELECT Title, Body, Tags, COUNT(*) as cnt_dup FROM data')
    con.close()
    print("Time taken to run this cell :", datetime.now() - start)
else:
    print("Please download the train.db file from drive or run the first to generate train.

```

Time taken to run this cell : 0:05:25.691910

In [6]:

```
df_no_dup.head()
# we can observe that there are duplicates
```

Out[6]:

	Title	Body	Tags	cnt_dup
0	Implementing Boundary Value Analysis of S...	<pre>&lt;pre&gt; &lt;code&gt;#include&lt;iosstream&gt;\n#include&lt;... </pre>	c++ c	1
1	Dynamic Datagrid Binding in Silverlight?	<p>I should do binding for datagrid dynamicall...</p>	c# silverlight data-binding	1
2	Dynamic Datagrid Binding in Silverlight?	<p>I should do binding for datagrid dynamicall...</p>	c# silverlight data-binding columns	1
3	java.lang.NoClassDefFoundError: javax/serv...	<p>I followed the guide in &lt;a href="http://sta...</p>	jsp jstl	1
4	java.sql.SQLException:[Microsoft][ODBC Dri...	<p>I use the following code&lt;/p&gt;\n\n&lt;pre&gt; &lt;code&gt;... </p>	java jdbc	2

In [7]:

```
print("number of duplicate questions :", num_rows['count(*)'].values[0]- df_no_dup.shape[0])
```

```
number of duplicate questions : 1827881 ( 30.292038906260256 % )
```

In [8]:

```
# number of times each question appeared in our database
df_no_dup.cnt_dup.value_counts()
```

Out[8]:

```
1    2656284
2    1272336
3     277575
4         90
5         25
6          5
Name: cnt_dup, dtype: int64
```

### 3.1.4 Removing the rows where no tags are present



In [13]:

```

sd=[]
start = datetime.now()
for i in range(df_no_dup.shape[0]):
    f=df_no_dup["Tags"][i]# no of characters==0
    if f==None:# when no tag given just remove that datapoint
        df_no_dup=df_no_dup.drop(i,axis=0)      # remove this datapoint
    else:
        d=len(df_no_dup["Tags"][i].split(" "))
        sd.append(d)

print(datetime.now()-start)

```

0:03:36.320016

In [14]:

```

start = datetime.now()
df_no_dup["tag_count"] = df_no_dup["Tags"].apply(lambda text: len(text.split(" ")))
# adding a new feature number of tags per question
print("Time taken to run this cell :", datetime.now() - start)
df_no_dup.head()

```

Time taken to run this cell : 0:00:02.329769

Out[14]:

	Title	Body	Tags	cnt_dup	t
0	Implementing Boundary Value Analysis of S...	<pre>#include<iosstream>\n#include<...	c++ c	1	
1	Dynamic Datagrid Binding in Silverlight?	<p>I should do binding for datagrid dynamicall...	c# silverlight data-binding	1	
2	Dynamic Datagrid Binding in Silverlight?	<p>I should do binding for datagrid dynamicall...	c# silverlight data-binding columns	1	
3	java.lang.NoClassDefFoundError: javax/serv...	<p>I followed the guide in <a href="http://sta...	jsp jstl	1	
4	java.sql.SQLException:[Microsoft][ODBC Dri...	<p>I use the following code</p>\n\n<pre><code>...	java jdbc	2	

In [15]:

```
# distribution of number of tags per question
df_no_dup.tag_count.value_counts()
```

Out[15]:

```
3    1206157
2    1111706
4     814996
1     568291
5     505158
Name: tag_count, dtype: int64
```

## Creating no Duplicate database

In [16]:

```
#Creating a new database with no duplicates
if not os.path.isfile('train_no_dup.db'):
    disk_dup = create_engine("sqlite:///train_no_dup.db")
    no_dup = pd.DataFrame(df_no_dup, columns=['Title', 'Body', 'Tags'])
    no_dup.to_sql('no_dup_train', disk_dup)
```

In [17]:

```
#This method seems more appropriate to work with this much data.
#creating the connection with database file.
if os.path.isfile('train_no_dup.db'):
    start = datetime.now()
    con = sqlite3.connect('train_no_dup.db')
    tag_data = pd.read_sql_query("""SELECT Tags FROM no_dup_train""", con)
    #Always remember to close the database
    con.close()

    # Let's now drop unwanted column.
    tag_data.drop(tag_data.index[0], inplace=True)
    #Printing first 5 columns from our data frame
    tag_data.head()
    print("Time taken to run this cell :", datetime.now() - start)
else:
    print("Please download the train.db file from drive or run the above cells to generate
```

Time taken to run this cell : 0:00:16.116694

## 3.2 Analysis of Tags

### 3.2.1 Total number of unique tags

In [18]:

```
# Importing & Initializing the "CountVectorizer" object, which
#is scikit-learn's bag of words tool.

#by default 'split()' will tokenize each tag using space.
vectorizer = CountVectorizer(tokenizer = lambda x: x.split())
# fit_transform() does two functions: First, it fits the model
# and learns the vocabulary; second, it transforms our training data
# into feature vectors. The input to fit_transform should be a list of strings.
tag_dtm = vectorizer.fit_transform(tag_data['Tags'])
#Returns:  X : array, [n_samples, n_features]
```

In [19]:

```
tag_data['Tags']
```

Out[19]:

```
1          c# silverlight data-binding
2      c# silverlight data-binding columns
3          jsp jstl
4          java jdbc
5      facebook api facebook-php-sdk
        ...
4206310      wordpress wordpress-plugin
4206311          php mysql text
4206312      php codeigniter character-encoding
4206313          php email outlook mime
4206314          html
Name: Tags, Length: 4206314, dtype: object
```

In [20]:

```
print("Number of data points :", tag_dtm.shape[0])
print("Number of unique tags :", tag_dtm.shape[1])
```

```
Number of data points : 4206314
Number of unique tags : 42048
```

In [21]:

```
#'get_feature_name()' gives us the vocabulary.
tags = vectorizer.get_feature_names()
#Lets look at the tags we have.
print("Some of the tags we have :", tags[:10])
```

```
Some of the tags we have : ['.a', '.app', '.asp.net-mvc', '.aspxauth', '.bas
h-profile', '.class-file', '.cs-file', '.doc', '.drv', '.ds-store']
```

### 3.2.3 Number of times a tag appeared

In [22]:

```
# https://stackoverflow.com/questions/15115765/how-to-access-sparse-matrix-elements
# Lets now store the document term matrix in a dictionary.
freqs = tag_dtm.sum(axis=0).A1
result = dict(zip(tags, freqs))
```

In [23]:

result

Out[23]:

```
{'.a': 18,
 '.app': 37,
 '.asp.net-mvc': 1,
 '.aspxauth': 21,
 '.bash-profile': 138,
 '.class-file': 53,
 '.cs-file': 14,
 '.doc': 47,
 '.drv': 1,
 '.ds-store': 8,
 '.each': 184,
 '.emf': 33,
 '.exe': 27,
 '.exe.config': 1,
 '.hgtags': 6,
 '.htaccess': 14884,
 '.htpasswd': 61,
 '.ico': 10.}
```

In [24]:

```
#Saving this dictionary to csv files.
if not os.path.isfile('data/tag_counts_dict_dtm.csv'):
    with open('data/tag_counts_dict_dtm.csv', 'w') as csv_file:
        writer = csv.writer(csv_file)
        for key, value in result.items():
            writer.writerow([key, value])
tag_df = pd.read_csv("data/tag_counts_dict_dtm.csv", names=['Tags', 'Counts'])
tag_df.head()
```

Out[24]:


	Tags	Counts
0	.a	18
1	.app	37
2	.asp.net-mvc	1
3	.aspxauth	21
4	.bash-profile	138

In [25]:

```
tag_df_sorted = tag_df.sort_values(['Counts'], ascending=False)
tag_counts = tag_df_sorted['Counts'].values
```


In [26]:

```
plt.plot(tag_counts)
plt.title("Distribution of number of times tag appeared questions")
plt.grid()
plt.xlabel("Tag number")
plt.ylabel("Number of times tag appeared")
plt.show()
```



In [27]:

```
plt.plot(tag_counts[0:10000])
plt.title('first 10k tags: Distribution of number of times tag appeared questions')
plt.grid()
plt.xlabel("Tag number")
plt.ylabel("Number of times tag appeared")
plt.show()
print(len(tag_counts[0:10000:25]), tag_counts[0:10000:25])
```




```
400 [331505  44829  22429  17728  13364  11162  10029   9148   8054   7151
    6466   5865   5370   4983   4526   4281   4144   3929   3750   3593
    3453   3299   3123   2989   2891   2738   2647   2527   2431   2331
    2259   2186   2097   2020   1959   1900   1828   1770   1723   1673
    1631   1574   1532   1479   1448   1406   1365   1328   1300   1266
    1245   1222   1197   1181   1158   1139   1121   1101   1076   1056
    1038   1023   1006   983    966   952    938   926    911    891
     882    869    856    841    830    816    804    789    779    770
     752    743    733    725    712    702    688    678    671    658
     650    643    634    627    616    607    598    589    583    577
     568    559    552    545    540    533    526    518    512    506
     500    495    490    485    480    477    469    465    457    450
     447    442    437    432    426    422    418    413    408    403
     398    393    388    385    381    378    374    370    367    365
     361    357    354    350    347    344    342    339    336    332
     330    326    323    319    315    312    309    307    304    301
     299    296    293    291    289    286    284    281    278    276
     275    272    270    268    265    262    260    258    256    254
     252    250    249    247    245    243    241    239    238    236
     234    233    232    230    228    226    224    222    220    219
     217    215    214    212    210    209    207    205    204    203
     201    200    199    198    196    194    193    192    191    189
     188    186    185    183    182    181    180    179    178    177
     175    174    172    171    170    169    168    167    166    165
     164    162    161    160    159    158    157    156    156    155
     154    153    152    151    150    149    149    148    147    146
     145    144    143    142    142    141    140    139    138    137
     137    136    135    134    134    133    132    131    130    130
     129    128    128    127    126    126    125    124    124    123
     123    122    122    121    120    120    119    118    118    117
     117    116    116    115    115    114    113    113    112    111
     111    110    109    109    108    108    107    106    106    106
     105    105    104    104    103    103    102    102    101    101
     100    100     99     99     98     98     97     97     96     96]
```

95	95	94	94	93	93	93	92	92	91
91	90	90	89	89	88	88	87	87	86
86	86	85	85	84	84	83	83	83	82
82	82	81	81	80	80	80	79	79	78
78	78	78	77	77	76	76	76	75	75
75	74	74	74	73	73	73	73	72	72]

In [28]:


```
plt.plot(tag_counts[0:1000])
plt.title('first 1k tags: Distribution of number of times tag appeared questions')
plt.grid()
plt.xlabel("Tag number")
plt.ylabel("Number of times tag appeared")
plt.show()
print(len(tag_counts[0:1000:5]), tag_counts[0:1000:5])
```



200	[331505	221533	122769	95160	62023	44829	37170	31897	26925	24537
22429	21820	20957	19758	18905	17728	15533	15097	14884	13703	
13364	13157	12407	11658	11228	11162	10863	10600	10350	10224	
10029	9884	9719	9411	9252	9148	9040	8617	8361	8163	
8054	7867	7702	7564	7274	7151	7052	6847	6656	6553	
6466	6291	6183	6093	5971	5865	5760	5577	5490	5411	
5370	5283	5207	5107	5066	4983	4891	4785	4658	4549	
4526	4487	4429	4335	4310	4281	4239	4228	4195	4159	
4144	4088	4050	4002	3957	3929	3874	3849	3818	3797	
3750	3703	3685	3658	3615	3593	3564	3521	3505	3483	
3453	3427	3396	3363	3326	3299	3272	3232	3196	3168	
3123	3094	3073	3050	3012	2989	2984	2953	2934	2903	
2891	2844	2819	2784	2754	2738	2726	2708	2681	2669	
2647	2621	2604	2594	2556	2527	2510	2482	2460	2444	
2431	2409	2395	2380	2363	2331	2312	2297	2290	2281	
2259	2246	2222	2211	2198	2186	2162	2142	2132	2107	
2097	2078	2057	2045	2036	2020	2011	1994	1971	1965	
1959	1952	1940	1932	1912	1900	1879	1865	1855	1841	
1828	1821	1813	1801	1782	1770	1760	1747	1741	1734	
1723	1707	1697	1688	1683	1673	1665	1656	1646	1639]	

In [29]:

```
plt.plot(tag_counts[0:500])
plt.title('first 500 tags: Distribution of number of times tag appeared questions')
plt.grid()
plt.xlabel("Tag number")
plt.ylabel("Number of times tag appeared")
plt.show()
print(len(tag_counts[0:500:5]), tag_counts[0:500:5])
```



```
100 [331505 221533 122769 95160 62023 44829 37170 31897 26925 24537
22429 21820 20957 19758 18905 17728 15533 15097 14884 13703
13364 13157 12407 11658 11228 11162 10863 10600 10350 10224
10029 9884 9719 9411 9252 9148 9040 8617 8361 8163
8054 7867 7702 7564 7274 7151 7052 6847 6656 6553
6466 6291 6183 6093 5971 5865 5760 5577 5490 5411
5370 5283 5207 5107 5066 4983 4891 4785 4658 4549
4526 4487 4429 4335 4310 4281 4239 4228 4195 4159
4144 4088 4050 4002 3957 3929 3874 3849 3818 3797
3750 3703 3685 3658 3615 3593 3564 3521 3505 3483]
```




In [30]:

```
plt.plot(tag_counts[0:100], c='b')
plt.scatter(x=list(range(0,100,5)), y=tag_counts[0:100:5], c='orange', label="quantiles with
# quantiles with 0.25 difference
plt.scatter(x=list(range(0,100,25)), y=tag_counts[0:100:25], c='m', label = "quantiles with

for x,y in zip(list(range(0,100,25)), tag_counts[0:100:25]):
    plt.annotate(s="({} , {})".format(x,y), xy=(x,y), xytext=(x-0.05, y+500))

plt.title('first 100 tags: Distribution of number of times tag appeared questions')
plt.grid()
plt.xlabel("Tag number")
plt.ylabel("Number of times tag appeared")
plt.legend()
plt.show()
print(len(tag_counts[0:100:5]), tag_counts[0:100:5])
```



```
20 [331505 221533 122769 95160 62023 44829 37170 31897 26925 24537
22429 21820 20957 19758 18905 17728 15533 15097 14884 13703]
```

In [31]:

```
# Store tags greater than 10K in one list
lst_tags_gt_10k = tag_df[tag_df.Counts>10000].Tags
#Print the length of the list
print ('{} Tags are used more than 10000 times'.format(len(lst_tags_gt_10k)))
# Store tags greater than 100K in one list
lst_tags_gt_100k = tag_df[tag_df.Counts>100000].Tags
#Print the length of the list.
print ('{} Tags are used more than 100000 times'.format(len(lst_tags_gt_100k)))
```

```
153 Tags are used more than 10000 times
14 Tags are used more than 100000 times
```

### Observations:

1. There are total 153 tags which are used more than 10000 times.
2. 14 tags are used more than 100000 times.
3. Most frequent tag (i.e. c#) is used 331505 times.
4. Since some tags occur much more frequently than others, Micro-averaged F1-score is the appropriate metric for this problem.

### 3.2.4 Tags Per Question

In [32]:

```
#Storing the count of tag in each question in list 'tag_count'
tag_quest_count = tag_dtm.sum(axis=1).tolist()
#Converting list of lists into single list, we will get [[3], [4], [2], [2], [3]] and we are
tag_quest_count=[int(j) for i in tag_quest_count for j in i]
print ('We have total {} datapoints.'.format(len(tag_quest_count)))

print(tag_quest_count[:5])
```

We have total 4206314 datapoints.

[3, 4, 2, 2, 3]

In [33]:

```
print( "Maximum number of tags per question: %d"%max(tag_quest_count))
print( "Minimum number of tags per question: %d"%min(tag_quest_count))
print( "Avg. number of tags per question: %f"% ((sum(tag_quest_count)*1.0)/len(tag_quest_count)))
```


Maximum number of tags per question: 5

Minimum number of tags per question: 1

Avg. number of tags per question: 2.899440

In [34]:

```
sns.countplot(tag_quest_count, palette='gist_rainbow')
plt.title("Number of tags in the questions ")
plt.xlabel("Number of Tags")
plt.ylabel("Number of questions")
plt.show()
```




#### Observations:

1. Maximum number of tags per question: 5
2. Minimum number of tags per question: 1
3. Avg. number of tags per question: 2.899
4. Most of the questions are having 2 or 3 tags



In [36]:

```
i=np.arange(30)
tag_df_sorted.head(30).plot(kind='bar')
plt.title('Frequency of top 20 tags')
plt.xticks(i, tag_df_sorted['Tags'])
plt.xlabel('Tags')
plt.ylabel('Counts')
plt.show()
```



### Observations:

1. Majority of the most frequent tags are programming language.
2. C# is the top most frequent programming language.
3. Android, IOS, Linux and windows are among the top most frequent operating systems.

## 3.3 Cleaning and preprocessing of Questions

### 3.3.1 Preprocessing

1. Sample 0.5M data points
2. Separate out code-snippets from Body
3. Remove Special characters from Question title and description (not in code)
4. Remove stop words (Except 'C')
5. Remove HTML Tags
6. Convert all the characters into small letters
7. Use SnowballStemmer to stem the words

In [37]:

```
def striphtml(data):  
    cleanr = re.compile('<.*?>')  
    cleantext = re.sub(cleanr, ' ', str(data))  
    return cleantext  
stop_words = set(stopwords.words('english'))  
stemmer = SnowballStemmer("english")
```

In [38]:

```

#http://www.sqlitetutorial.net/sqlite-python/create-tables/
def create_connection(db_file):
    """ create a database connection to the SQLite database
        specified by db_file
    :param db_file: database file
    :return: Connection object or None
    """
    try:
        conn = sqlite3.connect(db_file)
        return conn
    except Error as e:
        print(e)

    return None

def create_table(conn, create_table_sql):
    """ create a table from the create_table_sql statement
    :param conn: Connection object
    :param create_table_sql: a CREATE TABLE statement
    :return:
    """
    try:
        c = conn.cursor()
        c.execute(create_table_sql)
    except Error as e:
        print(e)

def checkTableExists(dbcon):
    cursr = dbcon.cursor()
    str = "select name from sqlite_master where type='table'"
    table_names = cursr.execute(str)
    print("Tables in the databse:")
    tables = table_names.fetchall()
    print(tables[0][0])
    return(len(tables))

def create_database_table(database, query):
    conn = create_connection(database)
    if conn is not None:
        create_table(conn, query)
        checkTableExists(conn)
    else:
        print("Error! cannot create the database connection.")
    conn.close()

sql_create_table = """CREATE TABLE IF NOT EXISTS QuestionsProcessed (question text NOT NULL
create_database_table("Processed.db", sql_create_table)

```

Tables in the databse:  
QuestionsProcessed

In [39]:

```
# http://www.sqlitetutorial.net/sqlite-delete/
# https://stackoverflow.com/questions/2279706/select-random-row-from-a-sqlite-table
start = datetime.now()
read_db = 'train_no_dup.db'
write_db = 'Processed.db'
if os.path.isfile(read_db):
    conn_r = create_connection(read_db)
    if conn_r is not None:
        reader = conn_r.cursor()
        reader.execute("SELECT Title, Body, Tags From no_dup_train ORDER BY RANDOM() LIMIT 1")

if os.path.isfile(write_db):
    conn_w = create_connection(write_db)
    if conn_w is not None:
        tables = checkTableExists(conn_w)
        writer = conn_w.cursor()
        if tables != 0:
            writer.execute("DELETE FROM QuestionsProcessed WHERE 1")
            print("Cleared All the rows")
print("Time taken to run this cell :", datetime.now() - start)
```

Tables in the databse:

QuestionsProcessed

Cleared All the rows

Time taken to run this cell : 0:03:32.619944

\_\_ we create a new data base to store the sampled and preprocessed questions \_\_

In [40]:

```
import nltk
nltk.download('punkt')
```

```
[nltk_data] Downloading package punkt to
[nltk_data] C:\Users\vansh\AppData\Roaming\nltk_data...
[nltk_data] Package punkt is already up-to-date!
```

Out[40]:

True

In [41]:

```
#http://www.bernzilla.com/2008/05/13/selecting-a-random-row-from-an-sqlite-table/

start = datetime.now()
preprocessed_data_list=[]
reader.fetchone()
questions_with_code=0
len_pre=0
len_post=0
questions_proccesed = 0
for row in reader:

    is_code = 0

    title, question, tags = row[0], row[1], row[2]

    if '<code>' in question:
        questions_with_code+=1
        is_code = 1
    x = len(question)+len(title)
    len_pre+=x

    code = str(re.findall(r'<code>(.*?)</code>', question, flags=re.DOTALL))

    question=re.sub('<code>(.*?)</code>', '', question, flags=re.MULTILINE|re.DOTALL)
    question=stripthtml(question.encode('utf-8'))

    title=title.encode('utf-8')

    question=str(title)+" "+str(question)
    question=re.sub(r'[^A-Za-z]+',' ',question)
    words=word_tokenize(str(question.lower()))

    #Removing all single letter and and stopwords from question exceptt for the letter 'c'
    question=' '.join(str(stemmer.stem(j)) for j in words if j not in stop_words and (len(j)

    len_post+=len(question)
    tup = (question,code,tags,x,len(question),is_code)
    questions_proccesed += 1
    writer.execute("insert into QuestionsProcessed(question,code,tags,words_pre,words_post,
if (questions_proccesed%100000==0):
    print("number of questions completed=",questions_proccesed)

no_dup_avg_len_pre=(len_pre*1.0)/questions_proccesed
no_dup_avg_len_post=(len_post*1.0)/questions_proccesed

print( "Avg. length of questions(Title+Body) before processing: %d"%no_dup_avg_len_pre)
print( "Avg. length of questions(Title+Body) after processing: %d"%no_dup_avg_len_post)
print( "Percent of questions containing code: %d"%((questions_with_code*100.0)/questions_pr

print("Time taken to run this cell :", datetime.now() - start)
```

```
number of questions completed= 100000
number of questions completed= 200000
number of questions completed= 300000
number of questions completed= 400000
Avg. length of questions(Title+Body) before processing: 1171
Avg. length of questions(Title+Body) after processing: 326
```



Percent of questions containing code: 57

Time taken to run this cell : 0:09:45.271985



In [42]:

```
# dont forget to close the connections, or else you will end up with Locks
conn_r.commit()
conn_w.commit()
conn_r.close()
conn_w.close()
```

In [43]:

```

if os.path.isfile(write_db):
    conn_r = create_connection(write_db)
    if conn_r is not None:
        reader = conn_r.cursor()
        reader.execute("SELECT question From QuestionsProcessed LIMIT 10")
        print("Questions after preprocessed")
        print('='*100)
        reader.fetchone()
        for row in reader:
            print(row)
            print('-'*100)
conn_r.commit()
conn_r.close()

```

Questions after preprocessed

```

=====
=====
('sql convers script delet bad data permiss bridg tabl goal write loop conve
rs go contain acucor securitycontain tabl clean permiss acucor securitypermi
ss bridg tabl base script care end two script complet fine notic differ two
clean queri vs saliva vs hair queri hair applic context secur queri saliva a
pplic context secur',)
-----
-----
('combin result multipl subqueri one set tri exclud row tabl base id tabl ni
tabl select result set like tri combin result subqueri one like filter big t
abl look group concat union kind seem find someth actual work anyon idea',)
-----
-----
('version percona db packag instal ubuntu natti narwhal new percona amp ubun
tu accord page percona avail use natti someon recommend version choos',)
-----
-----
('hide div view mvc develop mvc applic use razor syntax applic give comment
facil ad partial view load comment record db imag see comment box call run t
ime employe index view want remov div user click delet button work right ple
as check imag code applic updat script',)
-----
-----
('sudo anoth user without specifi usernam current tri creat sudoer file ran
someth figur end result look want user abl someth like instead run root like
script run user look sudoer file tri ad line like work specifi user easi way
script run specif user without specifi command line',)
-----
-----
('maximum size ellips net attempt use object draw ellips realli circl larg m
illion pixel though sinc pictur box draw reason size ought get either imag s
ometh approach line depend ellips locat use system draw graphic drawellips p
en rectanglef draw ellips success construct rectanglef describ dimens want e
xampl describ ellips much bigger actual visibl view space pass drawelips act
ual noth sinc part edg elips intersect visibl region instead tri draw ellips
bound produc method specifi ever throw microsoft document give indic inappro
pri argument big ellips draw far away center break',)
-----
-----
('move new line haskel updat follow function haskel must print sale week sal
e new line work way expect problem newlin charact code tri mani way new line
charact work expect everyth print line whichi want need help thank updat fol
low work compil error error come second line formatlin type decalar caus err

```

```
or need help',)
```

```
-----
('verita netbackup nbazd daemon wont start im new work verita netbackup inst
al server solari im plan use nbac problem nbazd daemon dont start go activ m
onitor daemon right click nbazd daemon mark start daemon dont anyth select d
aemon detail nbazd status stop start servic',)
-----
```

```
-----
('revers dns lookup command line like get list domain point certain ip addre
ss way get inform command line noth like host nslookup dig return hostnam ip
address help part want return edit inform nan exampl websit return inform ht
tp www domaintool com revers ip hostnam',)
-----
```

In [44]:

```
#Taking 0.5 Million entries to a dataframe.
write_db = 'Processed.db'
if os.path.isfile(write_db):
    conn_r = create_connection(write_db)
    if conn_r is not None:
        preprocessed_data = pd.read_sql_query("""SELECT question, Tags FROM QuestionsProces
conn_r.commit()
conn_r.close()
```

In [45]:

```
preprocessed_data.head()
```

Out[45]:

	question	tags
0	determin object lock synchron block java proce...	java synchronization
1	sql convers script delet bad data permiss brid...	sql-server-2008 tsql delete script
2	combin result multipl subqueri one set tri exc...	mysql set subquery
3	version percona db packag instal ubuntu natti ...	mysql ruby ubuntu natty percona
4	hide div view mvc develop mvc applic use razor...	javascript asp.net-mvc-3

In [46]:

```
print("number of data points in sample :", preprocessed_data.shape[0])
print("number of dimensions :", preprocessed_data.shape[1])
```

```
number of data points in sample : 499999
number of dimensions : 2
```

## 4. Machine Learning Models

### 4.1 Converting tags for multilabel problems

X	y1	y2	y3	y4
x1	0	1	1	0
x1	1	0	0	0
x1	0	1	0	0

In [47]:

```
# binary='true' will give a binary vectorizer
vectorizer = CountVectorizer(tokenizer = lambda x: x.split(), binary='true')
multilabel_y = vectorizer.fit_transform(preprocessed_data['tags'])
```

\_\_ We will sample the number of tags instead considering all of them (due to limitation of computing power) \_\_

In [48]:

```
def tags_to_choose(n):
    t = multilabel_y.sum(axis=0).tolist()[0]
    sorted_tags_i = sorted(range(len(t)), key=lambda i: t[i], reverse=True)
    multilabel_yn=multilabel_y[:,sorted_tags_i[:n]]
    return multilabel_yn

def questions_explained_fn(n):
    multilabel_yn = tags_to_choose(n)
    x= multilabel_yn.sum(axis=1)
    return (np.count_nonzero(x==0))
```

In [49]:

```
questions_explained = []
total_tags=multilabel_y.shape[1]
total_qs=preprocessed_data.shape[0]
for i in range(500, total_tags, 100):
    questions_explained.append(np.round(((total_qs-questions_explained_fn(i))/total_qs)*100))
```

In [50]:


questions\_explained

Out[50]:

```
[89.927,
 91.247,
 92.125,
 92.974,
 93.673,
 94.148,
 94.535,
 94.892,
 95.279,
 95.592,
 95.87,
 96.115,
 96.324,
 96.571,
 96.765,
 96.915,
 97.074,
 97.242.]
```

In [51]:

```
fig, ax = plt.subplots()
ax.plot(questions_explained)
xlabel = list(500+np.array(range(-50,450,50))*50)
ax.set_xticklabels(xlabel)
plt.xlabel("Number of tags")
plt.ylabel("Number Questions covered partially")
plt.grid()
plt.show()
# you can choose any number of tags based on your computing power, minimum is 50(it covers
print("with ",5500,"tags we are covering ",questions_explained[50],"% of questions")
```



with 5500 tags we are covering 99.048 % of questions

In [52]:

```

multilabel_yx = tags_to_choose(5500)
print("number of questions that are not covered :", questions_explained_fn(5500), "out of ",
number of questions that are not covered : 4760 out of 499999

```

In [53]:

```

print("Number of tags in sample :", multilabel_y.shape[1])
print("number of tags taken :", multilabel_yx.shape[1], "(", (multilabel_yx.shape[1]/multilab
Number of tags in sample : 30508
number of tags taken : 5500 ( 18.028058214238886 %)

```

\_\_ We consider top 15% tags which covers 99% of the questions \_\_

## 4.2 Split the data into test and train (80:20)

In [54]:

```

total_size=preprocessed_data.shape[0]
train_size=int(0.80*total_size)

x_train=preprocessed_data.head(train_size)
x_test=preprocessed_data.tail(total_size - train_size)

y_train = multilabel_yx[0:train_size,:]
y_test = multilabel_yx[train_size:total_size,:]

```

In [55]:

```

print("Number of data points in train data :", y_train.shape)
print("Number of data points in test data :", y_test.shape)

```

```

Number of data points in train data : (399999, 5500)
Number of data points in test data : (100000, 5500)

```

## 4.3 Featurizing data

In [56]:

```

start = datetime.now()
vectorizer = TfidfVectorizer(min_df=0.00009, max_features=200000, smooth_idf=True, norm="l2",
                             tokenizer = lambda x: x.split(), sublinear_tf=False, ngram_range=(1,1))
x_train_multilabel = vectorizer.fit_transform(x_train['question'])
x_test_multilabel = vectorizer.transform(x_test['question'])
print("Time taken to run this cell :", datetime.now() - start)

```

```

Time taken to run this cell : 0:02:36.205814

```

In [57]:

```
print("Dimensions of train data X:",x_train_multilabel.shape, "Y :",y_train.shape)
print("Dimensions of test data X:",x_test_multilabel.shape,"Y:",y_test.shape)
```

Dimensions of train data X: (399999, 89487) Y : (399999, 5500)

Dimensions of test data X: (100000, 89487) Y: (100000, 5500)

In [58]:

```
# https://www.analyticsvidhya.com/blog/2017/08/introduction-to-multi-label-classification/
#https://stats.stackexchange.com/questions/117796/scikit-multi-label-classification
# classifier = LabelPowerset(GaussianNB())
"""

from skmultilearn.adapt import MLkNN
classifier = MLkNN(k=21)

# train
classifier.fit(x_train_multilabel, y_train)

# predict
predictions = classifier.predict(x_test_multilabel)
print(accuracy_score(y_test,predictions))
print(metrics.f1_score(y_test, predictions, average = 'macro'))
print(metrics.f1_score(y_test, predictions, average = 'micro'))
print(metrics.hamming_loss(y_test,predictions))

"""

# we are getting memory error because the multilearn package
# is trying to convert the data into dense matrix
# -----
#MemoryError                                Traceback (most recent call last)
#<ipython-input-170-f0e7c7f3e0be> in <module>()
#----> classifier.fit(x_train_multilabel, y_train)
```

Out[58]:

```
"\nfrom skmultilearn.adapt import MLkNN\nnclassifier = MLkNN(k=21)\n\n# train\n\nclassifier.fit(x_train_multilabel, y_train)\n\n# predict\n\npredictions = cl
assifier.predict(x_test_multilabel)\n\nprint(accuracy_score(y_test,prediction
s))\n\nprint(metrics.f1_score(y_test, predictions, average = 'macro'))\n\nprint
(metrics.f1_score(y_test, predictions, average = 'micro'))\n\nprint(metrics.ha
mming_loss(y_test,predictions))\n\n"
```

In [59]:

```
x_train_multilabel.shape
```

Out[59]:

```
(399999, 89487)
```

In [60]:

```
x_test_multilabel.shape
```

Out[60]:

```
(100000, 89487)
```

## 4.4 Applying Logistic Regression with OneVsRest Classifier

In [64]:

```
classifier = OneVsRestClassifier(SGDClassifier(loss='log', alpha=0.00001, penalty='l1'), n_jobs=-1)
classifier.fit(x_train_multilabel, y_train)

predictions = classifier.predict(x_test_multilabel)

print("accuracy :", metrics.accuracy_score(y_test, predictions))
print("macro f1 score :", metrics.f1_score(y_test, predictions, average = 'macro'))
print("micro f1 score :", metrics.f1_score(y_test, predictions, average = 'micro'))
print("hamming loss :", metrics.hamming_loss(y_test, predictions))
print("Precision recall report :\n", metrics.classification_report(y_test, predictions))
```

```
accuracy : 0.08126
macro f1 score : 0.09291293571750123
micro f1 score : 0.37398320216916875
hamming loss : 0.0004130618181818182
Precision recall report :
```

	precision	recall	f1-score	support
0	0.61	0.22	0.32	7927
1	0.79	0.44	0.56	7236
2	0.82	0.54	0.65	6682
3	0.74	0.44	0.55	6298
4	0.95	0.76	0.84	5729
5	0.87	0.67	0.75	5244
6	0.73	0.31	0.43	3520
7	0.88	0.60	0.71	3150
8	0.69	0.38	0.49	3063
9	0.79	0.41	0.54	3055
10	0.84	0.60	0.70	2893
11	0.56	0.18	0.27	2889
12	0.54	0.20	0.35	2770

In [65]:

```
from sklearn.externals import joblib
joblib.dump(classifier, 'saved_models/lr_with_equal_weight_1.pkl')
```

Out[65]:

```
['saved_models/lr_with_equal_weight_1.pkl']
```

## 4.5 Modeling with less data points (0.5M data points) and more weight to title and 500 tags only.

In [66]:

```
sql_create_table = """CREATE TABLE IF NOT EXISTS QuestionsProcessed (question text NOT NULL, title text NOT NULL, tags text NOT NULL, PRIMARY KEY (question))"""
create_database_table("Titlmoreweight.db", sql_create_table)
```

Tables in the database:  
QuestionsProcessed



In [67]:

```
# http://www.sqlitetutorial.net/sqlite-delete/
# https://stackoverflow.com/questions/2279706/select-random-row-from-a-sqlite-table

read_db = 'train_no_dup.db'
write_db = 'Titlemoreweight.db'
train_datasize = 400000
if os.path.isfile(read_db):
    conn_r = create_connection(read_db)
    if conn_r is not None:
        reader = conn_r.cursor()
        # for selecting first 0.5M rows
        reader.execute("SELECT Title, Body, Tags From no_dup_train LIMIT 500001;")
        # for selecting random points
        #reader.execute("SELECT Title, Body, Tags From no_dup_train ORDER BY RANDOM() LIMIT 500001;")

if os.path.isfile(write_db):
    conn_w = create_connection(write_db)
    if conn_w is not None:
        tables = checkTableExists(conn_w)
        writer = conn_w.cursor()
        if tables != 0:
            writer.execute("DELETE FROM QuestionsProcessed WHERE 1")
            print("Cleared All the rows")
```

Tables in the database:  
QuestionsProcessed  
Cleared All the rows

## 4.5.1 Preprocessing of questions

1. Separate Code from Body
2. Remove Special characters from Question title and description (not in code)
3. **Give more weightage to title : Add title three times to the question**

- <li> Remove stop words (Except 'C') </li>
- <li> Remove HTML Tags </li>
- <li> Convert all the characters into small letters </li>
- <li> Use SnowballStemmer to stem the words </li>

In [68]:

```

#http://www.bernzilla.com/2008/05/13/selecting-a-random-row-from-an-sqlite-table/
start = datetime.now()
preprocessed_data_list=[]
reader.fetchone()
questions_with_code=0
len_pre=0
len_post=0
questions_proccesed = 0
for row in reader:

    is_code = 0

    title, question, tags = row[0], row[1], str(row[2])

    if '<code>' in question:
        questions_with_code+=1
        is_code = 1
    x = len(question)+len(title)
    len_pre+=x

    code = str(re.findall(r'<code>(.*?)</code>', question, flags=re.DOTALL))

    question=re.sub('<code>(.*?)</code>', '', question, flags=re.MULTILINE|re.DOTALL)
    question=striphtml(question.encode('utf-8'))

    title=title.encode('utf-8')

    # adding title three time to the data to increase its weight
    # add tags string to the training data

    question=str(title)+" "+str(title)+" "+str(title)+" "+question

#     if questions_proccesed<=train_datasize:
#         question=str(title)+" "+str(title)+" "+str(title)+" "+question+" "+str(tags)
#     else:
#         question=str(title)+" "+str(title)+" "+str(title)+" "+question

    question=re.sub(r'^A-Za-z0-9#+.\-]+', ' ', question)
    words=word_tokenize(str(question.lower()))

    #Removing all single letter and and stopwords from question exceptt for the letter 'c'
    question=' '.join(str(stemmer.stem(j)) for j in words if j not in stop_words and (len(j)

    len_post+=len(question)
    tup = (question,code,tags,x,len(question),is_code)
    questions_proccesed += 1
    writer.execute("insert into QuestionsProcessed(question,code,tags,words_pre,words_post,
if (questions_proccesed%100000==0):
    print("number of questions completed=",questions_proccesed)

no_dup_avg_len_pre=(len_pre*1.0)/questions_proccesed
no_dup_avg_len_post=(len_post*1.0)/questions_proccesed

print( "Avg. length of questions(Title+Body) before processing: %d"%no_dup_avg_len_pre)
print( "Avg. length of questions(Title+Body) after processing: %d"%no_dup_avg_len_post)
print( "Percent of questions containing code: %d"%((questions_with_code*100.0)/questions_pr

print("Time taken to run this cell :", datetime.now() - start)

```

```
number of questions completed= 100000  
number of questions completed= 200000  
number of questions completed= 300000  
number of questions completed= 400000  
number of questions completed= 500000  
Avg. length of questions(Title+Body) before processing: 1239  
Avg. length of questions(Title+Body) after processing: 424  
Percent of questions containing code: 57  
Time taken to run this cell : 0:13:05.637048
```

In [69]:

```
# never forget to close the conections or else we will end up with database locks  
conn_r.commit()  
conn_w.commit()  
conn_r.close()  
conn_w.close()
```

\_\_ Sample quesitons after preprocessing of data \_\_

In [70]:

```

if os.path.isfile(write_db):
    conn_r = create_connection(write_db)
    if conn_r is not None:
        reader = conn_r.cursor()
        reader.execute("SELECT question From QuestionsProcessed LIMIT 10")
        print("Questions after preprocessed")
        print('='*100)
        reader.fetchone()
        for row in reader:
            print(row)
            print('-'*100)
conn_r.commit()
conn_r.close()

```

Questions after preprocessed

```

=====
=====
('dynam datagrid bind silverlight dynam datagrid bind silverlight dynam data
grid bind silverlight bind datagrid dynam code wrote code debug code block s
eem bind correct grid come column form come grid column although necessari b
ind nthank repli advance..',)
-----
-----
('java.lang.noclassdeffounderror javax servlet jsp tagext taglibraryvalid ja
va.lang.noclassdeffounderror javax servlet jsp tagext taglibraryvalid java.l
ang.noclassdeffounderror javax servlet jsp tagext taglibraryvalid follow gui
d link instal jstl got follow error tri launch jsp page java.lang.noclassdef
founderror javax servlet jsp tagext taglibraryvalid taglib declar instal jst
l 1.1 tomcat webapp tri project work also tri version 1.2 jstl still messag
caus solv',)
-----
-----
('java.sql.sqllexcept microsoft odbc driver manag invalid descriptor index ja
va.sql.sqllexcept microsoft odbc driver manag invalid descriptor index java.s
ql.sqllexcept microsoft odbc driver manag invalid descriptor index use follow
code display caus solv',)
-----
-----
('better way updat feed fb php sdk better way updat feed fb php sdk better w
ay updat feed fb php sdk novic facebook api read mani tutori still confused.
i find post feed api method like correct second way use curl someth like way
better',)
-----
-----
('btnadd click event open two window record ad btnadd click event open two w
indow record ad btnadd click event open two window record ad open window sea
rch.aspx use code hav add button search.aspx nwhen insert record btnadd clic
k event open anoth window nafter insert record close window',)
-----
-----
('sql inject issu prevent correct form submiss php sql inject issu prevent c
orrect form submiss php sql inject issu prevent correct form submiss php che
ck everyth think make sure input field safe type sql inject good news safe b
ad news one tag mess form submiss place even touch life figur exact html use
templat file forgiv okay entir php script get execut see data post none foru
m field post problem use someth titl field none data get post current use pr
int post see submit noth work flawless statement though also mention script
work flawless local machin use host come across problem state list input tes
t mess',)

```

```
-----
('countabl subaddit lebesgu measur countabl subaddit lebesgu measur countabl
subaddit lebesgu measur let lbrace rbrace sequenc set sigma -algebra mathcal
want show left bigcup right leq sum left right countabl addit measur defin s
et sigma algebra mathcal think use monoton properti somewher proof start app
reci littl help nthank ad han answer make follow addit construct given han a
nswer clear bigcup bigcup cap emptyset neq left bigcup right left bigcup rig
ht sum left right also construct subset monoton left right leq left right fi
nal would sum leq sum result follow',)
-----
```

```
-----
('hql equival sql queri hql equival sql queri hql equival sql queri hql quer
i replac name class properti name error occur hql error',)
-----
```

```
-----
('undefin symbol architectur i386 objc class skpsmtpmessag referenc error un
defin symbol architectur i386 objc class skpsmtpmessag referenc error undefi
n symbol architectur i386 objc class skpsmtpmessag referenc error import fra
mework send email applic background import framework i.e skpsmtpmessag someb
odi suggest get error collect2 ld return exit status import framework correc
t sorc taken framework follow mfmcomposeviewcontrol question lock field u
pdat answer drag drop folder project click copi nthat',)
-----
```

\_\_ Saving Preprocessed data to a Database \_\_

In [71]:

```
#Taking 0.5 Million entries to a dataframe.
write_db = 'Titlemoreweight.db'
if os.path.isfile(write_db):
    conn_r = create_connection(write_db)
    if conn_r is not None:
        preprocessed_data = pd.read_sql_query("""SELECT question, Tags FROM QuestionsProces
conn_r.commit()
conn_r.close()
```

In [72]:

```
preprocessed_data.head()
```

Out[72]:

	question	tags
0	dynam datagrid bind silverlight dynam datagrid...	c# silverlight data-binding
1	dynam datagrid bind silverlight dynam datagrid...	c# silverlight data-binding columns
2	java.lang.noclassdeffounderror javax servlet j...	jsp jstl
3	java.sql.sqlexcept microsoft odbc driver manag...	java jdbc
4	better way updat feed fb php sdk better way up...	facebook api facebook-php-sdk

In [73]:

```
print("number of data points in sample :", preprocessed_data.shape[0])
print("number of dimensions :", preprocessed_data.shape[1])
```

number of data points in sample : 500000  
 number of dimensions : 2

\_\_ Converting string Tags to multilable output variables \_\_

In [74]:

```
vectorizer = CountVectorizer(tokenizer = lambda x: x.split(), binary='true')
multilabel_y = vectorizer.fit_transform(preprocessed_data['tags'])
```


\_\_ Selecting 500 Tags \_\_

In [75]:

```
questions_explained = []
total_tags=multilabel_y.shape[1]
total_qs=preprocessed_data.shape[0]
for i in range(500, total_tags, 100):
    questions_explained.append(np.round(((total_qs-questions_explained_fn(i))/total_qs)*100
```

In [76]:

```
fig, ax = plt.subplots()
ax.plot(questions_explained)
xlabel = list(500+np.array(range(-50,450,50))*50)
ax.set_xticklabels(xlabel)
plt.xlabel("Number of tags")
plt.ylabel("Number Questions covered partially")
plt.grid()
plt.show()
# you can choose any number of tags based on your computing power, minimun is 500(it covers
print("with ",5500,"tags we are covering ",questions_explained[50],"% of questions")
print("with ",500,"tags we are covering ",questions_explained[0],"% of questions")
```



with 5500 tags we are covering 99.157 % of questions  
 with 500 tags we are covering 90.956 % of questions

In [77]:

```
# we will be taking 500 tags
multilabel_yx = tags_to_choose(500)
print("number of questions that are not covered :", questions_explained_fn(500),"out of ",
```

number of questions that are not covered : 45221 out of 500000

In [78]:

```
x_train=preprocessed_data.head(train_datasize)
x_test=preprocessed_data.tail(preprocessed_data.shape[0] - 400000)

y_train = multilabel_yx[0:train_datasize,:]
y_test = multilabel_yx[train_datasize:preprocessed_data.shape[0],:]
```

In [79]:

```
print("Number of data points in train data :", y_train.shape)
print("Number of data points in test data :", y_test.shape)
```

Number of data points in train data : (400000, 500)

Number of data points in test data : (100000, 500)

## 4.5.2 Featurizing data with Tfidf vectorizer

In [80]:

```
start = datetime.now()
vectorizer = TfidfVectorizer(min_df=0.00009, max_features=200000, smooth_idf=True, norm="l2",
                             tokenizer = lambda x: x.split(), sublinear_tf=False, ngram_range=(1,1))
x_train_multilabel = vectorizer.fit_transform(x_train['question'])
x_test_multilabel = vectorizer.transform(x_test['question'])
print("Time taken to run this cell :", datetime.now() - start)
```

Time taken to run this cell : 0:03:58.877800

In [81]:

```
print("Dimensions of train data X:",x_train_multilabel.shape, "Y :",y_train.shape)
print("Dimensions of test data X:",x_test_multilabel.shape,"Y:",y_test.shape)
```

Dimensions of train data X: (400000, 94927) Y : (400000, 500)

Dimensions of test data X: (100000, 94927) Y: (100000, 500)

## 4.5.3 Applying Logistic Regression with OneVsRest Classifier

In [82]:

```

start = datetime.now()
classifier = OneVsRestClassifier(SGDClassifier(loss='log', alpha=0.00001, penalty='l1'), n_
classifier.fit(x_train_multilabel, y_train)
predictions = classifier.predict (x_test_multilabel)

print("Accuracy :",metrics.accuracy_score(y_test, predictions))
print("Hamming loss ",metrics.hamming_loss(y_test,predictions))

precision = precision_score(y_test, predictions, average='micro')
recall = recall_score(y_test, predictions, average='micro')
f1 = f1_score(y_test, predictions, average='micro')

print("Micro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision, recall, f1))

precision = precision_score(y_test, predictions, average='macro')
recall = recall_score(y_test, predictions, average='macro')
f1 = f1_score(y_test, predictions, average='macro')

print("Macro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision, recall, f1))

print (metrics.classification_report(y_test, predictions))
print("Time taken to run this cell :", datetime.now() - start)

```

```

Accuracy : 0.23614
Hamming loss  0.00278378
Micro-average quality numbers
Precision: 0.7208, Recall: 0.3251, F1-measure: 0.4481
Macro-average quality numbers
Precision: 0.5467, Recall: 0.2564, F1-measure: 0.3331

```

	precision	recall	f1-score	support
0	0.94	0.64	0.76	5519
1	0.69	0.26	0.38	8190
2	0.81	0.38	0.52	6529
3	0.81	0.43	0.56	3231
4	0.81	0.41	0.54	6430
5	0.82	0.34	0.48	2879
6	0.87	0.50	0.63	5086
7	0.88	0.54	0.67	4533
8	0.60	0.13	0.22	3000
9	0.81	0.52	0.63	2765
10	0.59	0.16	0.25	3051
11	0.60	0.22	0.35	3000

In [83]:

```

joblib.dump(classifier, 'saved_models/lr_with_more_title_weight.pkl')

```

Out[83]:

```

['lr_with_more_title_weight.pkl']

```

**Task 1: Use bag of words upto 4 grams and compute the micro f1 score with Logistic regression(OvR)**



In [85]:

```
start = datetime.now()
vectorizer = CountVectorizer(min_df=0.00009, max_features=200000, ngram_range=(1,4),tokeniz
x_train_multilabel = vectorizer.fit_transform(x_train['question'])
x_test_multilabel = vectorizer.transform(x_test['question'])
print("Time taken to run this cell :", datetime.now() - start)
```

Time taken to run this cell : 4:04:07.910078

In [86]:

```
print("Dimensions of train data X:",x_train_multilabel.shape, "Y :",y_train.shape)
print("Dimensions of test data X:",x_test_multilabel.shape,"Y:",y_test.shape)
```

Dimensions of train data X: (400000, 95585) Y : (400000, 500)

Dimensions of test data X: (100000, 95585) Y: (100000, 500)

## Logistic regression (BoW (1,4) gram without hyperparameter tuning

In [87]:

```

# this will be taking so much time try not to run it, download the lr_with_equal_weight.pkl
# This takes about 6-7 hours to run.
classifier = OneVsRestClassifier(SGDClassifier(loss='log', alpha=0.00001, penalty='l1'))
classifier.fit(x_train_multilabel, y_train)
predictions = classifier.predict(x_test_multilabel)
print("Accuracy :",metrics.accuracy_score(y_test, predictions))
print("Hamming loss ",metrics.hamming_loss(y_test,predictions))

precision = precision_score(y_test, predictions, average='micro')
recall = recall_score(y_test, predictions, average='micro')
f1 = f1_score(y_test, predictions, average='micro')

print("Micro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision, recall, f1))

precision = precision_score(y_test, predictions, average='macro')
recall = recall_score(y_test, predictions, average='macro')
f1 = f1_score(y_test, predictions, average='macro')

print("Macro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision, recall, f1))

print(metrics.classification_report(y_test, predictions))
print("Time taken to run this cell :", datetime.now() - start)

```

478	0.30	0.88	0.45	24
479	0.07	0.26	0.11	92
480	0.31	0.60	0.41	100
481	0.26	0.49	0.33	103
482	0.08	0.32	0.13	74
483	0.49	0.67	0.57	105
484	0.06	0.19	0.09	83
485	0.03	0.11	0.04	82
486	0.06	0.25	0.10	71
487	0.17	0.32	0.22	120
488	0.04	0.11	0.06	105
489	0.21	0.52	0.29	87
490	0.29	0.78	0.42	32
491	0.02	0.09	0.04	69
492	0.05	0.14	0.07	49
493	0.04	0.15	0.07	117
494	0.09	0.23	0.13	61
495	0.84	0.86	0.85	344
496	0.12	0.23	0.16	52
497	0.10	0.42	0.26	127

In [88]:

```
joblib.dump(classifier, 'lr_bow_without_hyp_tuned.pkl')
```

Out[88]:

```
['lr_bow_without_hyp_tuned.pkl']
```

## Task2: Hyperparameter tuning (GridsearchCV)

In [123]:

```
#https://stackoverflow.com/questions/12632992/gridsearch-for-an-estimator-inside-a-onevsres
from sklearn.multiclass import OneVsRestClassifier
from sklearn.model_selection import GridSearchCV

clf = OneVsRestClassifier(SGDClassifier(loss='log',penalty='l1'))

parameters = {
    "estimator__alpha": [0.0001,0.001,0.01,0.1,1,10]
}

model_tunning = GridSearchCV(clf, param_grid=parameters)

model_tunning.fit(x_train_multilabel, y_train)
```

Out[123]:

```
GridSearchCV(cv='warn', error_score='raise-deprecating',
             estimator=OneVsRestClassifier(estimator=SGDClassifier(alpha=0.0
001,
                                                                    average=F
                                                                    class_wi
                                                                    early_sto
                                                                    epsilon=
                                                                    eta0=0.0,
                                                                    fit_inter
                                                                    l1_ratio=
                                                                    learning_
                                                                    loss='lo
                                                                    max_iter=
                                                                    n_iter_no
                                                                    n_jobs=No
                                                                    penalty
                                                                    power_t=
                                                                    random_st
                                                                    shuffle=T
                                                                    tol=0.00
                                                                    validatio
                                                                    verbose=
                                                                    warm_star
                                                                    n_jobs=None),
```

```

iid='warn', n_jobs=None,
param_grid={'estimator__alpha': [0.0001, 0.001, 0.01, 0.1, 1, 1
0]},
pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
scoring=None, verbose=0)

```

In [125]:

```

print (model_tunning.best_score_)
print (model_tunning.best_params_)

```

```

0.189835
{'estimator__alpha': 0.001}

```

In [126]:

```

classifier = OneVsRestClassifier(SGDClassifier(loss='log', alpha=0.001, penalty='l1'))
classifier.fit(x_train_multilabel, y_train)
predictions = classifier.predict(x_test_multilabel)
print("Accuracy :",metrics.accuracy_score(y_test, predictions))
print("Hamming loss ",metrics.hamming_loss(y_test,predictions))

precision = precision_score(y_test, predictions, average='micro')
recall = recall_score(y_test, predictions, average='micro')
f1 = f1_score(y_test, predictions, average='micro')

print("Micro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision, recall, f1))

precision = precision_score(y_test, predictions, average='macro')
recall = recall_score(y_test, predictions, average='macro')
f1 = f1_score(y_test, predictions, average='macro')

print("Macro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision, recall, f1))

print (metrics.classification_report(y_test, predictions))
print("Time taken to run this cell :", datetime.now() - start)

```

```

Accuracy : 0.19534
Hamming loss 0.00310584
Micro-average quality numbers
Precision: 0.6068, Recall: 0.3028, F1-measure: 0.4040
Macro-average quality numbers
Precision: 0.4187, Recall: 0.2349, F1-measure: 0.2821

```

	precision	recall	f1-score	support
0	0.88	0.62	0.73	5519
1	0.58	0.18	0.28	8190
2	0.77	0.31	0.44	6529
3	0.79	0.36	0.49	3231
4	0.78	0.38	0.51	6430
5	0.77	0.30	0.43	2879
6	0.86	0.47	0.61	5086
7	0.87	0.51	0.65	4533
8	0.55	0.14	0.22	3000
9	0.79	0.48	0.60	2765
10	0.59	0.14	0.23	3051
11	0.71	0.32	0.42	3000

	precision	recall	f1-score	support
0	0.88	0.62	0.73	5519
1	0.58	0.18	0.28	8190
2	0.77	0.31	0.44	6529
3	0.79	0.36	0.49	3231
4	0.78	0.38	0.51	6430
5	0.77	0.30	0.43	2879
6	0.86	0.47	0.61	5086
7	0.87	0.51	0.65	4533
8	0.55	0.14	0.22	3000
9	0.79	0.48	0.60	2765
10	0.59	0.14	0.23	3051
11	0.71	0.32	0.42	3000

In [127]:

```
joblib.dump(classifier, 'lr_bow_with_hyp_tuned.pkl')
```

Out[127]:

```
['lr_bow_with_hyp_tuned.pkl']
```

**Task3: Try OneVsRestClassifier with Linear-SVM (SGDClassifier with loss-hinge)**

In [128]:

```

clf = OneVsRestClassifier(SGDClassifier(loss='hinge',penalty='l1'))

parameters = {
    "estimator__alpha": [0.0001,0.001,0.01,0.1,1,10]
}

model_tunning = GridSearchCV(clf, param_grid=parameters)

model_tunning.fit(x_train_multilabel, y_train)

```

Out[128]:

```

GridSearchCV(cv='warn', error_score='raise-deprecating',
             estimator=OneVsRestClassifier(estimator=SGDClassifier(alpha=
0.0001,
                                average
=False,
                                class_w
eight=None,
                                early_s
topping=False,
                                epsilon
=0.1,
                                eta0=0.
0,
                                fit_int
ercept=True,
                                l1_rati
o=0.15,
                                learnin
g_rate='optimal',
                                loss='h
inge',
                                max_ite
r=1000,
                                n_iter_
no_change=5,
                                n_jobs=
None,
                                penalty
='l1',
                                power_t
=0.5,
                                random_
state=None,
                                shuffle
=True,
                                tol=0.0
01,
                                validat
ion_fraction=0.1,
                                verbose
=0,
                                warm_st
art=False),
                                n_jobs=None),
             iid='warn', n_jobs=None,
             param_grid={'estimator__alpha': [0.0001, 0.001, 0.01, 0.1, 1,
10]}),

```

```
pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
e,
scoring=None, verbose=0)
```

In [129]:

```
print(model_tunning.best_score_)
print(model_tunning.best_params_)
```

```
0.1796775
{'estimator__alpha': 0.001}
```

In [130]:

```
classifier = OneVsRestClassifier(SGDClassifier(loss='hinge', alpha=0.001, penalty='l1'))
classifier.fit(x_train_multilabel, y_train)
predictions = classifier.predict(x_test_multilabel)
print("Accuracy :", metrics.accuracy_score(y_test, predictions))
print("Hamming loss ", metrics.hamming_loss(y_test, predictions))

precision = precision_score(y_test, predictions, average='micro')
recall = recall_score(y_test, predictions, average='micro')
f1 = f1_score(y_test, predictions, average='micro')

print("Micro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision, recall, f1))

precision = precision_score(y_test, predictions, average='macro')
recall = recall_score(y_test, predictions, average='macro')
f1 = f1_score(y_test, predictions, average='macro')

print("Macro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision, recall, f1))

print(metrics.classification_report(y_test, predictions))
print("Time taken to run this cell :", datetime.now() - start)
```

```
Accuracy : 0.19507
Hamming loss 0.00310382
Micro-average quality numbers
Precision: 0.6090, Recall: 0.2994, F1-measure: 0.4014
Macro-average quality numbers
Precision: 0.3197, Recall: 0.2267, F1-measure: 0.2492
```

	precision	recall	f1-score	support
0	0.84	0.63	0.72	5519
1	0.52	0.18	0.26	8190
2	0.80	0.30	0.44	6529
3	0.72	0.40	0.52	3231
4	0.83	0.32	0.46	6430
5	0.71	0.40	0.51	2879
6	0.79	0.54	0.64	5086
7	0.81	0.60	0.69	4533
8	0.52	0.18	0.27	3000
9	0.71	0.53	0.60	2765
10	0.32	0.01	0.02	3051
11	0.60	0.24	0.45	3000

In [131]:

```
joblib.dump(classifier, 'lr_svm.pkl')
```

Out[131]:

```
['lr_svm.pkl']
```

## Conclusion:

In [133]:

```
from prettytable import PrettyTable
t = PrettyTable()
t.field_names= ("Featurization method", "Model", "Macro Averaged F1 Score", "micro f1 score"
t.add_row(["tf-idf", "Logistic Regression", 0.09, 0.37, 0.0004])
t.add_row(["tf-idf", "Logistic Regression", 0.33, 0.44, 0.002])
t.add_row(["Bow", "Logistic Regression", 0.27, 0.36, 0.005])
t.add_row(["Bow", "Logistic Regression", 0.28, 0.40, 0.003])
t.add_row(["Bow", "Liner SVM classifier", 0.24, 0.40, 0.003])
print(t.get_string(titles = "Obeservations"))
```

```
+-----+-----+-----+-----+
| Featurization method | Model | Macro Averaged F1 Score | micro f1 score | hamming loss |
+-----+-----+-----+-----+
| tf-idf | Logistic Regression | 0.09 | 0.37 | 0.0004 |
| tf-idf | Logistic Regression | 0.33 | 0.44 | 0.002 |
| Bow | Logistic Regression | 0.27 | 0.36 | 0.005 |
| Bow | Logistic Regression | 0.28 | 0.40 | 0.003 |
| Bow | Liner SVM classifier | 0.24 | 0.40 | 0.003 |
+-----+-----+-----+-----+
```

## Procedure:

- Load the data to the sqlite database.
- Removing the duplicates rows and loading the data in a new database.
- Analysis on tags (Frequency of each tag).
- Text preprocessing removing HTML tags, stopwords except "c" and stemming.
- Finding tags the covers 99.08% questions --5500 tags.
- Applied on Models:
  - M1. Logistic regression(OvR) tfidf featurization (5500 tags)
  - M2. Logistic regression(OvR) tfidf featurization and (more weight to title & 5500 tags)
- Adding more weight to title and Finding tags the covers 90.0% questions-- 500 tags
- Applied on Models:
  - M3. Logistic regression(OvR) BoW featurization (more weight to title & 500 tags)



M4. Logistic regression(OvR) BOW featurization with Hyperparameter tuning (more weight to title & 500 tags)

M5. Linear SVM BOW featurization with Hyperparameter tuning (more weight to title & 500 tags)

- Summerizing all the models using pretty table